

# Buffered vs Unbuffered Channels

## The Critical Decision

Choosing between buffered and unbuffered channels affects:

- Synchronization guarantees
- Throughput and latency
- Goroutine coupling
- Deadlock likelihood
- Resource consumption

This is not a performance optimization question. This is a correctness question.

## Unbuffered Channels: Synchronous Rendezvous

```
ch := make(chan int) // No buffer, capacity 0
```

### Behavior

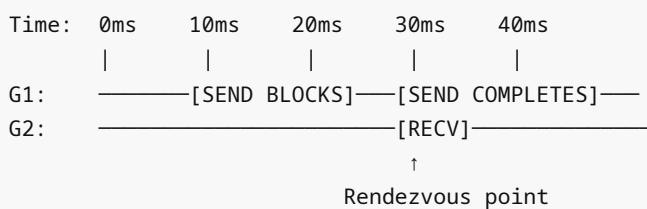
**Send and receive must meet**—neither can proceed without the other.

```
ch := make(chan int)

// Goroutine 1
ch <- 42 // Blocks until goroutine 2 receives

// Goroutine 2
val := <-ch // Blocks until goroutine 1 sends
```

### Execution Timeline



Both goroutines **synchronize at the channel operation**.

### Happens-Before Guarantee

For unbuffered channels:

```
Send happens-before receive completes
```

**Example:**

```

var data string
done := make(chan bool)

go func() {
    data = "hello"      // (1) Write
    done <- true        // (2) Send
}()

<-done           // (3) Receive
fmt.Println(data) // (4) Read

// Guaranteed order: (1) → (2) → (3) → (4)
// Reading `data` is safe

```

## When to Use Unbuffered

**Use unbuffered when:**

- You need **exact synchronization** between goroutines
- You want to **signal events** (completion, errors)
- You're **transferring ownership** and want confirmation of receipt
- You want **backpressure** (sender waits for receiver)
- You're implementing **request-response** patterns

### Example: Event Signaling

```

func worker(ready chan struct{}) {
    // Initialize
    doSetup()
    close(ready) // Signal "I'm ready"
}

func main() {
    ready := make(chan struct{}) // Unbuffered
    go worker(ready)
    <-ready // Wait until worker is ready
    // Now safe to proceed
}

```

## Buffered Channels: Asynchronous Communication

```
ch := make(chan int, 10) // Buffer capacity 10
```

### Behavior

**Send blocks only when buffer full; receive blocks only when buffer empty.**

```
ch := make(chan int, 2)
```

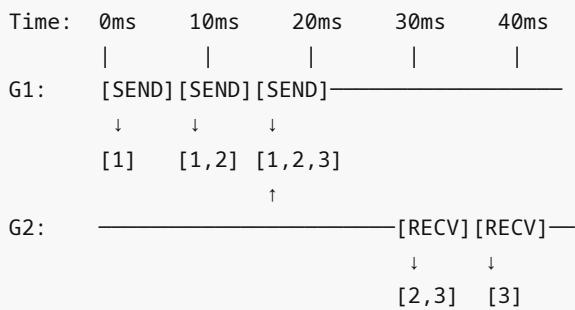
```

ch <- 1 // Returns immediately (buffer: [1])
ch <- 2 // Returns immediately (buffer: [1, 2])
ch <- 3 // BLOCKS (buffer full)

val := <-ch // Receives 1 (buffer: [2])
// Now ch <- 3 can proceed

```

## Execution Timeline



Goroutines **decouple**—sender doesn't wait for receiver (until buffer full).

## Happens-Before Guarantee

For buffered channels:

Send happens-before corresponding receive starts

**Note:** Different from unbuffered (send happens-before receive **completes**).

### Example:

```

var data string
done := make(chan bool, 1) // Buffered

go func() {
    data = "hello"      // (1) Write
    done <- true        // (2) Send (returns immediately)
}()

<-done           // (3) Receive
fmt.Println(data) // (4) Read

// Guaranteed order: (1) → (2) happens-before (3) → (4)
// Still safe to read `data`

```

## When to Use Buffered

### Use buffered when:

- Sender and receiver run at **different rates** (producer faster/slower than consumer)
- You want to **decouple goroutines** (avoid blocking sender on slow receiver)
- You're implementing a **work queue** with bounded size

- You want **burst handling** (accept spikes without blocking)
- You're using **buffered returns** (size 1) to prevent leaks

#### Example: Work Queue

```
func workQueue(capacity int) chan Job {
    jobs := make(chan Job, capacity) // Buffered

    // Workers can process at their own pace
    for i := 0; i < 10; i++ {
        go func() {
            for job := range jobs {
                process(job)
            }
        }()
    }

    return jobs
}
```

## Buffer Size: The Critical Choice

### Buffer Size = 0 (Unbuffered)

```
ch := make(chan int)
```

- **Tightest coupling:** Every send waits for receive
- **Strongest sync:** Explicit rendezvous
- **Slowest:** Maximum blocking
- **Simplest reasoning:** Happens-before is obvious

### Buffer Size = 1

```
ch := make(chan int, 1)
```

**Most common buffered size.** Why?

### Use case 1: Async result returns (leak prevention)

```
func compute() <-chan Result {
    ch := make(chan Result, 1) // Buffer size 1

    go func() {
        result := expensiveComputation()
        ch <- result // Never blocks (buffer has space)
    }()
}

return ch
```

```
// Even if caller never receives, goroutine completes
}
```

### Use case 2: Latest value pattern

```
type Monitor struct {
    updates chan State // Buffer size 1
}

func (m *Monitor) NotifyStateChange(state State) {
    select {
    case m.updates <- state: // Send if space
    default:                // Drop old value if full
        <-m.updates          // Remove old
        m.updates <- state   // Send new
    }
}
```

### Buffer Size = NumWorkers

```
workers := 10
jobs := make(chan Job, workers)
```

**Rationale:** Buffer matches concurrency level.

- Producer can submit up to `NumWorkers` jobs without blocking
- Workers never starve (jobs available)
- Bounded memory: Max buffered = worker count

#### Example:

```
func processParallel(items []Item) {
    workers := runtime.NumCPU()
    jobs := make(chan Item, workers)
    results := make(chan Result, workers)

    // Start workers
    for i := 0; i < workers; i++ {
        go func() {
            for item := range jobs {
                results <- process(item)
            }
        }()
    }

    // Send jobs
    go func() {
        for _, item := range items {
            jobs <- item
        }
    }
```

```

    close(jobs)
}()

// Collect results...
}

```

## Buffer Size = len(items)

```

items := getItems()
ch := make(chan Item, len(items))

```

### When to use:

- You know **exact number of items** upfront
- Want to send all **without blocking**
- Short-lived channel (not long-running queue)

### Example:

```

func scatterGather(items []Item) []Result {
    results := make(chan Result, len(items)) // Exact size

    for _, item := range items {
        go func(it Item) {
            results <- process(it) // Never blocks
        }(item)
    }

    // Collect exactly len(items) results
    collected := make([]Result, 0, len(items))
    for i := 0; i < len(items); i++ {
        collected = append(collected, <-results)
    }

    return collected
}

```

## Buffer Size = Large (100+)

```

ch := make(chan Message, 1000)

```

### When to use:

- High-throughput systems
- Smoothing out **bursty traffic**
- Decoupling fast producer from slow consumer

**Warning:** Large buffers hide problems.

- Delays detection of slow consumers
- Increases memory usage

- Can hide deadlocks (buffer never fills during testing)

**Rule of thumb:** If you need buffer > 100, reconsider your design. You might need:

- Dropping policy (discard old messages)
- Multiple consumers (scale out)
- Backpressure (slow down producer)

## Comparison Table

Aspect	Unbuffered	Buffered (Size N)
<b>Blocking</b>	Send blocks until recv	Send blocks when full
<b>Synchronization</b>	Explicit rendezvous	Async up to N items
<b>Latency</b>	Higher (blocking)	Lower (non-blocking until full)
<b>Throughput</b>	Lower (tight coupling)	Higher (decoupling)
<b>Memory</b>	None (no buffer)	$O(N)$ per channel
<b>Backpressure</b>	Immediate	Delayed (kicks in when full)
<b>Deadlock risk</b>	Higher (tight coupling)	Lower (loose coupling)
<b>Reasoning</b>	Simpler (explicit sync)	Harder (async effects)

## Common Misunderstandings

### Myth 1: "Buffered channels are faster"

**Truth:** They can improve **throughput** (more messages/sec) if sender/receiver speeds differ, but individual operations take the same time (~50-100ns).

**Example where buffered is faster:**

```
// Unbuffered: Send blocks until recv
func slowConsumer() {
    ch := make(chan int)

    go func() {
        for i := 0; i < 1000; i++ {
            ch <- i // Blocks on every send
        }
    }()

    for i := 0; i < 1000; i++ {
        <-ch
        time.Sleep(time.Millisecond) // Slow consumer
    }
}

// Total time: ~1000ms (send waits for slow recv)

// Buffered: Send doesn't block until full
```

```

func slowConsumerBuffered() {
    ch := make(chan int, 100)

    go func() {
        for i := 0; i < 1000; i++ {
            ch <- i // Only blocks when buffer full
        }
    }()

    for i := 0; i < 1000; i++ {
        <-ch
        time.Sleep(time.Millisecond) // Slow consumer
    }
}

// Total time: Still ~1000ms, but sender finishes early

```

## Myth 2: "Buffered channels prevent deadlocks"

**Truth:** They can **hide** deadlocks by delaying them.

```

// Unbuffered: Deadlocks immediately
func unbufferedDeadlock() {
    ch := make(chan int)
    ch <- 1 // DEADLOCK: no receiver
}

// Buffered: Deadlocks only when buffer full
func bufferedDeadlock() {
    ch := make(chan int, 5)
    for i := 0; i < 10; i++ {
        ch <- i // DEADLOCK at i=5: buffer full, no receiver
    }
}

```

Buffered channels **delay** the deadlock, making it harder to detect in small-scale tests.

## Myth 3: "Always use buffered size 1 to prevent leaks"

**Truth:** This works for single-result patterns but isn't a universal solution.

```

// Good: Single result
func compute() <-chan int {
    ch := make(chan int, 1)
    go func() {
        ch <- expensiveCalc() // Won't leak even if not received
    }()
    return ch
}

// BAD: Multiple results
func computeMany() <-chan int {

```

```

ch := make(chan int, 1) // Buffer size 1
go func() {
    for i := 0; i < 100; i++ {
        ch <- i // Will block at 2nd send if not received
    }
}()
return ch
}
// Still leaks if receiver stops early

```

**Better for multiple results:** Use context cancellation.

## Real-World Failure: Buffer Size Matters

**Company:** Social media platform (2020)

**What happened:**

During a traffic spike, HTTP handlers started timing out. CPU usage was normal. Memory usage exploded.

**Root cause:**

```

type Handler struct {
    notificationQueue chan Notification // Unbuffered
}

func (h *Handler) handleRequest(w http.ResponseWriter, r *http.Request) {
    // Process request
    notification := createNotification(r)

    h.notificationQueue <- notification // Blocks if consumer slow
    // Handler blocked → can't handle new requests

    w.WriteHeader(http.StatusOK)
}

```

**Problem:**

- Notification consumer was slow (DB writes)
- Handlers blocked sending notifications
- No handlers available for new requests
- HTTP timeout errors

**Fix:**

```

type Handler struct {
    notificationQueue chan Notification // Buffered (size 1000)
}

func (h *Handler) handleRequest(w http.ResponseWriter, r *http.Request) {
    notification := createNotification(r)

    select {

```

```

    case h.notificationQueue <- notification:
        // Sent successfully
    default:
        // Queue full, drop notification (log for monitoring)
        metrics.Inc("notifications_dropped")
    }

    w.WriteHeader(http.StatusOK)
}

```

#### Lessons:

1. Unbuffered channels couple goroutines tightly
2. Buffering decouples sender from slow receiver
3. Always have a drop policy when buffer fills
4. Monitor queue depth in production

## Decision Tree: Buffered or Unbuffered?

```

Do you need exact synchronization (rendezvous)?
├ YES → Unbuffered
|   └ Examples: event signaling, handshake, barrier
|
└ NO → Buffered
  |
  ├ Single result return?
  |   └ YES → Buffer size 1
  |
  ├ Known number of items?
  |   └ YES → Buffer size = len(items)
  |
  ├ Work queue with N workers?
  |   └ YES → Buffer size = N
  |
  ├ High throughput, bursty traffic?
  |   └ YES → Larger buffer (10-1000)
  |       └ + Implement drop policy
  |
  └ Unsure?
      └ Start with unbuffered (safer)
          Monitor in production
          Add buffering if needed

```

## Interview Traps

### Trap 1: "Buffered channels are always better"

**Wrong.** Buffering introduces complexity and can hide bugs.

### Correct answer:

"Unbuffered channels provide stronger synchronization guarantees and simpler reasoning. Buffered channels

decouple sender/receiver, improving throughput when they run at different speeds, but can hide deadlocks and coordination issues. Use unbuffered by default; add buffering when profiling shows it helps."

### Trap 2: "Buffer size should be large to avoid blocking"

**Wrong.** Large buffers consume memory and hide slow consumers.

**Correct answer:**

"Buffer size should match your concurrency model: size 1 for async returns, size N for N workers, or bounded based on memory constraints. Large buffers (100+) can hide slow consumer problems and delay backpressure. Better to fix the slow consumer or implement a drop policy."

### Trap 3: "Unbuffered channels are slower"

**Misleading.** Individual operations take the same time; throughput differs.

**Correct answer:**

"Both unbuffered and buffered channel operations take ~50-100ns. Unbuffered channels block sender until receiver is ready, which can reduce throughput if sender/receiver speeds differ. Buffered channels improve throughput by decoupling, but don't make individual operations faster."

### Trap 4: "This unbuffered deadlock would be fixed with a buffer"

```
ch := make(chan int)
ch <- 1 // Deadlock
```

**Lazy answer:** "Add buffer: `make(chan int, 1)`"

**Correct answer:**

"This is a design issue, not a buffering issue. The problem is sending without a receiver. Buffering would delay the deadlock, not fix it. Correct solutions: 1) Use a goroutine for the send, 2) Use a receiver on the other end, 3) Redesign to avoid this pattern."

## Key Takeaways

1. **Unbuffered = synchronous rendezvous** (tight coupling)
2. **Buffered = async up to capacity** (loose coupling)
3. **Buffer size 1** is most common for async returns
4. **Buffer size N** matches N workers pattern
5. **Large buffers hide problems** (slow consumers, deadlocks)
6. **Start unbuffered** (simpler), add buffering when needed
7. **Always implement drop policy** for bounded buffers
8. **Happens-before guarantees differ** (send-before-recv-completes vs send-before-recv-starts)

## What You Should Be Thinking Now

- "How do I wait on multiple channels?"
- "How do I implement timeouts?"
- "What's the select statement?"
- "How do I handle non-blocking channel operations?"

Next: [select.md](#) - Multiplexing channels with select.

## **Exercises (Do These Before Moving On)**

1. Write two versions of a program: one with unbuffered, one with buffered channels. Measure throughput difference.
2. Create a deadlock with an unbuffered channel. "Fix" it with a buffer. Explain why this is a band-aid, not a fix.
3. Implement a work queue with buffer size matching worker count.
4. Write code that uses buffer size 1 for an async result return. Verify it doesn't leak.
5. Implement a notification queue with a drop policy when buffer is full.

Don't continue until you can explain: "When is unbuffered better than buffered, despite potential blocking?"