

# Common Misconceptions About Concurrency and Memory

## The "Works On My Machine" Fallacy

**Misconception:** "My code runs fine in tests, so it's correct."

**Why it's wrong:**

1. **Timing-dependent:** Race conditions are non-deterministic
2. **CPU-dependent:** x86 has stronger memory ordering than ARM
3. **Load-dependent:** Low contention hides races
4. **Optimization-dependent:** Debug builds can hide races present in optimized builds

```
var ready bool
var data int

func writer() {
    data = 42
    ready = true
}

func reader() {
    for !ready {} // Spin
    print(data)
}

// Runs fine 99.99% of time
// But CAN print 0 (data race)
```

**In tests:** Usually prints 42 (lucky timing).

**In production:** Occasionally prints 0 (different CPU, load, optimization).

**Real example:** A trading system ran for 6 months without issues, then a single race condition caused \$2.3M loss when market volatility caused different timing.

**Truth:** Tests cannot prove absence of races. Must reason about happens-before.

---

## The "Atomic Types" Fallacy

**Misconception:** "Reads and writes to small types (bool, int) are atomic, so no synchronization needed."

**Why it's wrong:**

Even if individual reads/writes are atomic (no tearing), **without synchronization there's no visibility guarantee.**

```
var flag bool
```

```
// Goroutine 1
flag = true

// Goroutine 2
for !flag {} // Can spin forever!
print("done")
```

### What happens:

1. G1 writes `flag=true` to its CPU cache
2. G2 reads `flag` from its own CPU cache (cached `false`)
3. **G2's cache never invalidated** → spins forever

**This is NOT a bug in Go.** This is how modern CPUs work (cache coherency is not instantaneous).

**Fix:** Use `atomic.LoadBool` / `atomic.StoreBool`.

```
var flag atomic.Bool

// Goroutine 1
flag.Store(true) // Inserts memory barrier

// Goroutine 2
for !flag.Load() {} // Inserts memory barrier
print("done")
```

**Truth: Small types != atomic operations.** Atomicity ≠ visibility. Use sync/atomic.

**Real example:** Video streaming service had goroutines spin indefinitely waiting for boolean flag, causing memory leak. Worked in dev (fewer CPUs), failed in production.

## The "Single Writer" Fallacy

**Misconception: "If only one goroutine writes, I don't need synchronization."**

**Why it's wrong:**

Even with single writer, **readers need synchronization to see writes.**

```
var config *Config

// Only this goroutine writes (single writer)
func reloadConfig() {
    newConfig := loadFromFile() // (1)
    config = newConfig          // (2) Write
}

// Many goroutines read
func handleRequest() {
    cfg := config // (3) Read
    if cfg != nil {
        use(cfg) // Might see partially constructed Config
```

```
    }
}
```

**Problem:** No happens-before between (2) and (3).

- Reader can see pointer but not fields (partial construction)
- Reader might see old pointer indefinitely (caching)

**Fix:** Use atomic.Pointer or mutex.

```
var config atomic.Pointer[Config]

func reloadConfig() {
    newConfig := loadFromFile()
    config.Store(newConfig) // Atomic store
}

func handleRequest() {
    cfg := config.Load() // Atomic load
    if cfg != nil {
        use(cfg)
    }
}
```

**Truth:** Multiple readers need synchronization with writer even if writer is single.

## The "Immutable Data" Fallacy

**Misconception:** "Immutable data doesn't need synchronization."

**Why it's wrong:**

Even if data is immutable after construction, **publication of the data needs synchronization**.

```
type Config struct {
    Host string
    Port int
}

var config *Config

func init() {
    config = &Config{ // (1) Construct
        Host: "localhost",
        Port: 8080,
    }
}

func main() {
    go func() {
        if config != nil {
            connect(config.Host, config.Port) // (2) Use
        }
    }()
}
```

```
    }
    }()
}
```

**Problem:** No happens-before between init and goroutine.

**Wait, doesn't init happen-before main?**

Yes: init →hb main starts.

**But:** main starts →hb go func()... does NOT mean init →hb inside func().

**Correct reasoning:**

- init →hb main starts (guaranteed)
- main starts →hb "go func()" statement (guaranteed)
- "go func()" statement →hb func executes (guaranteed)
- **Chain: init →hb (2) ✓ Actually safe!**

**But this is NOT safe:**

```
var config *Config

func Load() {
    config = &Config{...} // (1)
}

func Use() {
    if config != nil {
        use(config) // (2)
    }
}

// In main:
go Load()
go Use()

// NO happens-before between (1) and (2) → race
```

**Truth: Immutable data still requires synchronized publication.**

## The "Happens-Before Means Time" Fallacy

**Misconception: "Happens-before means happens earlier in wall-clock time."**

**Why it's wrong:**

Happens-before is about **causality and visibility**, not time.

```
var a, b int
ch := make(chan int, 1)

// Goroutine 1
```

```

ch <- 0      // (1) Happens at time T
a = 1        // (2) Happens at time T+10

// Goroutine 2
<-ch         // (3) Happens at time T+1
print(a)     // (4) Happens at time T+2

```

**In wall-clock time:** (3) and (4) happen BEFORE (2).

**But happens-before:** (1) →hb (3), NOT (2) →hb (4).

**Result:** (4) can print 0 even though it executed after (2) in real time.

**Truth: Happens-before is about synchronization, not time.** Operations can happen later in time but not be visible.

---

## The "Volatile Is Enough" Fallacy (From Other Languages)

**Misconception: "Go doesn't have `volatile`, but I can simulate it."**

**Why it's wrong:**

**In C/C++:**

```
volatile int flag = 0;
```

- Prevents compiler from optimizing away reads/writes
- Does NOT insert memory barriers (usually)
- Does NOT guarantee visibility across CPUs

**In Java:**

```
volatile int flag = 0;
```

- Inserts memory barriers (acquire/release semantics)
- Guarantees visibility

**In Go:**

- No `volatile` keyword
- Must use `sync/atomic` for visibility

**Attempting to simulate `volatile`:**

```

//go:noinline
func readFlag() bool { return flag }

// WRONG: This prevents inlining, not caching

```

**Truth: Go requires `sync/atomic` for visibility.** No way to simulate `volatile`.

---

## The "Mutexes Protect Data" Fallacy

**Misconception:** "Mutexes protect data structures."

**Why it's misleading:**

Mutexes don't know about data. **You protect data by establishing conventions.**

```
var mu sync.Mutex
var count int

func increment() {
    mu.Lock()
    count++
    mu.Unlock()
}

func read() int {
    // WRONG: Forgot to lock
    return count // Data race!
}
```

**Mutex didn't protect count.** You must lock in ALL accesses.

**Correct version:**

```
func read() int {
    mu.Lock()
    defer mu.Unlock()
    return count // Safe
}
```

**Truth: Mutexes protect critical sections, not data.** You must ensure all accesses are inside critical sections.

**Real example:** Payment processor had mutex protecting balance updates but forgot to protect balance reads in logging. Rare race caused double-spend when reader saw partial update.

---

## The "Race-Free Equals Correct" Fallacy

**Misconception:** "If my code has no data races, it's correct."

**Why it's wrong:**

**Race-free** doesn't mean **correct logic**.

```
var (
    balance int
    mu      sync.Mutex
)

func withdraw(amount int) bool {
```

```

mu.Lock()
defer mu.Unlock()

if balance >= amount {
    // Simulated processing time
    time.Sleep(time.Millisecond)
    balance -= amount
    return true
}
return false
}

// Goroutine 1: withdraw(100)
// Goroutine 2: withdraw(100)
// Balance: 150

// Both withdrawals succeed! Balance = -50

```

**No data race** (all accesses protected by mutex).

**But incorrect** (allows overdraft).

**Fix:** Check-then-act must be atomic at business-logic level, not just memory level.

**Truth:** Race-freedom is necessary but not sufficient for correctness.

---

## The "Buffered Channels Don't Block" Fallacy

**Misconception:** "Buffered channels never block, unbuffered channels always block."

**Why it's wrong:**

**Buffered channels block when full:**

```

ch := make(chan int, 2) // Buffer size 2

ch <- 1 // Doesn't block (buffer: 1/2)
ch <- 2 // Doesn't block (buffer: 2/2)
ch <- 3 // BLOCKS forever (buffer full, no receiver)

```

**Unbuffered channels don't always block:**

```

ch := make(chan int)

// Goroutine 1
ch <- 0 // Blocks until receiver ready

// Goroutine 2
<-ch // Receiver ready, sender unblocks immediately

```

**Truth:** All channel sends can block if no receiver or buffer full. All channel receives block if no sender or buffer empty.

---

## The "Goroutines Are Free" Fallacy

**Misconception:** "Goroutines are so lightweight, I can create millions."

**Why it's wrong:**

Goroutines are cheap, not free.

**Cost per goroutine:**

- 2 KB initial stack (can grow)
- Goroutine struct (~few hundred bytes)
- Scheduler overhead

**1 million goroutines:**

- Stack: 2 GB minimum
- Runtime structures: ~500 MB
- Scheduler contention

**Real example:** Web scraper created 1 million goroutines, consumed 3 GB, then crashed with OOM when stacks grew.

**Fix:** Use worker pool pattern (bounded concurrency).

```
const maxWorkers = 100

sem := make(chan struct{}, maxWorkers)

for url := range urls {
    sem <- struct{}{} // Acquire
    go func(u string) {
        defer func() { <-sem }()
        fetch(u)
    }(url)
}
```

**Truth:** Goroutines are cheap, but not free. Bound concurrency in production.

---

## The "nil Channel Is Broken" Fallacy

**Misconception:** "Sending/receiving on nil channel is a bug."

**Why it's misleading:**

Nil channel behavior is **intentional and useful**.

```
var ch chan int // nil

ch <- 0 // Blocks forever
<-ch    // Blocks forever

// In select:
```

```

select {
case ch <- 0: // Never selected (blocks forever)
case <-ch:    // Never selected
case <-other: // This can be selected
}

```

### Useful pattern: Disabling select cases

```

func merge(a, b <-chan int) <-chan int {
out := make(chan int)
go func() {
    defer close(out)
    for a != nil || b != nil {
        select {
        case v, ok := <-a:
            if !ok {
                a = nil // Disable this case
                continue
            }
            out <- v
        case v, ok := <-b:
            if !ok {
                b = nil // Disable this case
                continue
            }
            out <- v
        }
    }
}()
return out
}

```

**Truth: Nil channels are a feature, not a bug.** Use them to disable select cases.

## The "Close Detects All Goroutines Done" Fallacy

**Misconception: "close(ch) means all senders are done."**

**Why it's wrong:**

You can close a channel while senders are still active → **panic**.

```

ch := make(chan int)

// Goroutine 1
go func() {
    time.Sleep(time.Millisecond)
    ch <- 1 // Might panic if ch closed
}()

```

```
// Goroutine 2
close(ch) // Close immediately
```

**Panic:**

```
panic: send on closed channel
```

**Correct pattern:**

1. **Only sender closes** (or use sync.WaitGroup to coordinate)
2. Never close from receiver

```
ch := make(chan int)
var wg sync.WaitGroup

wg.Add(1)
go func() {
    defer wg.Done()
    ch <- 1
}()

go func() {
    wg.Wait()      // Wait for all senders
    close(ch)     // Safe to close
}()
```

**Truth: close() doesn't detect completion; you must ensure completion first.**

## The "Context Cancel Stops Goroutines" Fallacy

**Misconception: "ctx.Cancel() stops goroutines."**

**Why it's wrong:**

Context cancellation is **cooperative**. Goroutines must check context.

```
ctx, cancel := context.WithCancel(context.Background())

go func() {
    for {
        work() // Doesn't check context → never stops
    }
}()

cancel() // Goroutine keeps running!
```

**Fix:** Check context in loop.

```
go func() {
    for {
```

```

    select {
        case <-ctx.Done():
            return // Goroutine exits
        default:
            work()
    }
}

```

**Truth:** Context cancellation is a signal, not a kill command. Goroutines must cooperate.

---

## The "Race Detector Finds All Races" Fallacy

**Misconception:** "If race detector doesn't report races, my code is race-free."

**Why it's wrong:**

Race detector is **dynamic analysis** → only finds races in executed code paths.

```

func process(condition bool) {
    if condition {
        // Path A (uses shared variable)
    } else {
        // Path B (uses shared variable)
    }
}

// Test only exercises Path A
// Race in Path B not detected

```

**Coverage-dependent:** Need to execute all code paths under all interleavings (impossible).

**Truth:** Race detector is necessary but not sufficient. Must also reason about happens-before.

---

## The "Defer Runs Before Return" Fallacy

**Misconception:** "defer runs before function returns."

**Truth:** Defer runs before function returns **but after return value is evaluated**.

```

func mystery() (result int) {
    defer func() {
        result++ // Modifies return value
    }()
    return 0 // Returns 0, then defer runs, result becomes 1
}

fmt.Println(mystery()) // Prints 1

```

**But:**

```

func mystery2() int {
    result := 0
    defer func() {
        result++ // Modifies local variable, not return value
    }()
    return result // Returns 0 (value copied before defer)
}

fmt.Println(mystery2()) // Prints 0

```

**In concurrency:**

```

func send(ch chan int, v int) {
    defer close(ch) // Wrong order!
    ch <- v
}

// Receiver might miss value if close() happens first

```

**Truth: defer executes in LIFO order after return value evaluation.** Be careful with ordering.

---

## Interview Traps

### Trap 1: "Code review looks fine, no obvious races"

**Wrong assumption:** Visual inspection finds races.

**Correct answer:** "Visual inspection can spot obvious races, but subtle races require reasoning about happens-before relationships. All shared memory accesses need synchronization or happens-before ordering. I would also run race detector and write concurrency tests."

### Trap 2: "This worked in 10,000 test runs"

**Wrong assumption:** Passing tests prove correctness.

**Correct answer:** "Passing tests increase confidence but don't prove absence of races. Race conditions are non-deterministic and depend on timing, CPU architecture, and load. Must verify happens-before relationships and use race detector."

### Trap 3: "I'll just add time.Sleep to fix the timing issue"

**Wrong:** time.Sleep is not synchronization.

**Correct answer:** "time.Sleep doesn't provide happens-before guarantees and just masks races. The correct fix is to use proper synchronization: channels for communication, mutexes for shared state, or atomic operations for flags."

### Trap 4: "Reads don't need protection, only writes"

**Wrong:** Readers need synchronization to see writes.

**Correct answer:** "Both reads and writes need synchronization to establish happens-before. Without it, reads might see stale values indefinitely due to CPU caching, even if writes are completed."

---

## Key Takeaways

1. **Tests cannot prove absence of races** (non-deterministic)
  2. **"Atomic" types ≠ sync/atomic operations** (no visibility without barriers)
  3. **Single writer still needs reader synchronization** (visibility is bidirectional)
  4. **Publication of immutable data needs synchronization** (construction vs. publication)
  5. **Happens-before ≠ time ordering** (causality, not clocks)
  6. **Mutexes protect critical sections, not data** (you must enforce convention)
  7. **Race-free ≠ correct** (no races doesn't mean correct logic)
  8. **All channels can block** (buffered when full, unbuffered when no receiver)
  9. **Goroutines cost memory** (stack + runtime structures)
  10. **Race detector is necessary, not sufficient** (dynamic analysis, not proof)
- 

## What You Should Be Thinking Now

- "What mistakes am I making in my current code?"
- "How do I verify my happens-before reasoning?"
- "What patterns should I always use?"

Next: [Section 03: Classic Problems](#) - Learn how race conditions manifest in real code.

---

## Exercises

1. Find 5 misconceptions you held before reading this. Write down what you learned.
2. Review your recent code: Find 3 places where you made assumptions about:
  - Visibility without synchronization
  - Mutexes protecting data
  - Tests proving correctness
3. Write a program that "works" in tests but has a race. Run with `-race` to confirm.
4. Explain to a colleague why "it works on my machine" is not proof of correctness.
5. Create a checklist for code review: What questions should you ask about every shared variable?

Don't continue until you can: "Identify and debunk dangerous misconceptions about concurrent code."