

# Memory Visibility vs. Execution Ordering

## The Core Confusion

Most engineers conflate two distinct concepts:

1. **Execution ordering:** The sequence in which operations complete
2. **Memory visibility:** When one CPU sees effects of another CPU's operations

**Critical insight:** Operations can execute in order but still have visibility problems.

## What is Memory Visibility?

**Memory visibility** is about **when writes by one goroutine become observable by reads in another goroutine.**

```
var a int

// CPU 1 (Goroutine 1)
a = 42 // Write to L1 cache

// CPU 2 (Goroutine 2)
print(a) // Read from L1 cache

// Question: Does CPU 2 see a=42?
// Answer: Maybe. Depends on cache coherency protocol and memory barriers.
```

**Without synchronization, CPU 2 might read stale value from its cache indefinitely.**

## What is Execution Ordering?

**Execution ordering** is about the sequence in which instructions execute within a goroutine.

```
// Source code order
a := 1 // (1)
b := 2 // (2)
c := a + b // (3)

// Compiler can reorder to:
b := 2 // (2)
a := 1 // (1)
c := a + b // (3)

// Why? (1) and (2) are independent.
// Observable behavior within goroutine unchanged.
```

**Important:** Reordering is legal if single-goroutine semantics are preserved.

## Why CPUs and Compilers Reorder

## Compiler Reordering

Compilers optimize code by reordering instructions.

### Example: Register allocation

```
// Source
a = 1
b = 2
c = a

// Compiler sees: c depends on a, not b.
// Optimized:
a = 1
c = a // Use a immediately (a might be in register)
b = 2 // Delay b assignment
```

**Rule:** Compiler preserves **as-if-serial semantics** (behavior as observed within goroutine).

**Problem:** Other goroutines can see reordered effects.

## CPU Reordering (Store Buffer)

Modern CPUs use **store buffers** to improve write performance.

### What happens on write:

1. CPU writes to store buffer (fast, ~1 cycle)
2. Value eventually written to L1 cache (~10 cycles)
3. Cache line eventually synced to other CPUs (~100+ cycles)

**Implication:** Writes are not immediately visible to other CPUs.

```
// CPU 1
WRITE a = 1 // Enter store buffer
WRITE b = 1 // Enter store buffer

// CPU 2 (simultaneously)
READ b        // Might see b=1 (if flushed from store buffer)
READ a        // Might see a=0 (if not yet flushed)

// Result: See b=1 but a=0 (reordered from perspective of CPU 2)
```

## CPU Reordering (Out-of-Order Execution)

CPUs execute instructions out-of-order for pipeline efficiency.

```
// Instructions
LOAD R1, [x] // (1) Memory load (slow, ~100 cycles)
LOAD R2, [y] // (2) Memory load (slow)
ADD R3, R1, R2 // (3) Depends on (1) and (2)
```

```
// CPU might execute:  
(2) before (1) to parallelize memory loads
```

If another CPU reads intermediate states, sees reordered effects.

## The Four Types of Memory Reordering

CPUs can reorder memory operations in four ways:

Reordering Type	Description	Example
<b>Store-Store</b>	Write B before Write A	a=1; b=2 → b=2; a=1
<b>Load-Load</b>	Read B before Read A	x=a; y=b → y=b; x=a
<b>Load-Store</b>	Write B before Read A	x=a; b=2 → b=2; x=a
<b>Store-Load</b>	Read B before Write A	a=1; x=b → x=b; a=1

**x86/x64 guarantees:**

- ✓ Store-Store: No reordering
- ✓ Load-Load: No reordering
- ✓ Load-Store: No reordering
- ✗ Store-Load: **CAN reorder** (most expensive to prevent)

**ARM/ARM64:**

- ✗ All four can reorder (weakly ordered)

**Go memory model abstracts over architectures:** Assumes weak ordering (like ARM). Code correct on ARM works on x86.

## Example: Dekker's Algorithm (Classic Failure)

Dekker's algorithm is a mutual exclusion algorithm without locks.

```
var flag1, flag2 int

// Goroutine 1
flag1 = 1           // (1) "I want to enter"
if flag2 == 0 {    // (2) "Is other goroutine out?"
    // Critical section
}

// Goroutine 2
flag2 = 1           // (3) "I want to enter"
if flag1 == 0 {    // (4) "Is other goroutine out?"
    // Critical section
}
```

**Intended behavior:** At least one sees the other's flag.

**With reordering:**

Time	G1	G2
1	flag2 == 0 (read)	flag1 == 0 (read)
2	flag1 = 1 (write)	flag2 = 1 (write)

**Both reads before both writes (load-load and store-store reordering).**

**Result:** Both goroutines enter critical section → race!

**Why it fails:** No happens-before. Compiler/CPU can reorder independent loads and stores.

## Example: Write Buffering Bug

```
var a, b int

// Goroutine 1
a = 1 // (1)
b = 1 // (2)

// Goroutine 2
for b == 0 {} // (3) Spin until b=1
print(a) // (4)

// Expected: G2 sees b=1, then a must be 1
// Reality: Can print a=0
```

**Execution timeline:**

1. CPU 1 executes (1): a=1 enters store buffer
2. CPU 1 executes (2): b=1 enters store buffer
3. **b=1 flushes to memory first** (arbitrary order)
4. CPU 2 sees b=1, exits loop
5. CPU 2 reads a (still 0 in CPU 2's cache, a=1 not yet visible)

**Result:** G2 sees b=1 but a=0 (reordered visibility).

## Memory Barriers (How Synchronization Works)

**Memory barrier:** CPU instruction that enforces ordering.

### Types of Barriers

1. **Store barrier (SFENCE on x86):** All stores before barrier complete before stores after barrier
2. **Load barrier (LFENCE on x86):** All loads before barrier complete before loads after barrier
3. **Full barrier (MFENCE on x86):** All memory ops before barrier complete before ops after barrier

**Go synchronization primitives insert barriers automatically.**

### How Mutex Inserts Barriers

```

mu.Lock()      // Inserts ACQUIRE barrier
a = 1          // Protected write
mu.Unlock()    // Inserts RELEASE barrier

```

#### **ACQUIRE barrier (Lock):**

- Prevents operations AFTER Lock from moving BEFORE Lock
- Ensures you see all writes from previous Unlock

#### **RELEASE barrier (Unlock):**

- Prevents operations BEFORE Unlock from moving AFTER Unlock
- Ensures your writes visible to next Lock

**Result:** Lock/Unlock create happens-before.

### **How Channel Send Inserts Barriers**

```

a = 1      // (1)
ch <- 0    // (2) RELEASE barrier

<-ch      // (3) ACQUIRE barrier
print(a)   // (4)

```

**RELEASE (send):** Flushes store buffer, ensures (1) visible before (2) completes.

**ACQUIRE (receive):** Ensures all writes before send are visible after receive.

**Result:** (1) →hb (4)

### **Example: Double-Checked Locking (Why It Fails)**

```

type Singleton struct {
    instance *Instance
    mu       sync.Mutex
}

func (s *Singleton) Instance() *Instance {
    if s.instance != nil { // (1) Read without lock
        return s.instance // (2) Return
    }

    s.mu.Lock()
    if s.instance == nil {
        s.instance = &Instance{ // (3) Write
            Field1: value1,    // (4)
            Field2: value2,    // (5)
        }
    }
    s.mu.Unlock() // (6) RELEASE barrier
}

```

```
    return s.instance  
}
```

### The bug:

**Goroutine A:** (3) writes pointer, (4)(5) initialize fields, (6) unlocks.

**Goroutine B:** (1) reads pointer (not protected by lock).

**Problem:** Even if B sees pointer non-nil, no happens-before between (4)(5) and (1).

### Reordering scenario:

1. A allocates memory for Instance
2. A writes pointer to s.instance (pointer now non-nil) — (3)
3. **B reads pointer at (1), sees non-nil**
4. **B returns partially constructed instance** — (2)
5. A writes Field1 — (4)
6. A writes Field2 — (5)

**Result:** B uses instance with zero-value fields → crash or incorrect behavior.

**Why?** No ACQUIRE barrier at (1). Compiler/CPU can reorder (3) to be visible before (4)(5).

## Visibility vs. Ordering in Practice

### Case 1: Sequential Consistency (Strong Guarantee)

**Sequential consistency:** All operations appear in some global order consistent with each goroutine's program order.

**Go does NOT guarantee sequential consistency** (too expensive).

### Example that violates sequential consistency:

```
var a, b, x, y int  
  
// Goroutine 1  
a = 1  
x = b  
  
// Goroutine 2  
b = 1  
y = a  
  
// Possible: x=0 and y=0  
// Violates sequential consistency
```

**If sequentially consistent, one of these must be true:**

- G1's writes before G2's: x=0, y=1
- G2's writes before G1's: x=1, y=0
- Interleaved: x=1, y=1 or x=0, y=1 or x=1, y=0

**But x=0, y=0 is impossible with sequential consistency.**

### With reordering:

- G1 reads b before writing a (load-store reorder)
- G2 reads a before writing b (load-store reorder)
- Result: Both reads see 0

This is legal in Go's memory model (no synchronization).

### Case 2: Release-Acquire Ordering (What Go Provides)

Go's synchronization primitives guarantee release-acquire ordering:

- **Release:** Stores before Unlock/Send/Done visible to acquirer
- **Acquire:** Acquirer sees all stores from releaser

This is weaker than sequential consistency but sufficient for correctness.

```
var a, b int
var mu sync.Mutex

// Goroutine 1
mu.Lock()
a = 1
b = 1
mu.Unlock() // RELEASE

// Goroutine 2
mu.Lock() // ACQUIRE
print(a, b) // Sees a=1, b=1
mu.Unlock()
```

## Real-World Failure: Linux Kernel RCU Bug (2006)

**RCU (Read-Copy-Update):** Lock-free data structure in Linux kernel.

**Bug:** Missing memory barrier in RCU update path.

```
// Simplified
void publish(node *new_node) {
    new_node->data = value; // (1)
    list->head = new_node; // (2) Publish
}

node *read() {
    node *n = list->head; // (3) Read pointer
    return n->data; // (4) Read data
}
```

**Without barrier between (1) and (2):** CPU can make (2) visible before (1).

**Result:** Reader sees pointer to uninitialized node → crash.

**Fix:** Insert `smp_wmb()` (write memory barrier) between (1) and (2).

```

new_node->data = value; // (1)
smp_wmb();           // Write barrier
list->head = new_node; // (2)

```

**In Go:** This would be prevented by using channels or mutexes (automatic barriers).

## Real-World Failure: Java HashMap Race (Common)

**Java's HashMap is not thread-safe.** Concurrent access without synchronization causes visibility problems.

```

// Thread 1
map.put("key", "value"); // Writes to internal array

// Thread 2
String v = map.get("key"); // Reads internal array
// Can return null even after put() completed

```

**Why?** No happens-before. Thread 2's CPU cache might not see Thread 1's writes.

**Same in Go with non-concurrent map:**

```

m := make(map[string]int)

go func() {
    m["key"] = 42 // Write
}()

go func() {
    v := m["key"] // Read - can see 0 or panic
}()

```

**Go detects this with race detector:**

```

WARNING: DATA RACE
Write by goroutine 1
Read by goroutine 2

```

## How to Reason About Visibility

### Rule 1: Synchronization Provides Visibility

If operations are synchronized (happens-before), writes are visible.

```

var a int
var mu sync.Mutex

// Writer
mu.Lock()
a = 1
mu.Unlock() // RELEASE: a written to memory

```

```
// Reader
mu.Lock()    // ACQUIRE: a read from memory
v := a        // Sees a=1
mu.Unlock()
```

## Rule 2: Without Synchronization, No Visibility Guarantee

```
var a int

// Writer
a = 1 // Might stay in CPU cache

// Reader
v := a // Might read from different CPU's cache (stale value)
```

## Rule 3: Atomics Provide Visibility

```
var a int
var flag int32

// Writer
a = 1
atomic.StoreInt32(&flag, 1) // RELEASE: a visible

// Reader
for atomic.LoadInt32(&flag) == 0 {} // ACQUIRE
v := a // Sees a=1
```

## Performance Implications

**Memory barriers are expensive:**

Operation	x86 Latency	ARM Latency
Register read	~1 cycle	~1 cycle
L1 cache read	~4 cycles	~4 cycles
Store buffer flush	~10 cycles	~10 cycles
<b>Memory barrier</b>	~20-100 cycles	~50-200 cycles
Cache line sync	~100-300 cycles	~100-500 cycles

**Why slow?** Barrier flushes pipelines, waits for pending memory ops, ensures cache coherency.

**Implication:** Minimize synchronization in hot paths.

## Interview Traps

## **Trap 1: "Reads are free, only writes need synchronization"**

**Wrong.** Both reads and writes need synchronization to ensure visibility.

**Correct:** "Without synchronization, reads might see stale values indefinitely due to CPU caching. Both directions (write→read and read→write) need happens-before."

## **Trap 2: "volatile keyword in other languages is like atomic"**

**Wrong.** Volatile in C/C++/Java only prevents compiler reordering, not CPU reordering (mostly).

**Go has no volatile.** Use sync/atomic for visibility guarantees.

**Correct:** "In Go, use sync/atomic for atomic operations with memory visibility guarantees. sync/atomic operations include necessary memory barriers to ensure visibility across goroutines."

## **Trap 3: "If execution order is A then B, reads see A before B"**

**Wrong.** Execution order within a goroutine doesn't guarantee visibility order in another goroutine.

**Correct:** "Execution order (program order) is only guaranteed within a single goroutine. For visibility across goroutines, must establish happens-before through synchronization. CPUs can have different views of memory due to caching."

## **Trap 4: "x86 has strong memory model, so my code is safe"**

**Wrong.** Code might work on x86 but fail on ARM.

**Correct:** "Go's memory model assumes weak ordering (like ARM). Code must be correct under weak ordering to be portable. Never rely on architecture-specific behavior. Use race detector to find such bugs."

## **Key Takeaways**

1. **Visibility ≠ Ordering:** Operations can execute in order but have visibility problems
2. **CPUs cache memory:** Without barriers, writes stay in cache
3. **CPUs reorder operations:** For performance, within as-if-serial constraint
4. **Compilers reorder operations:** Within single-goroutine semantics
5. **Synchronization inserts barriers:** Lock/Channel/Atomic enforce ordering and visibility
6. **Go assumes weak memory model:** Correct code works on x86 and ARM
7. **Always use race detector:** Catches visibility bugs

## **What You Should Be Thinking Now**

- "How do I know if I need synchronization?"
- "What are the common misconceptions about memory?"
- "Why do tests pass but production fails?"

**Next:** [common-misconceptions.md](#) - Debunking dangerous myths about memory and concurrency.

---

## **Exercises**

1. Explain why this can print "y=1, x=0":

```
var x, y int
go func() { x=1; y=1 }()
go func() { if y==1 { print(x) } }()

```

2. Draw the store buffer timeline for:

```
a=1; b=1
// Other goroutine reads b then a
```

3. Explain why double-checked locking fails using visibility arguments (not ordering).

4. What memory barriers does this insert?

```
mu.Lock()
data = value
mu.Unlock()
```

5. Why does x86 code sometimes work without barriers but fail on ARM?

Don't continue until you can: "Distinguish between execution ordering and memory visibility in concurrent code."