

False Sharing

What is False Sharing?

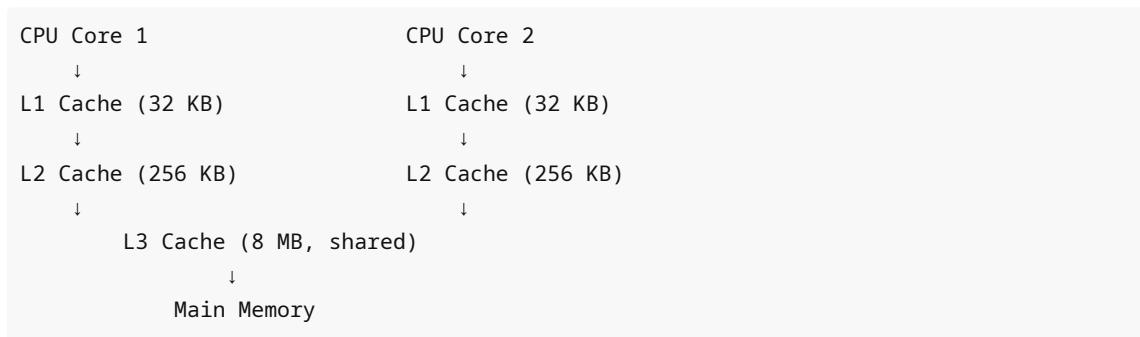
False sharing: Performance degradation when multiple CPUs access different variables that reside on the same cache line, causing unnecessary cache coherency traffic.

Critical insight: Even though goroutines access DIFFERENT variables (no logical sharing), they share the SAME cache line (hardware sharing).

Cache line: Smallest unit of memory transferred between CPU cache and main memory (typically 64 bytes on x86/ARM).

How CPU Caches Work

Cache Hierarchy



Latencies (approximate):

- L1: ~~4 cycles~~ (1ns)
- L2: ~~12 cycles~~ (3ns)
- L3: ~~40 cycles~~ (10ns)
- Main memory: ~~200 cycles~~ (50ns)

Cache Coherency (MESI Protocol)

When one CPU writes to a cache line, protocol ensures other CPUs' caches are invalidated.

States:

- Modified: This CPU has exclusive, modified copy
- Exclusive: This CPU has exclusive, clean copy
- Shared: Multiple CPUs have clean copies
- Invalid: Cache line is stale

Example:

Initial: Variable x in Shared state in CPU1 and CPU2 caches

CPU1 writes x:
1. CPU1 cache line → Modified
2. CPU2 cache line → Invalid (invalidation message sent)
3. CPU2 must reload from memory on next access

False sharing: CPUs invalidate each other's caches even when accessing different variables.

Example 1: Counter Array False Sharing

```
type Counter struct {
    a int64 // Offset 0-7
    b int64 // Offset 8-15
    // Both in same 64-byte cache line
}

var counter Counter

// Goroutine 1 on CPU 1
func incrementA() {
    for i := 0; i < 1000000; i++ {
        atomic.AddInt64(&counter.a, 1)
    }
}

// Goroutine 2 on CPU 2
func incrementB() {
    for i := 0; i < 1000000; i++ {
        atomic.AddInt64(&counter.b, 1)
    }
}

// False sharing:
// CPU1 writes counter.a → invalidates CPU2's cache line
// CPU2 writes counter.b → invalidates CPU1's cache line
// Constant cache line bouncing → slow
```

Benchmark:

```
func BenchmarkFalseSharing(b *testing.B) {
    var counter Counter
    var wg sync.WaitGroup

    wg.Add(2)
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&counter.a, 1)
        }
    }()
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&counter.b, 1)
        }
    }()
}
```

```

        wg.Wait()
}

// Result: ~50ns per operation (cache line bouncing)

```

Fix: Padding

```

type Counter struct {
    a int64
    _ [56]byte // Padding: 64 - 8 = 56 bytes
    b int64
    // a and b now in different cache lines
}

// Result: ~5ns per operation (10x faster!)

```

Example 2: Slice Element False Sharing

```

type Item struct {
    value int64
}

var items [8]Item

// 8 goroutines, each updates its own item
for i := 0; i < 8; i++ {
    go func(idx int) {
        for j := 0; j < 1000000; j++ {
            atomic.AddInt64(&items[idx].value, 1)
        }
    }(i)
}

// False sharing:
// Each Item is 8 bytes
// 8 Items = 64 bytes = 1 cache line
// All goroutines contend for same cache line

```

Benchmark:

```

func BenchmarkSliceFalseSharing(b *testing.B) {
    var items [8]Item
    var wg sync.WaitGroup

    for i := 0; i < 8; i++ {
        wg.Add(1)
        go func(idx int) {
            defer wg.Done()
            for j := 0; j < b.N; j++ {

```

```

        atomic.AddInt64(&items[idx].value, 1)
    }
} (i)
}
wg.Wait()
}

// Result: ~200ns per operation (severe contention)

```

Fix: Padding

```

type Item struct {
    value int64
    _     [56]byte // Padding
}

// Result: ~10ns per operation (20x faster!)

```

Example 3: Per-Goroutine State False Sharing

```

type Worker struct {
    id    int
    count int64
}

var workers [4]Worker

// Each worker increments its own count
for i := 0; i < 4; i++ {
    go func(w *Worker) {
        for j := 0; j < 1000000; j++ {
            atomic.AddInt64(&w.count, 1)
        }
    }(&workers[i])
}

// False sharing: All Worker structs in same cache lines

```

Fix: Use sync.Pool or separate slices

```

// Option 1: Padding
type Worker struct {
    id    int
    count int64
    _     [48]byte // 64 - 16 = 48
}

// Option 2: Separate arrays
var workerIDs [4]int

```

```
var workerCounts [4]int64
var _ [8]int64 // Padding between arrays
```

Cache Line Size in Go

```
// Get cache line size
import "unsafe"

const CacheLinePad = 64 // Typical cache line size

// Or use runtime
import "runtime"

func getCacheLineSize() int {
    // Not directly exposed, assume 64 bytes
    return 64
}
```

Real-World Failure: High-Frequency Trading (2019)

Company: Financial trading firm

Issue: Latency spikes in order execution

Simplified scenario:

```
type OrderBook struct {
    bidCount int64 // Offset 0
    askCount int64 // Offset 8
    // Both in same cache line
}

var book OrderBook

// Thread 1: Update bids
func updateBids() {
    for {
        atomic.AddInt64(&book.bidCount, 1)
        // Process bid
    }
}

// Thread 2: Update asks
func updateAsks() {
    for {
        atomic.AddInt64(&book.askCount, 1)
        // Process ask
    }
}
```

```
// False sharing: Cache line bouncing between threads  
// Latency: 100ns → 2000ns (20x slower)
```

Impact: Missed trading opportunities (microseconds matter).

Fix: Padding

```
type OrderBook struct {  
    bidCount int64  
    _ [56]byte  
    askCount int64  
    _ [56]byte  
}  
  
// Latency: 100ns (consistent)
```

Detecting False Sharing

1. Performance Profiling

```
go test -bench=. -cpuprofile=cpu.out  
go tool pprof -http=:8080 cpu.out
```

Look for:

- High contention on atomic operations
- Cache misses (requires `perf` on Linux)

2. Linux `perf` Tool

```
# Record cache misses  
perf record -e cache-misses,cache-references ./program  
  
# Analyze  
perf report
```

Metrics:

- **cache-misses:** High value indicates false sharing
- **cache-references:** Total cache accesses
- **Miss rate:** cache-misses / cache-references

False sharing symptoms:

- Miss rate > 10%
- High L1/L2 miss rates
- CPU time in atomic operations

3. Manual Inspection

Check struct layouts:

```
go build -gcflags="-m" # Escape analysis
```

Look for:

- Small structs (< 64 bytes) accessed by multiple goroutines
- Arrays of small structs
- Adjacent atomic variables

Preventing False Sharing

Strategy 1: Padding

```
type PaddedCounter struct {
    value int64
    _     [56]byte // 64 - 8 = 56
}
```

Trade-off: Memory usage vs. performance.

Strategy 2: Alignment

```
// Ensure 64-byte alignment
type AlignedCounter struct {
    _     [0]uint64 // Force 8-byte alignment
    value int64
    _     [7]uint64 // Pad to 64 bytes
}
```

Strategy 3: Separate Caching

```
// Don't share data
type LocalCounter struct {
    local map[int64]int64 // Goroutine ID → count
}

func (c *LocalCounter) Increment() {
    id := goroutineID()
    c.local[id]++
}

func (c *LocalCounter) Total() int64 {
    sum := int64(0)
    for _, v := range c.local {
        sum += v
    }
    return sum
}
```

Strategy 4: Use sync.Pool

```
var pool = sync.Pool{
    New: func() interface{} {
        return &Counter{}
    },
}

func work() {
    c := pool.Get().(*Counter)
    defer pool.Put(c)

    // Each goroutine gets own counter (no sharing)
    c.value++
}
```

When to Worry About False Sharing

Optimize if:

1. High contention on atomic variables
2. Benchmarks show cache-line bouncing
3. Multiple goroutines update adjacent data
4. Latency-critical code (< 100ns per operation)

Don't optimize if:

1. Low contention
2. Data not "hot" (infrequent access)
3. Readability/maintainability matters more
4. Unsure (profile first!)

Premature optimization is root of all evil. Profile, then optimize.

Benchmarking False Sharing

```
package main

import (
    "sync"
    "sync/atomic"
    "testing"
)

// Without padding
type UnpaddedCounter struct {
    a, b int64
}

// With padding
type PaddedCounter struct {
```

```
a int64
_ [56]byte
b int64
}

func BenchmarkUnpadded(b *testing.B) {
    var c UnpaddedCounter
    var wg sync.WaitGroup

    wg.Add(2)
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&c.a, 1)
        }
    }()
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&c.b, 1)
        }
    }()
    wg.Wait()
}

func BenchmarkPadded(b *testing.B) {
    var c PaddedCounter
    var wg sync.WaitGroup

    wg.Add(2)
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&c.a, 1)
        }
    }()
    go func() {
        defer wg.Done()
        for i := 0; i < b.N; i++ {
            atomic.AddInt64(&c.b, 1)
        }
    }()
    wg.Wait()
}

// Run:
// go test -bench=. -benchmem
// BenchmarkUnpadded-8      2000000000 ~80 ns/op
// BenchmarkPadded-8       2000000000 ~8 ns/op
```

Interview Traps

Trap 1: "No data race, so no problem"

Wrong: "Goroutines access different variables, no synchronization needed."

Correct: "Even without logical sharing, false sharing causes performance degradation due to cache-line bouncing. Need padding to separate hot variables."

Trap 2: "Padding fixes all performance problems"

Wrong: "Add padding everywhere."

Correct: "Padding trades memory for performance. Only optimize hot paths with high contention. Profile first. Padding increases struct size, can hurt cache locality for other access patterns."

Trap 3: "Cache line is 32 bytes"

Wrong: Assumes outdated architecture.

Correct: "Modern x86/ARM CPUs have 64-byte cache lines. Always use 64 bytes for padding. Cache line size can vary (128 bytes on some systems), but 64 is safe default."

Trap 4: "Atomics eliminate false sharing"

Wrong: "atomic.AddInt64 is thread-safe, so no false sharing."

Correct: "Atomics prevent data races but don't prevent false sharing. Atomics on adjacent variables still cause cache-line invalidation. Need padding or separate cache lines."

Key Takeaways

1. False sharing = hardware sharing, not logical sharing
2. Cache line is typically 64 bytes
3. Cache coherency causes invalidation (MESI protocol)
4. Padding separates hot variables into different cache lines
5. Profile before optimizing (cache-misses with perf)
6. Trade memory for performance (padding increases size)
7. Affects latency-critical code (< 100ns operations)
8. Go's atomic operations still have false sharing

Exercises

1. Write benchmark comparing padded vs. unpadded counters. Measure speedup.
2. Create array of 8 counters. Benchmark with/without padding between elements.
3. Use Linux `perf` to measure cache-miss rate on padded vs. unpadded.
4. Calculate memory overhead of padding 1000 structs with 64-byte padding.
5. Explain why padding might HURT performance in some cases (hint: cache locality for sequential access).

Next: [Section 04: Patterns](#) - Learn production-ready concurrency patterns.

Congratulations!

You've completed **Section 03: Classic Problems**. You now understand:

- Race conditions (data races vs. logic bugs)
- Deadlocks (circular wait conditions)
- Livelocks (active but no progress)
- Starvation (resource denial)
- False sharing (cache-line bouncing)

These are the foundational bugs you'll encounter in concurrent systems. Always keep them in mind when writing goroutines and channels.

Next up: Patterns for building robust, scalable concurrent systems.