

Bounded Parallelism

What is Bounded Parallelism?

Problem: Processing N items, but can only handle M concurrent operations.

Examples:

- Process 10,000 files, but only 100 concurrent file handles
- Make 1,000 API calls, but rate limit is 10/sec
- Database batch: 50,000 inserts, but pool has 25 connections

Solution: Limit concurrency while still processing all items.

Semaphore Pattern

```
func processBatchWithSemaphore(items []Item, maxConcurrent int) error {
    sem := make(chan struct{}, maxConcurrent)
    var wg sync.WaitGroup
    errCh := make(chan error, 1)

    for _, item := range items {
        wg.Add(1)

        // Acquire semaphore
        sem <- struct{}{}

        go func(i Item) {
            defer wg.Done()
            defer func() { <-sem }() // Release semaphore

            if err := process(i); err != nil {
                select {
                case errCh <- err:
                default:
                }
            }
        }(item)
    }

    wg.Wait()
    close(errCh)

    // Return first error
    for err := range errCh {
        return err
    }

    return nil
}
```

Worker Pool Pattern

```
func processBatchWithWorkerPool(items []Item, numWorkers int) error {
    tasks := make(chan Item, numWorkers*2)
    errors := make(chan error, 1)
    var wg sync.WaitGroup

    // Start workers
    for i := 0; i < numWorkers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            for item := range tasks {
                if err := process(item); err != nil {
                    select {
                    case errors <- err:
                    default:
                    }
                    return
                }
            }
        }()
    }

    // Send tasks
    for _, item := range items {
        tasks <- item
    }
    close(tasks)

    // Wait for completion
    wg.Wait()
    close(errors)

    // Check errors
    for err := range errors {
        return err
    }

    return nil
}
```

Rate-Limited Batch Processing

```
import "golang.org/x/time/rate"

func processBatchWithRateLimit(items []Item, rps int) error {
    limiter := rate.NewLimiter(rate.Limit(rps), rps)
```

```

for _, item := range items {
    // Wait for rate limiter
    if err := limiter.Wait(context.Background()); err != nil {
        return err
    }

    if err := process(item); err != nil {
        return err
    }
}

return nil
}

// Concurrent + rate limited:
func processBatchConcurrentRateLimited(items []Item, workers, rps int) error {
    limiter := rate.NewLimiter(rate.Limit(rps), rps)
    sem := make(chan struct{}, workers)
    var wg sync.WaitGroup

    for _, item := range items {
        // Wait for rate limit
        if err := limiter.Wait(context.Background()); err != nil {
            return err
        }

        // Acquire semaphore
        sem <- struct{}{}
        wg.Add(1)

        go func(i Item) {
            defer wg.Done()
            defer func() { <-sem }()
            process(i)
        }(item)
    }

    wg.Wait()
    return nil
}

```

Batching with Channels

```

func processBatchWithChannels(items []Item, batchSize int) error {
    batches := make(chan []Item, 10)

    // Send batches
    go func() {
        defer close(batches)

```

```

        for i := 0; i < len(items); i += batchSize {
            end := i + batchSize
            if end > len(items) {
                end = len(items)
            }

            batch := items[i:end]
            batches <- batch
        }
    }()
}

// Process batches concurrently
var wg sync.WaitGroup
numWorkers := runtime.NumCPU()

for i := 0; i < numWorkers; i++ {
    wg.Add(1)
    go func() {
        defer wg.Done()
        for batch := range batches {
            processBatch(batch)
        }
    }()
}
}

wg.Wait()
return nil
}

```

Real Example: Image Processing

```

func processImages(imageFiles []string, maxConcurrent int) error {
    sem := make(chan struct{}, maxConcurrent)
    var wg sync.WaitGroup

    for _, filename := range imageFiles {
        wg.Add(1)
        sem <- struct{}{}
    }

    go func(fn string) {
        defer wg.Done()
        defer func() { <-sem }()

        // Load image
        img, err := loadImage(fn)
        if err != nil {
            log.Printf("Failed to load %s: %v", fn, err)
            return
        }
    }
}

```

```

    // Resize
    resized := resize(img, 800, 600)

    // Save
    outFile := strings.Replace(fn, ".jpg", "_thumb.jpg", 1)
    if err := saveImage(outFile, resized); err != nil {
        log.Printf("Failed to save %s: %v", outFile, err)
    }
}

wg.Wait()
return nil
}

```

Real Example: API Batch Calls

```

func fetchUsersInBatches(userIDs []int, batchSize, maxConcurrent int) ([]*User,
error) {
    type result struct {
        users []*User
        err   error
    }

    results := make(chan result, (len(userIDs)/batchSize)+1)
    sem := make(chan struct{}, maxConcurrent)
    var wg sync.WaitGroup

    // Process in batches
    for i := 0; i < len(userIDs); i += batchSize {
        end := i + batchSize
        if end > len(userIDs) {
            end = len(userIDs)
        }

        batch := userIDs[i:end]

        wg.Add(1)
        sem <- struct{}{}

        go func(ids []int) {
            defer wg.Done()
            defer func() { <-sem }()
            // API call with batch
            users, err := api.FetchUsers(ids)
            results <- result{users: users, err: err}
        }(batch)
    }
}

```

```
// Wait and close
go func() {
    wg.Wait()
    close(results)
}()

// Collect results
var allUsers []*User
for r := range results {
    if r.err != nil {
        return nil, r.err
    }
    allUsers = append(allUsers, r.users...)
}

return allUsers, nil
}
```

Bounded Parallelism with Context

```
func processBatchWithContext(ctx context.Context, items []Item, workers int) error {
    tasks := make(chan Item, workers)
    errors := make(chan error, 1)
    var wg sync.WaitGroup

    // Start workers
    for i := 0; i < workers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()

            for {
                select {
                case <-ctx.Done():
                    return

                case item, ok := <-tasks:
                    if !ok {
                        return
                    }

                    if err := process(ctx, item); err != nil {
                        select {
                        case errors <- err:
                        default:
                            }
                        return
                    }
                }
            }
        }
    }

    // Wait for all workers to finish
    wg.Wait()
    close(tasks)
    close(errors)

    var err error
    for {
        select {
        case err = <-errors:
            return err
        }
    }
}
```

```

        }
    }()
}

// Send tasks
go func() {
    defer close(tasks)

    for _, item := range items {
        select {
        case <-ctx.Done():
            return
        case tasks <- item:
        }
    }
}()

// Wait for completion
wg.Wait()
close(errors)

// Check for errors
select {
case err := <-errors:
    return err
default:
    return ctx.Err()
}
}
}

```

Dynamic Worker Count

Adjust workers based on system load.

```

type DynamicPool struct {
    tasks      chan Task
    minWorkers int
    maxWorkers int
    currentWorkers atomic.Int32
    wg         sync.WaitGroup
    quit       chan struct{}
}

func (dp *DynamicPool) Start() {
    // Start minimum workers
    for i := 0; i < dp.minWorkers; i++ {
        dp.addWorker()
    }

    // Monitor queue depth
    go dp.scaleWorkers()
}

```

```

}

func (dp *DynamicPool) scaleWorkers() {
    ticker := time.NewTicker(5 * time.Second)
    defer ticker.Stop()

    for {
        select {
        case <-ticker.C:
            queueDepth := len(dp.tasks)
            workers := int(dp.currentWorkers.Load())

            // Scale up if queue backing up
            if queueDepth > workers*2 && workers < dp.maxWorkers {
                dp.addWorker()
                log.Printf("Scaled up to %d workers", workers+1)
            }

            // Scale down if idle
            if queueDepth == 0 && workers > dp.minWorkers {
                // Signal one worker to stop (implementation omitted)
            }
        case <-dp.quit:
            return
        }
    }
}

func (dp *DynamicPool) addWorker() {
    dp.currentWorkers.Add(1)
    dp.wg.Add(1)

    go func() {
        defer dp.wg.Done()
        defer dp.currentWorkers.Add(-1)

        for {
            select {
            case task := <-dp.tasks:
                task.Execute()
            case <-dp.quit:
                return
            }
        }
    }()
}

```

Parallel Map/Reduce

```

func parallelMap(items []Item, workers int, fn func(Item) Result) []Result {
    tasks := make(chan Item, len(items))
    results := make(chan Result, len(items))

    // Start workers
    var wg sync.WaitGroup
    for i := 0; i < workers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            for item := range tasks {
                results <- fn(item)
            }
        }()
    }

    // Send tasks
    for _, item := range items {
        tasks <- item
    }
    close(tasks)

    // Wait and close results
    go func() {
        wg.Wait()
        close(results)
    }()

    // Collect results
    var output []Result
    for result := range results {
        output = append(output, result)
    }

    return output
}

func parallelReduce(items []Item, workers int, fn func(Item, Item) Item) Item {
    if len(items) == 0 {
        var zero Item
        return zero
    }

    // Divide into chunks
    chunkSize := (len(items) + workers - 1) / workers
    results := make(chan Item, workers)

    for i := 0; i < len(items); i += chunkSize {
        end := i + chunkSize
        if end > len(items) {
            end = len(items)
        }
        go func(i, end int) {
            for item := range tasks {
                results <- fn(item, item)
            }
        }(i, end)
    }

    var sum Item
    for result := range results {
        sum = fn(sum, result)
    }

    return sum
}

```

```

    }

    chunk := items[i:end]

    go func(c []Item) {
        acc := c[0]
        for _, item := range c[1:] {
            acc = fn(acc, item)
        }
        results <- acc
    }(chunk)
}

// Reduce results
partials := make([]Item, 0, workers)
for i := 0; i < workers && i < len(items); i++ {
    partials = append(partials, <-results)
}

result := partials[0]
for _, partial := range partials[1:] {
    result = fn(result, partial)
}

return result
}

```

Choosing Worker Count

```

// CPU-bound work:
workers := runtime.NumCPU()

// I/O-bound work (network, disk):
workers := runtime.NumCPU() * 10

// Memory-intensive:
workers := runtime.NumCPU() / 2

// Rate-limited API:
workers := rateLimit / avgRequestsPerSecond

// Database operations:
workers := db.Stats().MaxOpenConnections

// Benchmark to find optimal:
func BenchmarkWorkers(b *testing.B) {
    workerCounts := []int{1, 2, 4, 8, 16, 32}

    for _, workers := range workerCounts {
        b.Run(fmt.Sprintf("workers=%d", workers), func(b *testing.B) {

```

```

        for i := 0; i < b.N; i++ {
            processBatch(items, workers)
        }
    })
}
}

```

Common Mistakes

Mistake 1: Unbounded Concurrency

```

// WRONG: Creates 1 million goroutines
for _, item := range items { // 1 million items
    go process(item) // 1 million goroutines!
}

// Fix: Bounded concurrency
sem := make(chan struct{}, 100)
for _, item := range items {
    sem <- struct{}{}
    go func(i Item) {
        defer func() { <-sem }()
        process(i)
    }(item)
}

```

Mistake 2: Wrong Worker Count

```

// WRONG: Way too many workers for CPU-bound work
workers := 1000 // But only 8 CPUs!

// Fix: Match to resource
workers := runtime.NumCPU()

```

Mistake 3: Not Waiting for Completion

```

// WRONG: Returns before processing done
func processBatch(items []Item) {
    for _, item := range items {
        go process(item)
    }
    return // WRONG: Doesn't wait!
}

// Fix: Wait for goroutines
func processBatch(items []Item) {
    var wg sync.WaitGroup
    for _, item := range items {

```

```

wg.Add(1)
go func(i Item) {
    defer wg.Done()
    process(i)
}(item)
}
wg.Wait()
}

```

Performance Benchmarks

```

func BenchmarkBoundedVsUnbounded(b *testing.B) {
    items := makeTestItems(1000)

    b.Run("unbounded", func(b *testing.B) {
        for i := 0; i < b.N; i++ {
            var wg sync.WaitGroup
            for _, item := range items {
                wg.Add(1)
                go func(i Item) {
                    defer wg.Done()
                    process(i)
                }(item)
            }
            wg.Wait()
        }
    })

    b.Run("bounded-100", func(b *testing.B) {
        for i := 0; i < b.N; i++ {
            processBatchWithSemaphore(items, 100)
        }
    })
}

// Results (typical):
// unbounded:    High memory, context switch overhead
// bounded-100:   Lower memory, better throughput

```

Interview Questions

Q: "Semaphore vs. worker pool for bounded parallelism?"

"Semaphore: goroutine per task, acquire/release pattern, simple for short tasks, higher memory (goroutine per item). Worker pool: fixed goroutines, task queue, better for long-running, lower memory (fixed goroutines). Choose semaphore for simple short tasks, worker pool for sustained load or long tasks."

Q: "How do you choose optimal worker count?"

"Depends on bottleneck. CPU-bound: NumCPU (avoid oversubscription). I/O-bound: NumCPU * 10-100 (goroutines block on I/O). Memory-intensive: Lower (avoid OOM). Rate-limited: Match rate limit. Database: Match connection pool size. Best practice: Benchmark different worker counts, measure throughput and latency."

Q: "What happens if you don't bound concurrency?"

"Memory exhaustion (each goroutine ~2KB), GC pressure (tracking millions of goroutines), scheduler overhead (context switching), resource exhaustion (file handles, connections), system instability. Example: Processing 1M items unbounded = 2GB + overhead. With 100 workers = 200KB."

Key Takeaways

1. Always bound concurrency for large batches
2. Semaphore: simple, goroutine per task
3. Worker pool: efficient, fixed goroutines
4. CPU-bound: workers = NumCPU
5. I/O-bound: workers = NumCPU * 10-100
6. Rate limiting complements bounded parallelism
7. Use WaitGroup to wait for completion
8. Context for cancellation
9. Benchmark to find optimal worker count
10. Monitor queue depth and latency

Exercises

1. Implement bounded batch processor with semaphore, benchmark different limits.
2. Build worker pool with dynamic worker count (scales 1-100 based on queue).
3. Create rate-limited batch API client (100 req/sec, 10 concurrent).
4. Benchmark: Compare semaphore vs. worker pool for 10,000 items.
5. Implement parallel map/reduce with bounded workers.

Next: [avoiding-goroutine-leaks.md](#) - Detecting and preventing goroutine leaks.