

# Concurrent File I/O

## File System Basics

### Key constraints:

- **Single disk head:** Physical disks can't truly parallelize reads from different locations
- **OS buffering:** Operating system caches frequently accessed files
- **SSDs change game:** Solid-state drives benefit more from concurrency
- **Network filesystems:** NFS,S3 can benefit significantly from parallelism

### When concurrency helps:

- Multiple files (each file independent)
- Network I/O (S3, HTTP downloads)
- CPU-bound processing of file data
- SSDs with multiple channels

### When concurrency doesn't help:

- Single large file on spinning disk
- Sequential read patterns

## Reading Files Concurrently

### Multiple Files in Parallel

```
func readMultipleFiles(filenames []string) (map[string][]byte, error) {
    type result struct {
        filename string
        data     []byte
        err      error
    }

    results := make(chan result, len(filenames))

    // Fan-out: Read files concurrently
    for _, filename := range filenames {
        go func(fn string) {
            data, err := os.ReadFile(fn)
            results <- result{filename: fn, data: data, err: err}
        }(filename)
    }

    // Fan-in: Collect results
    fileData := make(map[string][]byte)
    for range filenames {
        r := <-results
        if r.err != nil {
            return nil, fmt.Errorf("%s: %w", r.filename, r.err)
        }
        fileData[r.filename] = r.data
    }
}
```

```
    }

    return fileData, nil
}
```

## Processing Large File in Chunks

```
func processLargeFile(filename string, chunkSize int) error {
    file, err := os.Open(filename)
    if err != nil {
        return err
    }
    defer file.Close()

    chunks := make(chan []byte, 10)
    results := make(chan error, 1)

    // Start processor goroutines
    var wg sync.WaitGroup
    numWorkers := runtime.NumCPU()

    for i := 0; i < numWorkers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            for chunk := range chunks {
                if err := processChunk(chunk); err != nil {
                    select {
                    case results <- err:
                    default:
                    }
                    return
                }
            }
        }()
    }

    // Read file in chunks
    buffer := make([]byte, chunkSize)
    for {
        n, err := file.Read(buffer)
        if n > 0 {
            // Copy buffer (reader reuses it)
            chunk := make([]byte, n)
            copy(chunk, buffer[:n])
            chunks <- chunk
        }
        if err == io.EOF {
            break
        }
    }
}
```

```

        if err != nil {
            close(chunks)
            return err
        }
    }

close(chunks)
wg.Wait()

select {
case err := <-results:
    return err
default:
    return nil
}
}

```

## Buffered Reading

```

func readFileBuffered(filename string) error {
    file, err := os.Open(filename)
    if err != nil {
        return err
    }
    defer file.Close()

    // Buffered reader (default 4KB buffer)
    reader := bufio.NewReader(file)

    for {
        line, err := reader.ReadString('\n')
        if err == io.EOF {
            break
        }
        if err != nil {
            return err
        }

        processLine(line)
    }

    return nil
}

```

## Writing Files Concurrently

### Multiple Files

```

func writeMultipleFiles(data map[string][]byte) error {
    errCh := make(chan error, len(data))
    var wg sync.WaitGroup

    for filename, content := range data {
        wg.Add(1)
        go func(fn string, d []byte) {
            defer wg.Done()

            if err := os.WriteFile(fn, d, 0644); err != nil {
                errCh <- fmt.Errorf("%s: %w", fn, err)
            }
        }(filename, content)
    }

    wg.Wait()
    close(errCh)

    // Check for errors
    for err := range errCh {
        return err // Return first error
    }

    return nil
}

```

## Buffered Writing

```

func writeFileBuffered(filename string, lines []string) error {
    file, err := os.Create(filename)
    if err != nil {
        return err
    }
    defer file.Close()

    // Buffered writer (default 4KB buffer)
    writer := bufio.NewWriter(file)
    defer writer.Flush() // Important!

    for _, line := range lines {
        if _, err := writer.WriteString(line + "\n"); err != nil {
            return err
        }
    }

    return writer.Flush()
}

```

## Concurrent Writes to Single File (with Synchronization)

```

type ConcurrentWriter struct {
    file *os.File
    mu   sync.Mutex
}

func NewConcurrentWriter(filename string) (*ConcurrentWriter, error) {
    file, err := os.Create(filename)
    if err != nil {
        return nil, err
    }

    return &ConcurrentWriter{file: file}, nil
}

func (cw *ConcurrentWriter) Write(data []byte) error {
    cw.mu.Lock()
    defer cw.mu.Unlock()

    _, err := cw.file.Write(data)
    return err
}

func (cw *ConcurrentWriter) Close() error {
    cw.mu.Lock()
    defer cw.mu.Unlock()

    return cw.file.Close()
}

// Usage:
writer, _ := NewConcurrentWriter("output.txt")
defer writer.Close()

var wg sync.WaitGroup
for i := 0; i < 10; i++ {
    wg.Add(1)
    go func(n int) {
        defer wg.Done()
        data := fmt.Sprintf("Line %d\n", n)
        writer.Write([]byte(data))
    }(i)
}
wg.Wait()

```

## Directory Traversal

### Sequential Walk

```

func walkDirectory(root string) ([]string, error) {
    var files []string

```

```

    err := filepath.Walk(root, func(path string, info os.FileInfo, err error) error
{
    if err != nil {
        return err
    }

    if !info.IsDir() {
        files = append(files, path)
    }

    return nil
})

return files, err
}

```

## Concurrent Walk

```

func walkDirectoryConcurrent(root string) ([]string, error) {
    var (
        mu     sync.Mutex
        files []string
        wg     sync.WaitGroup
    )

    err := filepath.Walk(root, func(path string, info os.FileInfo, err error) error
{
    if err != nil {
        return err
    }

    if !info.IsDir() {
        wg.Add(1)
        go func(p string) {
            defer wg.Done()

            // Process file (e.g., read, parse)
            processFile(p)

            // Add to results
            mu.Lock()
            files = append(files, p)
            mu.Unlock()
        }(path)
    }
}

    return nil
})

wg.Wait()

```

```
    return files, err
}
```

## Parallel Walk with Worker Pool

```
func walkWithWorkerPool(root string, workers int) error {
    type task struct {
        path string
        info os.FileInfo
    }

    tasks := make(chan task, 100)
    var wg sync.WaitGroup

    // Start workers
    for i := 0; i < workers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            for t := range tasks {
                processFile(t.path, t.info)
            }
        }()
    }

    // Walk directory
    err := filepath.Walk(root, func(path string, info os.FileInfo, err error) error {
        if err != nil {
            return err
        }

        if !info.IsDir() {
            tasks <- task{path: path, info: info}
        }

        return nil
    })
}

close(tasks)
wg.Wait()

return err
}
```

## Pipeline: Read → Process → Write

```
func processFilePipeline(inputFiles []string, outputDir string) error {
    // Stage 1: Read files
```

```

fileChans := make(chan fileData, len(inputFiles))

go func() {
    defer close(fileChans)
    for _, filename := range inputFiles {
        data, err := os.ReadFile(filename)
        if err != nil {
            log.Printf("Error reading %s: %v", filename, err)
            continue
        }
        fileChans <- fileData{filename: filename, data: data}
    }
}()

// Stage 2: Process (multiple workers)
processed := make(chan fileData, 10)
var wg sync.WaitGroup
numWorkers := runtime.NumCPU()

for i := 0; i < numWorkers; i++ {
    wg.Add(1)
    go func() {
        defer wg.Done()
        for fd := range fileChans {
            // Process data
            result := transform(fd.data)
            processed <- fileData{
                filename: filepath.Join(outputDir, filepath.Base(fd.filename)),
                data:     result,
            }
        }
    }()
}

// Close processed channel when all workers done
go func() {
    wg.Wait()
    close(processed)
}()

// Stage 3: Write results
for fd := range processed {
    if err := os.WriteFile(fd.filename, fd.data, 0644); err != nil {
        return err
    }
}

return nil
}

type fileData struct {
    filename string
}

```

```
    data      []byte
}
```

## Real Example: Log File Analyzer

```
type LogAnalyzer struct {
    workers int
}

func (la *LogAnalyzer) Analyze(logFiles []string) (*Stats, error) {
    lines := make(chan string, 1000)
    results := make(chan Stats, la.workers)

    var wg sync.WaitGroup

    // Workers: Process lines
    for i := 0; i < la.workers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()

            stats := Stats{}
            for line := range lines {
                stats.Process(line)
            }

            results <- stats
        }()
    }

    // Read all log files
    for _, filename := range logFiles {
        file, err := os.Open(filename)
        if err != nil {
            close(lines)
            return nil, err
        }

        scanner := bufio.NewScanner(file)
        for scanner.Scan() {
            lines <- scanner.Text()
        }

        file.Close()
    }

    close(lines)
    wg.Wait()
    close(results)
```

```
// Merge results
finalStats := Stats{}
for stats := range results {
    finalStats.Merge(stats)
}

return &finalStats, nil
}

type Stats struct {
    ErrorCount    int
    WarningCount int
    TotalLines    int
}

func (s *Stats) Process(line string) {
    s.TotalLines++
    if strings.Contains(line, "ERROR") {
        s.ErrorCount++
    } else if strings.Contains(line, "WARN") {
        s.WarningCount++
    }
}

func (s *Stats) Merge(other Stats) {
    s.ErrorCount += other.ErrorCount
    s.WarningCount += other.WarningCount
    s.TotalLines += other.TotalLines
}
```

# File Watching

```

        go processFile(event.Name)

    case event.Op&fsnotify.Write == fsnotify.Write:
        log.Println("Modified:", event.Name)

    case event.Op&fsnotify.Remove == fsnotify.Remove:
        log.Println("Removed:", event.Name)
    }

    case err := <-watcher.Errors:
        log.Println("Error:", err)
    }
}

err = watcher.Add(dir)
if err != nil {
    return err
}

<-done
return nil
}

```

## Common Mistakes

### Mistake 1: Forgetting to Flush Buffered Writer

```

// WRONG: Data lost
writer := bufio.NewWriter(file)
writer.WriteString("data")
file.Close() // Doesn't flush buffer!

// Fix:
writer := bufio.NewWriter(file)
writer.WriteString("data")
writer.Flush() // Flush before closing
file.Close()

```

### Mistake 2: Not Checking Scanner Errors

```

// WRONG: Ignores errors
scanner := bufio.NewScanner(file)
for scanner.Scan() {
    process(scanner.Text())
}
// Missing scanner.Err() check!

// Fix:

```

```

scanner := bufio.NewScanner(file)
for scanner.Scan() {
    process(scanner.Text())
}
if err := scanner.Err(); err != nil {
    return err
}

```

### Mistake 3: Race on Shared Slice

```

// WRONG: Race condition
var files []string

filepath.Walk(root, func(path string, info os.FileInfo, err error) error {
    go func() {
        files = append(files, path) // RACE!
    }()
    return nil
})

// Fix: Use mutex
var (
    files []string
    mu     sync.Mutex
)

filepath.Walk(root, func(path string, info os.FileInfo, err error) error {
    go func() {
        mu.Lock()
        files = append(files, path)
        mu.Unlock()
    }()
    return nil
})

```

### Performance Tips

1. **Use buffered I/O:** `bufio.Reader/Writer` for small reads/writes
2. **Parallelize multiple files** (easy wins)
3. **For single file:** Process after reading, don't try to parallelize reads
4. **Worker pools for processing:** Match to NumCPU
5. **SSDs benefit from concurrency** more than spinning disks

### Interview Questions

**Q: "When does concurrent file I/O help?"**

"Helps when: multiple independent files, network filesystems (S3, NFS), CPU-bound processing after read, SSDs. Doesn't help: single file sequential read on spinning disk (single head can't parallelize). Pattern: read file sequentially, parallelize processing in worker pool."

### **Q: "How do you handle concurrent writes to same file?"**

"File writes aren't thread-safe. Solutions: 1) Mutex around writes (serializes, simple), 2) Channel to single writer goroutine (serializes, clean), 3) Append atomic writes (limited use), 4) Write to separate files then merge.  
Choose: mutex if simple, channel if better separation."

### **Q: "What's the difference between os.ReadFile and bufio?"**

"os.ReadFile: reads entire file to memory, single call. bufio.Reader: reads in chunks, preserves memory. Use ReadFile for small files (< MB), bufio for large files or line-by-line processing. bufio uses buffer (default 4KB), reduces syscalls."

## **Key Takeaways**

1. **Concurrent file reads help for multiple files**
2. **Single file: read sequential, parallelize processing**
3. **Always flush buffered writers before closing**
4. **Check scanner.Err() after loop**
5. **Concurrent writes need synchronization (mutex/channel)**
6. **Use bufio for large files**
7. **Worker pool for processing file contents**
8. **Directory traversal benefits from parallelism**
9. **SSDs benefit more from concurrency than spinning disks**
10. **File watching enables real-time processing**

## **Exercises**

1. Build concurrent directory walker with worker pool (N workers).
2. Implement log analyzer that processes multiple GB files in parallel.
3. Create pipeline: read CSV → parse → transform → write JSON.
4. Benchmark: Compare sequential vs. concurrent file processing (10 files, 100 files).
5. Build file watcher that processes new files as they arrive.

**Next:** [graceful-shutdown.md](#) - Cleanly shutting down concurrent systems.