

Benchmarking Concurrent Code

Go Benchmark Basics

```
func BenchmarkFunction(b *testing.B) {
    for i := 0; i < b.N; i++ {
        function()
    }
}

// Run: go test -bench=.
// Output: BenchmarkFunction-8 1000000 1234 ns/op
//           Name          cpus  iterations  time/op
```

Key metrics:

- **ns/op**: Time per operation
- **B/op**: Bytes allocated per operation
- **allocs/op**: Allocations per operation

Benchmarking Concurrent Operations

```
func BenchmarkConcurrentCounter(b *testing.B) {
    var counter atomic.Int64

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            counter.Add(1)
        }
    })
}

// RunParallel: Runs function in GOMAXPROCS goroutines
// pb.Next(): Returns true while benchmark running
```

Comparing Mutex vs. Atomic

```
func BenchmarkMutexIncrement(b *testing.B) {
    var (
        counter int
        mu      sync.Mutex
    )

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            mu.Lock()
            counter++
            mu.Unlock()
        }
    })
}
```

```

        }
    })
}

func BenchmarkAtomicIncrement(b *testing.B) {
    var counter atomic.Int64

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            counter.Add(1)
        }
    })
}

// Results (typical on 8-core machine):
// BenchmarkMutexIncrement-8      20000000      85 ns/op
// BenchmarkAtomicIncrement-8      50000000      30 ns/op
//
// Atomic is ~3x faster for simple increment

```

Benchmarking Channels

```

func BenchmarkUnbufferedChannel(b *testing.B) {
    ch := make(chan int)

    go func() {
        for range ch {
        }
    }()
}

b.ResetTimer()

for i := 0; i < b.N; i++ {
    ch <- i
}

close(ch)
}

func BenchmarkBufferedChannel(b *testing.B) {
    ch := make(chan int, 100)

    go func() {
        for range ch {
        }
    }()
}

b.ResetTimer()

for i := 0; i < b.N; i++ {

```

```

        ch <- i
    }

    close(ch)
}

// Results:
// BenchmarkUnbufferedChannel-8      5000000      312 ns/op
// BenchmarkBufferedChannel-8       20000000      65 ns/op
//
// Buffered channels ~5x faster (less synchronization)

```

Benchmarking Different Concurrency Levels

```

func BenchmarkWorkerPool(b *testing.B) {
    workerCounts := []int{1, 2, 4, 8, 16, 32}

    for _, workers := range workerCounts {
        b.Run(fmt.Sprintf("workers=%d", workers), func(b *testing.B) {
            pool := NewWorkerPool(workers)
            defer pool.Shutdown()

            b.ResetTimer()

            for i := 0; i < b.N; i++ {
                pool.Submit(func() {
                    doWork()
                })
            }
        })

        pool.Wait()
    }
}

// Output:
// BenchmarkWorkerPool/workers=1-8      100000      15234 ns/op
// BenchmarkWorkerPool/workers=2-8      200000      7891 ns/op
// BenchmarkWorkerPool/workers=4-8      400000      4123 ns/op
// BenchmarkWorkerPool/workers=8-8      500000      3456 ns/op
// BenchmarkWorkerPool/workers=16-8     450000      3789 ns/op
// BenchmarkWorkerPool/workers=32-8     400000      4234 ns/op
//
// Optimal: 8 workers (matches CPU cores)

```

Memory Allocation Benchmarks

```

func BenchmarkMapWithMutex(b *testing.B) {
    m := make(map[string]int)
    var mu sync.RWMutex

    b.ReportAllocs()
    b.ResetTimer()

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            mu.Lock()
            m["key"] = 42
            mu.Unlock()
        }
    })
}

func BenchmarkSyncMap(b *testing.B) {
    var m sync.Map

    b.ReportAllocs()
    b.ResetTimer()

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            m.Store("key", 42)
        }
    })
}

// Run: go test -bench=. -benchmem
//
// BenchmarkMapWithMutex-8      10000000      145 ns/op      0 B/op      0 allocs/op
// BenchmarkSyncMap-8           20000000      75 ns/op      8 B/op      1 allocs/op
//
// sync.Map has allocation overhead but faster under contention

```

Benchmarking Lock Contention

```

func BenchmarkLockContention(b *testing.B) {
    contentionLevels := []int{1, 10, 100}

    for _, workDuration := range contentionLevels {
        b.Run(fmt.Sprintf("work=%dns", workDuration), func(b *testing.B) {
            var (
                counter int
                mu      sync.Mutex
            )

            b.RunParallel(func(pb *testing.PB) {

```

```

        for pb.Next() {
            mu.Lock()

                // Simulate work under lock
                start := time.Now()
                for time.Since(start) < time.Duration(workDuration) {
                }

                counter++
                mu.Unlock()
            }
        }
    }
}

// Output shows contention impact:
// BenchmarkLockContention/work=1ns-8      500000000    35 ns/op
// BenchmarkLockContention/work=10ns-8     300000000    45 ns/op
// BenchmarkLockContention/work=100ns-8    100000000   145 ns/op
//
// Longer critical sections → more contention → slower

```

Comparing Synchronization Primitives

```

func BenchmarkMutex(b *testing.B) {
    var (
        data int
        mu   sync.Mutex
    )

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            mu.Lock()
            data++
            mu.Unlock()
        }
    })
}

func BenchmarkRWMutex(b *testing.B) {
    var (
        data int
        mu   sync.RWMutex
    )

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            mu.Lock()
            data++
        }
    })
}

```

```

        mu.Unlock()
    }
})
}

func BenchmarkAtomic(b *testing.B) {
    var data atomic.Int64

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            data.Add(1)
        }
    })
}

func BenchmarkChannel(b *testing.B) {
    ch := make(chan int, 100)

    go func() {
        for range ch {
        }
    }()

    b.ResetTimer()

    for i := 0; i < b.N; i++ {
        ch <- 1
    }

    close(ch)
}

// Results (typical):
// BenchmarkMutex-8      20000000    85 ns/op
// BenchmarkRWMutex-8    15000000   110 ns/op  (slower for writes)
// BenchmarkAtomic-8     50000000    30 ns/op  (fastest)
// BenchmarkChannel-8    10000000   165 ns/op  (slowest, most flexible)

```

Read-Heavy Workload Benchmark

```

func BenchmarkRWMutexReadHeavy(b *testing.B) {
    var (
        data int
        mu   sync.RWMutex
    )

    // 90% reads, 10% writes
    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            if rand.Intn(10) == 0 {

```

```

        // Write
        mu.Lock()
        data++
        mu.Unlock()
    } else {
        // Read
        mu.RLock()
        _ = data
        mu.RUnlock()
    }
}
}

func BenchmarkMutexReadHeavy(b *testing.B) {
    var (
        data int
        mu   sync.Mutex
    )

    b.RunParallel(func(pb *testing.PB) {
        for pb.Next() {
            if rand.Intn(10) == 0 {
                mu.Lock()
                data++
                mu.Unlock()
            } else {
                mu.Lock()
                _ = data
                mu.Unlock()
            }
        }
    })
}

// Results for read-heavy:
// BenchmarkRWMutexReadHeavy-8      50000000      35 ns/op  (wins!)
// BenchmarkMutexReadHeavy-8        20000000      85 ns/op
//
// RWMutex better when reads >> writes

```

Benchmarking with Profiling

```

# CPU profile
go test -bench=. -cpuprofile=cpu.prof
go tool pprof cpu.prof

# Memory profile
go test -bench=. -memprofile=mem.prof
go tool pprof mem.prof

```

```
# Blocking profile
go test -bench=. -blockprofile=block.prof
go tool pprof block.prof
```

Interpreting Benchmark Results

```
// Example output:
// BenchmarkFunction-8      1000000      1234 ns/op      456 B/op      7 allocs/op

// -8: GOMAXPROCS (CPU cores used)
// 1000000: Number of iterations
// 1234 ns/op: Time per operation
// 456 B/op: Bytes allocated per operation
// 7 allocs/op: Number of allocations per operation
```

Lower is better for all metrics!

Benchmark Comparison Tool

```
# Run baseline
go test -bench=. > old.txt

# Make changes

# Run new benchmark
go test -bench=. > new.txt

# Compare
go install golang.org/x/perf/cmd/benchstat@latest
benchstat old.txt new.txt

# Output:
# name      old time/op  new time/op  delta
# Function-8    1.23µs ± 2%   0.85µs ± 1%  -30.89% (p=0.000 n=10+10)
```

CI/CD Benchmark Regression Detection

```
// Test: Benchmarks must not regress
func TestBenchmarkRegression(t *testing.T) {
    if testing.Short() {
        t.Skip()
    }

    // Run benchmark
    result := testing.Benchmark(BenchmarkFunction)
```

```

// Check against baseline (e.g., 1000 ns/op)
baseline := int64(1000)
threshold := baseline * 110 / 100 // Allow 10% regression

nsPerOp := result.NsPerOp()
if nsPerOp > threshold {
    t.Errorf("Benchmark regressed: %d ns/op (threshold: %d ns/op)",
        nsPerOp, threshold)
}
}

```

Real-World Example: Cache Benchmark

```

type Cache interface {
    Get(key string) (interface{}, bool)
    Set(key string, value interface{})
}

func BenchmarkCache(b *testing.B) {
    caches := map[string]Cache{
        "mutex": NewMutexCache(),
        "sync.Map": NewSyncMapCache(),
        "sharded": NewShardedCache(16),
    }

    for name, cache := range caches {
        b.Run(name, func(b *testing.B) {
            // Warm up
            for i := 0; i < 1000; i++ {
                cache.Set(fmt.Sprint(i), i)
            }

            b.ResetTimer()
            b.ReportAllocs()

            b.RunParallel(func(pb *testing.PB) {
                i := 0
                for pb.Next() {
                    key := fmt.Sprint(i % 1000)
                    if i%10 == 0 {
                        cache.Set(key, i)
                    } else {
                        cache.Get(key)
                    }
                    i++
                }
            })
        })
    }
}

```

```

// Output:
// BenchmarkCache/mutex-8      5000000    345 ns/op    32 B/op   2 allocs/op
// BenchmarkCache/sync.Map-8    10000000   156 ns/op    48 B/op   3 allocs/op
// BenchmarkCache/sharded-8     15000000   98 ns/op    32 B/op   2 allocs/op
//
// Sharded cache wins! (reduces lock contention)

```

Sub-benchmarks for Different Scenarios

```

func BenchmarkWorkerPool(b *testing.B) {
    taskDurations := []time.Duration{
        0,                      // No work
        10 * time.Microsecond,  // Light
        100 * time.Microsecond, // Medium
        time.Millisecond,       // Heavy
    }

    for _, duration := range taskDurations {
        b.Run(fmt.Sprintf("task=%v", duration), func(b *testing.B) {
            pool := NewWorkerPool(runtime.NumCPU())
            defer poolShutdown()

            b.ResetTimer()

            for i := 0; i < b.N; i++ {
                pool.Submit(func() {
                    time.Sleep(duration)
                })
            }

            pool.Wait()
        })
    }
}

```

Common Mistakes in Benchmarks

Mistake #1: Not Resetting Timer

```

// WRONG: Setup time included in benchmark
func BenchmarkWrong(b *testing.B) {
    cache := setupExpensiveCache() // Counted!

    for i := 0; i < b.N; i++ {
        cache.Get("key")
    }
}

```

```
// Right: Reset timer after setup
func BenchmarkRight(b *testing.B) {
    cache := setupExpensiveCache()
    b.ResetTimer() // Reset!

    for i := 0; i < b.N; i++ {
        cache.Get("key")
    }
}
```

Mistake #2: Compiler Optimization Removing Code

```
// WRONG: Compiler may optimize away unused result
func BenchmarkWrong(b *testing.B) {
    for i := 0; i < b.N; i++ {
        compute() // Result unused, might be optimized out
    }
}

// Right: Store result to prevent optimization
var result int

func BenchmarkRight(b *testing.B) {
    var r int
    for i := 0; i < b.N; i++ {
        r = compute()
    }
    result = r // Assign to package-level var
}
```

Mistake #3: Benchmarking Wrong Thing

```
// WRONG: Benchmarks goroutine creation, not actual work
func BenchmarkWrong(b *testing.B) {
    for i := 0; i < b.N; i++ {
        go doWork() // Just spawns goroutine
    }
}

// Right: Wait for work completion
func BenchmarkRight(b *testing.B) {
    for i := 0; i < b.N; i++ {
        done := make(chan struct{})
        go func() {
            doWork()
            close(done)
        }()
        <-done
    }
}
```

```
    }  
}
```

Interview Questions

Q: "How do you benchmark concurrent code?"

"Use `b.RunParallel()` which runs benchmark function in `GOMAXPROCS` goroutines. Pattern: `b.RunParallel(func(pb *testing.PB) { for pb.Next() { operation() } })`. Measures real concurrent performance with contention. Also benchmark with different concurrency levels using sub-benchmarks to find optimal worker count."

Q: "What's the difference between ns/op, B/op, and allocs/op?"

"`ns/op`: nanoseconds per operation, measures time/performance. `B/op`: bytes allocated per operation, measures memory usage. `allocs/op`: number of allocations per operation, measures GC pressure. Lower is better for all. Optimize in order: correctness → performance (`ns/op`) → memory (`B/op`) → GC pressure (`allocs/op`)."

Q: "How do you detect performance regressions?"

"Use `benchstat` to compare old vs new benchmarks: `go test -bench=. > old.txt`, make changes, `go test -bench=. > new.txt`, `benchstat old.txt new.txt`. Shows statistical significance. In CI: Run benchmark in tests, fail if > threshold (e.g., 10% regression). Store baseline, alert on regression."

Q: "When should you use RWMutex vs Mutex?"

"RWMutex better when reads >> writes (>5:1 ratio). Benchmark both: Mutex has lower overhead for writes, RWMutex allows concurrent reads. Pattern: Benchmark with realistic read/write ratio. If 90% reads, RWMutex wins. If 50/50, Mutex often faster. Always measure, don't guess."

Key Takeaways

1. Use `b.RunParallel()` for concurrent benchmarks
2. Always `b.ResetTimer()` after setup
3. Assign results to prevent compiler optimization
4. Use `-benchmem` to measure allocations
5. Compare with `benchstat` for statistical significance
6. Benchmark different concurrency levels
7. Profile with `-cpuprofile`, `-memprofile`
8. Lock contention kills performance
9. Atomic fastest, channels slowest, mutexes middle
10. Always benchmark, don't assume

Exercises

1. Benchmark: Compare atomic, mutex, RWMutex, channel for counter.
2. Find optimal worker count: Benchmark 1, 2, 4, 8, 16, 32 workers.
3. Benchmark your cache: Test mutex, sync.Map, sharded approaches.
4. Measure lock contention: Vary critical section duration, plot results.

5. Use `benchstat`: Make optimization, prove it's better statistically.

Next: [tracing.md](#) - Using Go's execution tracer to visualize concurrency.