

Concurrency Boundaries

What are Concurrency Boundaries?

Concurrency boundary: Point where sequential code transitions to concurrent execution.

Key question: "Where should I spawn goroutines?"

Bad answer: "Everywhere!" (leads to chaos)

Good answer: "At natural system boundaries" (maintains clarity)

Identifying Natural Boundaries

Boundary 1: Network I/O

Example: HTTP Server

```
// Natural boundary: One goroutine per request
func (s *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {
    // This already runs in a goroutine (net/http does it)

    // Process request sequentially
    user := authenticate(r)
    data := fetchData(user)
    response := render(data)
    w.Write(response)
}
```

Why natural:

- Each request independent
- I/O-bound (waiting for network)
- Clear lifecycle (request → response)

Boundary 2: Independent Tasks

Example: Batch Processing

```
func processBatch(items []Item) []Result {
    // Boundary: Each item independent
    resultChan := make(chan Result, len(items))

    for _, item := range items {
        item := item
        go func() {
            resultChan <- process(item) // Concurrent
        }()
    }

    // Collect results
    results := make([]Result, 0, len(items))
    for i := 0; i < len(items); i++ {
```

```

        results = append(results, <-resultChan)
    }

    return results
}

```

Why natural:

- No dependencies between items
- CPU-bound work benefits from parallelism
- Each task self-contained

Boundary 3: Long-Running Services

Example: Background Workers

```

type Application struct {
    httpServer *http.Server
    taskQueue  *TaskQueue
    cacheSync   *CacheSync
    healthCheck *HealthCheck
}

func (a *Application) Start() {
    // Boundary: Each service independent
    go a.httpServer.ListenAndServe()
    go a.taskQueue.Run()
    go a.cacheSync.Run()
    go a.healthCheck.Run()
}

```

Why natural:

- Each service has its own lifecycle
- Run indefinitely
- Independent responsibilities

Boundary 4: Fan-Out Operations

Example: Multi-API Aggregation

```

func getProfile(userID string) (*Profile, error) {
    // Boundary: Fetch from multiple sources concurrently

    var userInfo UserInfo
    var posts []Post
    var friends []User
    var errInfo, errPosts, errFriends error

    var wg sync.WaitGroup
    wg.Add(3)

    go func() {

```

```

    defer wg.Done()
    userInfo, errInfo = fetchUserInfo(userID)
}()

go func() {
    defer wg.Done()
    posts, errPosts = fetchPosts(userID)
}()

go func() {
    defer wg.Done()
    friends, errFriends = fetchFriends(userID)
}()

wg.Wait()

if errInfo != nil {
    return nil, errInfo
}

return &Profile{
    User:    userInfo,
    Posts:   posts,
    Friends: friends,
}, nil
}

```

Why natural:

- Independent API calls
- I/O-bound (network requests)
- Results combined at end

Where NOT to Add Concurrency

Anti-Pattern 1: Fine-Grained Concurrency

```

// BAD: Goroutine for every tiny operation
func calculateTotal(items []int) int {
    results := make(chan int, len(items))

    for _, item := range items {
        go func(item int) {
            results <- item * 2 // Goroutine overhead > work!
        }(item)
    }

    total := 0
    for i := 0; i < len(items); i++ {
        total += <-results
    }
}

```

```

        return total
    }

// GOOD: Sequential for small work
func calculateTotal(items []int) int {
    total := 0
    for _, item := range items {
        total += item * 2
    }
    return total
}

```

Rule: Goroutine overhead (~2KB stack + scheduling) should be < work time.

Anti-Pattern 2: Hidden Concurrency in Libraries

```

// BAD: Library spawns goroutines internally
type Cache struct{}

func (c *Cache) Get(key string) string {
    result := make(chan string)
    go func() {
        // Fetch from cache
        result <- fetch(key)
    }()
    return <-result // Caller waits anyway!
}

// GOOD: Let caller decide concurrency
func (c *Cache) Get(key string) string {
    return fetch(key) // Sequential
}

```

Rule: Don't hide goroutines in synchronous-looking APIs.

Anti-Pattern 3: Premature Optimization

```

// BAD: Adding concurrency "just in case"
func readConfig() Config {
    configChan := make(chan Config)
    errChan := make(chan error)

    go func() {
        // Read config (happens once at startup)
        data, err := ioutil.ReadFile("config.json")
        if err != nil {
            errChan <- err
            return
        }
        var cfg Config

```

```

        json.Unmarshal(data, &cfg)
        configChan <- cfg
    }()
}

select {
case cfg := <-configChan:
    return cfg
case err := <-errChan:
    log.Fatal(err)
}
}

// GOOD: Sequential is simpler
func readConfig() Config {
    data, err := ioutil.ReadFile("config.json")
    if err != nil {
        log.Fatal(err)
    }
    var cfg Config
    json.Unmarshal(data, &cfg)
    return cfg
}

```

Rule: Add concurrency when proven necessary, not speculatively.

Boundary Design Patterns

Pattern 1: API Layer Boundary

Concurrency starts at API handlers.

```

// HTTP handler already concurrent (net/http)
func (h *Handler) ServeHTTP(w http.ResponseWriter, r *http.Request) {
    // This runs in its own goroutine

    // Business logic sequential
    result := h.service.Process(r.Context(), extractData(r))

    writeJSON(w, result)
}

```

Pattern 2: Repository Boundary

Data access points where I/O happens.

```

type UserRepository struct {
    db *sql.DB
}

// Sequential interface
func (r *UserRepository) GetByID(ctx context.Context, id string) (*User, error) {

```

```

// Let caller decide if concurrent
return queryUser(ctx, r.db, id)
}

// Caller controls concurrency
func getMultipleUsers(repo *UserRepository, ids []string) []*User {
    results := make(chan *User, len(ids))

    for _, id := range ids {
        id := id
        go func() {
            user, _ := repo.GetByID(context.Background(), id)
            results <- user
        }()
    }

    users := make([]*User, 0, len(ids))
    for i := 0; i < len(ids); i++ {
        users = append(users, <-results)
    }
    return users
}

```

Pattern 3: Service Layer Boundary

Business logic orchestrates concurrent operations.

```

type OrderService struct {
    inventory InventoryService
    payment   PaymentService
    shipping  ShippingService
    notifier  NotificationService
}

func (s *OrderService) PlaceOrder(ctx context.Context, order Order) error {
    // Sequential critical path
    if err := s.inventory.Reserve(ctx, order.Items); err != nil {
        return err
    }

    if err := s.payment.Charge(ctx, order.Payment); err != nil {
        s.inventory.Release(ctx, order.Items)
        return err
    }

    // Concurrent non-critical operations
    var wg sync.WaitGroup

    wg.Add(2)
    go func() {
        defer wg.Done()

```

```

        s.shipping.Schedule(ctx, order)
    }()

    go func() {
        defer wg.Done()
        s.notifier.SendConfirmation(ctx, order.UserID)
    }()
}

wg.Wait()

return nil
}

```

Pattern 4: Worker Pool Boundary

Controlled concurrency for CPU-bound work.

```

type ImageProcessor struct {
    workers int
    tasks   chan Task
}

func NewImageProcessor(workers int) *ImageProcessor {
    p := &ImageProcessor{
        workers: workers,
        tasks:   make(chan Task, workers*2),
    }

    // Boundary: Fixed number of workers
    for i := 0; i < workers; i++ {
        go p.worker()
    }

    return p
}

func (p *ImageProcessor) Process(img Image) Result {
    // Caller submits, worker pool handles concurrency
    result := make(chan Result)
    p.tasks <- Task{Image: img, Result: result}
    return <-result
}

func (p *ImageProcessor) worker() {
    for task := range p.tasks {
        task.Result <- processImage(task.Image)
    }
}

```

Minimizing Shared State

Technique 1: Immutable Configuration

```
// Shared read-only config (no locks needed)
type Config struct {
    Timeout time.Duration
    MaxRetry int
}

var cfg atomic.Value // Holds *Config

func init() {
    cfg.Store(&Config{Timeout: time.Second, MaxRetry: 3})
}

func getConfig() *Config {
    return cfg.Load().(*Config)
}

// Safe to use across goroutines
func handler() {
    config := getConfig()
    // Use config.Timeout, config.MaxRetry
}
```

Technique 2: Message Passing

```
// Each goroutine owns its state
type Account struct {
    balance int
    ops     chan Operation
}

type Operation struct {
    Type   OpType
    Amount int
    Result chan int
}

func (a *Account) Run() {
    for op := range a.ops {
        switch op.Type {
        case Deposit:
            a.balance += op.Amount
            op.Result <- a.balance
        case Withdraw:
            a.balance -= op.Amount
            op.Result <- a.balance
        }
    }
}
```

```
// No shared state, no locks
func (a *Account) Deposit(amount int) int {
    result := make(chan int)
    a.ops <- Operation{Type: Deposit, Amount: amount, Result: result}
    return <-result
}
```

Technique 3: Partitioning

```
// Shard data to reduce contention
type ShardedCache struct {
    shards []*CacheShard
}

type CacheShard struct {
    mu    sync.RWMutex
    data map[string]string
}

func (sc *ShardedCache) getShard(key string) *CacheShard {
    hash := fnv32(key)
    return sc.shards[hash%uint32(len(sc.shards))]
}

func (sc *ShardedCache) Get(key string) string {
    shard := sc.getShard(key)
    shard.mu.RLock()
    defer shard.mu.RUnlock()
    return shard.data[key]
}

// Each shard independent, reduces lock contention
```

Real Example: Web Crawler

```
type Crawler struct {
    // Boundary 1: Fetcher goroutines
    fetchWorkers int

    // Boundary 2: Parser goroutines
    parseWorkers int

    // Internal state (not shared)
    urlQueue chan string
    htmlQueue chan HTML
    results   chan Result

    // Shared state (protected)
```

```

    visited sync.Map
}

func NewCrawler(fetchWorkers, parseWorkers int) *Crawler {
    c := &Crawler{
        fetchWorkers: fetchWorkers,
        parseWorkers: parseWorkers,
        urlQueue:     make(chan string, 100),
        htmlQueue:    make(chan HTML, 100),
        results:      make(chan Result, 100),
    }

    // Boundary: Start workers
    for i := 0; i < fetchWorkers; i++ {
        go c.fetcher()
    }

    for i := 0; i < parseWorkers; i++ {
        go c.parser()
    }

    return c
}

func (c *Crawler) fetcher() {
    for url := range c.urlQueue {
        // Check if visited (shared state, but sync.Map handles concurrency)
        if _, seen := c.visited.LoadOrStore(url, true); seen {
            continue
        }

        // Fetch (I/O-bound, benefits from concurrency)
        html := fetch(url)
        c.htmlQueue <- html
    }
}

func (c *Crawler) parser() {
    for html := range c.htmlQueue {
        // Parse (CPU-bound, benefits from parallelism)
        result := parse(html)
        c.results <- result

        // Extract links
        for _, link := range result.Links {
            c.urlQueue <- link
        }
    }
}

func (c *Crawler) Crawl(startURL string) <-chan Result {
    c.urlQueue <- startURL
}

```

```
        return c.results
    }
```

Boundaries:

1. **API boundary:** Crawl() spawns system
2. **I/O boundary:** Fetcher goroutines (network)
3. **CPU boundary:** Parser goroutines (processing)

State management:

- urlQueue , htmlQueue , results : Private channels (no locks)
- visited : Shared sync.Map (internally synchronized)

Testing Concurrency Boundaries

```
func TestCrawler(t *testing.T) {
    // Mock HTTP responses
    server := httptest.NewServer(http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
        w.Write([]byte("<a href='/page2'>Link</a>"))
    }))
    defer server.Close()

    crawler := NewCrawler(2, 2)
    results := crawler.Crawl(server.URL)

    // Collect results
    timeout := time.After(time.Second)
    found := 0

    for {
        select {
        case <-results:
            found++
        case <-timeout:
            if found < 2 {
                t.Errorf("Expected at least 2 results, got %d", found)
            }
            return
        }
    }
}
```

Boundary Design Checklist

- Concurrency at natural I/O or CPU boundaries
- Each boundary has clear ownership
- Shared state minimized or eliminated
- Goroutine count controlled (not unbounded)

- Each goroutine has exit condition
- Boundaries documented in code
- Testing covers concurrent execution
- Errors propagated across boundaries
- Graceful shutdown at each boundary
- Monitoring/observability at boundaries

Interview Questions

Q: "Where should concurrency start in a web service?"

"Primary boundary: HTTP handler (net/http already spawns goroutine per request). Secondary boundaries: 1) Fan-out for parallel API calls (fetch user, posts, friends concurrently). 2) Worker pools for CPU-intensive tasks (image processing, data transformation). 3) Background services (cache sync, health checks). Don't add concurrency in: Business logic (keep sequential), data access layer (caller controls), utility functions. Rule: Concurrency at I/O and CPU boundaries, not everywhere."

Q: "How do you minimize shared state?"

"1) Ownership: Each goroutine owns its data. 2) Message passing: Communicate via channels, not shared memory. 3) Immutable data: Read-only config safe to share. 4) Atomic operations: Simple counters/flags. 5) Partitioning: Shard data to reduce contention (16 shards = 1/16 lock contention). Example: Instead of shared cache with mutex, use actor pattern where cache goroutine owns state, others send requests via channel."

Q: "When is concurrency premature?"

"When: 1) Operation runs once (startup config loading). 2) Work smaller than goroutine overhead (~ $1\mu s$). 3) No I/O or CPU parallelism (sequential dependencies). 4) Not measured as bottleneck. Example: Don't spawn goroutine to read 10-line config file. Do spawn for: HTTP endpoints, parallel API calls, batch processing, long-running services. Measure first, optimize second."

Q: "How do you control goroutine proliferation?"

"1) Worker pools: Fixed number of goroutines, tasks queued. 2) Semaphores: Limit concurrent operations (buffered channel of size N). 3) Rate limiters: Control creation rate (golang.org/x/time/rate). 4) Context: Cancel goroutines when work unnecessary. 5) WaitGroups: Track and wait for completion. Example: Instead of `go fetch(url)` for 10000 URLs (10000 goroutines), use worker pool with `runtime.NumCPU()*10` workers."

Key Takeaways

1. Place boundaries at natural I/O or CPU points
2. Don't hide goroutines in synchronous APIs
3. Minimize shared state across boundaries
4. Control goroutine count (worker pools, semaphores)
5. Each boundary needs clear ownership
6. Caller should control concurrency, not library
7. Add concurrency for I/O latency or CPU parallelism
8. Sequential is simpler, use when work is fast
9. Document boundaries in code
10. Test with realistic concurrency levels

Exercises

1. Design boundaries for e-commerce checkout flow.
2. Identify unnecessary concurrency in codebas and remove it.
3. Convert unbounded goroutine creation to worker pool.
4. Build system with 3 clear concurrency boundaries, document each.
5. Refactor shared state to message-passing architecture.

Next: [choosing-primitives.md](#) - Decision tree for selecting concurrency primitives.