

Job Queue Project

Learning Objectives

Build a production-ready job queue to learn:

- Worker pool pattern with priority queues
- Graceful shutdown with in-flight tracking
- Job persistence and at-least-once delivery
- Retry logic with exponential backoff
- Preventing worker starvation

Requirements

Functional Requirements

1. Worker Pool

- Fixed number of workers processing jobs
- Bounded job queue (backpressure when full)
- Concurrent job execution

2. Priority Levels

- High, Medium, Low priority jobs
- High priority processed first
- No starvation (low priority eventually runs)

3. Job Lifecycle

- Enqueue (blocking if queue full)
- Processing (worker picks job)
- Completion (success or failure)
- Retry (on failure with backoff)

4. Persistence (Final only)

- Jobs persisted to disk
- Survive process restart
- At-least-once delivery guarantee

Non-Functional Requirements

1. Performance

- Process 10,000+ jobs/sec
- Low latency (<10ms queue → worker)
- Efficient memory usage

2. Reliability

- No job loss on crash (with persistence)
- Retry failed jobs (configurable attempts)
- Graceful shutdown (wait for in-flight)

3. Observability

- Metrics (queued, processing, completed, failed)
- Job status tracking
- Worker utilization

Three Implementations

1. Naive Implementation (naive/job_queue.go)

Approach:

- Single unbuffered channel for all jobs
- No priorities (FIFO only)
- No persistence (in-memory only)
- No retry logic

Problems (Intentional):

- **No backpressure:** Unbounded queue can OOM
- **No priorities:** Critical jobs wait behind low-priority
- **No persistence:** Jobs lost on crash
- **Poor shutdown:** Workers may exit with jobs in-flight
- **No retry:** Failed jobs are lost

Expected Issues:

- Memory grows unbounded with fast producers
- Can't prioritize urgent work
- Not production-ready (data loss)

2. Improved Implementation (improved/job_queue.go)

Fixes:

- **Bounded queues:** Separate channel per priority
- **Priority selection:** High > Medium > Low with select
- **Retry logic:** Exponential backoff (1s, 2s, 4s, 8s)
- **Graceful shutdown:** WaitGroup tracks in-flight jobs
- **Metrics:** Atomic counters for observability

Improvements:

- Backpressure prevents OOM
- Priorities work correctly
- Failed jobs retry automatically
- Clean shutdown guaranteed

Remaining Issues:

- Still no persistence (in-memory only)
- Priority starvation possible (high priority floods)
- No way to query job status
- Retry state lost on restart

3. Final Implementation (final/job_queue.go)

Production-Ready:

- **Persistence:** Jobs persisted to disk (at-least-once)
- **Anti-starvation:** Token-based fairness (1 low per 10 high)
- **Job tracking:** Query status by ID
- **Comprehensive metrics:** Per-priority counters
- **Configurable:** Workers, retries, timeouts

Key Features:

```
// Job with metadata
type Job struct {
    ID      string
    Priority Priority
    Payload []byte
    Retries int
    Created time.Time
}

// Enqueue with backpressure
err := jq.Enqueue(ctx, job)

// Query status
status := jq.Status(jobID) // Queued, Processing, Completed, Failed

// Metrics
m := jq.Metrics()
fmt.Printf("Completed: %d, Failed: %d\n", m.Completed, m.Failed)
```

Performance:

- Throughput: ~50k jobs/sec (4 workers)
- Latency: p50=1ms, p99=10ms
- Persistence overhead: ~10% (batched writes)

Success Criteria

1. Understand worker pool:

- How many workers for CPU-bound vs I/O-bound jobs?
- What happens if workers > CPU cores?
- How to detect worker starvation?

2. Master priority queues:

- Why separate channels per priority?
- How to prevent starvation?
- Trade-off: fairness vs throughput

3. Implement persistence:

- At-least-once vs at-most-once vs exactly-once
- Where to persist (before processing? after?)
- How to recover on restart?

4. Handle failures:

- Exponential backoff vs fixed delay
- Max retries before dead letter queue
- Idempotent job handlers

Usage Examples

Naive Implementation

```
jq := naive.NewJobQueue(4) // 4 workers
defer jq.Close()

// Enqueue job (blocks if internal buffer full)
jq.Enqueue(naive.Job{
    ID:      "job1",
    Handler: func() error { return processJob() },
})
```

Final Implementation

```
jq := final.NewJobQueue(final.Config{
    NumWorkers:   4,
    QueueSize:   1000, // per priority
    MaxRetries:  3,
    PersistPath: "./jobs.db",
    ShutdownTimeout: 30 * time.Second,
})
defer jq.Close()

// Enqueue with priority and context
ctx, cancel := context.WithTimeout(context.Background(), 5*time.Second)
defer cancel()

job := final.Job{
    ID:      uuid.New().String(),
    Priority: final.HighPriority,
    Payload: []byte(`{"action": "send_email"}`),
}

if err := jq.Enqueue(ctx, job); err != nil {
    log.Fatal(err)
}

// Query status
status := jq.Status(job.ID)
fmt.Printf("Job %s: %s\n", job.ID, status)

// Wait for completion
ticker := time.NewTicker(100 * time.Millisecond)
for range ticker.C {
    if status := jq.Status(job.ID); status == final.Completed {
```

```
        break
    }
}
```

Testing Strategy

Unit Tests

```
cd final
go test -v
go test -race
go test -cover
```

Test Cases

1. **TestEnqueue** - Verify job queued successfully
2. **TestPriority** - High priority processed first
3. **TestRetry** - Failed jobs retry with backoff
4. **TestGracefulShutdown** - In-flight jobs complete
5. **TestPersistence** - Jobs survive restart
6. **TestBackpressure** - Enqueue blocks when full
7. **TestMetrics** - Counters accurate

Stress Test

```
func TestStress(t *testing.T) {
    jq := NewJobQueue(Config{NumWorkers: 10, QueueSize: 1000})
    defer jq.Close()

    var wg sync.WaitGroup
    numJobs := 100000

    for i := 0; i < numJobs; i++ {
        wg.Add(1)
        go func(id int) {
            defer wg.Done()
            job := Job{
                ID: fmt.Sprintf("job%d", id),
                Priority: Priority(id % 3),
                Payload: []byte("data"),
            }
            jq.Enqueue(context.Background(), job)
        }(i)
    }

    wg.Wait()

    // Wait for completion
    time.Sleep(5 * time.Second)
```

```

m := jq.Metrics()
if m.Completed != uint64(numJobs) {
    t.Errorf("Expected %d completed, got %d", numJobs, m.Completed)
}
}

```

Real-World Scenarios

1. Background Email Sender

```

jq := final.NewJobQueue(final.Config{
    NumWorkers: 10, // I/O-bound, more workers
    QueueSize: 10000,
    MaxRetries: 3,
})

// Enqueue email job
job := final.Job{
    ID:      emailID,
    Priority: final.MediumPriority,
    Payload: marshalEmail(to, subject, body),
}

jq.EnqueueWithHandler(context.Background(), job, func(payload []byte) error {
    return sendEmail(payload)
})

```

2. Image Processing Pipeline

```

jq := final.NewJobQueue(final.Config{
    NumWorkers: runtime.NumCPU(), // CPU-bound
    QueueSize: 1000,
})

// High priority for user-uploaded images
job := final.Job{
    ID:      imageID,
    Priority: final.HighPriority,
    Payload: imageData,
}

jq.EnqueueWithHandler(ctx, job, func(payload []byte) error {
    return processImage(payload) // Resize, compress, etc.
})

```

3. Webhook Delivery System

```

jq := final.NewJobQueue(final.Config{
    NumWorkers: 20, // I/O-bound (network requests)
    QueueSize: 50000,
    MaxRetries: 5, // Webhooks can fail temporarily
    PersistPath: "/var/lib/webhooks.db",
})

// Critical webhooks get high priority
priority := final.MediumPriority
if webhook.IsCritical {
    priority = final.HighPriority
}

job := final.Job{
    ID:        webhookID,
    Priority: priority,
    Payload:  marshalWebhook(webhook),
}

jq.EnqueueWithHandler(ctx, job, func(payload []byte) error {
    return deliverWebhook(payload)
})

```

Common Pitfalls

Pitfall 1: Too Many Workers (CPU-Bound)

```

// ✗ WRONG: 100 workers for CPU-bound work
jq := NewJobQueue(Config{NumWorkers: 100})

// ✅ CORRECT: Match CPU cores
jq := NewJobQueue(Config{NumWorkers: runtime.NumCPU()})

```

Pitfall 2: Priority Starvation

```

// ✗ WRONG: Always pick high priority
select {
case job := <-highQueue:
    process(job)
case job := <-lowQueue: // Never selected if high queue busy
    process(job)
}

// ✅ CORRECT: Token-based fairness
tokens := 10
for {
    if tokens > 0 {
        select {
            case job := <-highQueue:

```

```

        process(job)
        tokens--
    default:
        // High queue empty, reset tokens
        tokens = 10
    }
} else {
    // Process 1 low priority, then reset
    select {
        case job := <-lowQueue:
            process(job)
            tokens = 10
        default:
            tokens = 10
    }
}
}

```

Pitfall 3: Forgetting WaitGroup.Add

```

// ❌ WRONG: Add inside goroutine (race)
go func() {
    wg.Add(1) // Main might call Wait() before this
    defer wg.Done()
    work()
}()

// ✅ CORRECT: Add before spawning
wg.Add(1)
go func() {
    defer wg.Done()
    work()
}()

```

Pitfall 4: Not Handling Context Cancellation

```

// ❌ WRONG: Enqueue ignores context
func (jq *JobQueue) Enqueue(ctx context.Context, job Job) {
    jq.queue <- job // Blocks forever if queue full
}

// ✅ CORRECT: Respect context
func (jq *JobQueue) Enqueue(ctx context.Context, job Job) error {
    select {
        case jq.queue <- job:
            return nil
        case <-ctx.Done():
            return ctx.Err()
    }
}

```

```
    }  
}
```

Pitfall 5: Losing Jobs on Shutdown

```
// ❌ WRONG: Close queue immediately  
close(jq.queue)  
  
// ✅ CORRECT: Wait for in-flight jobs  
jq.shutdown() // Signal workers to stop  
jq.wg.Wait() // Wait for in-flight jobs  
close(jq.queue)
```

Interview Discussion Points

1. How many workers for I/O-bound vs CPU-bound?

- CPU-bound: NumCPU (more causes contention)
- I/O-bound: $\text{NumCPU} \times 10-100$ (waiting for I/O)
- Benchmark to find optimal

2. How to prevent priority starvation?

- Token-based: Process N high, then 1 low
- Weighted: Pick high 90%, low 10%
- Age-based: Increase priority over time

3. At-least-once vs at-most-once vs exactly-once?

- At-least-once: Persist before processing (may process twice on crash)
- At-most-once: Process then persist (may lose on crash)
- Exactly-once: Distributed transaction (complex, slow)
- Most systems use at-least-once + idempotent handlers

4. How to recover jobs after crash?

- Persist job state to disk/database
- On startup, read all Queued/Processing jobs
- Re-enqueue for processing
- Use WriteAheadLog for durability

5. Why exponential backoff?

- Avoids thundering herd (all retries at once)
- Gives failing service time to recover
- Formula: $\text{delay} = \min(\text{maxDelay}, \text{baseDelay} \times 2^{\text{attempt}})$

Next Steps

After mastering job-queue:

1. **cache/** - Concurrent LRU cache with sharding
2. **web-crawler/** - Bounded concurrency and politeness
3. **connection-pool/** - Database pool with circuit breaker

4. **pub-sub/** - Topic-based messaging with fan-out

Build production-ready concurrent systems! 