

# Netflix System Design

## Video Streaming Platform at Global Scale

(New Grad / Junior Engineer Interview Focus | TypeScript/Node.js + Go + PostgreSQL)

## 0 Executive Summary

### System Purpose

Netflix is a **video streaming platform** that delivers TV shows, movies, and original content to 260 million subscribers worldwide. Users can watch unlimited video on-demand, pause and resume across devices, download content for offline viewing, and receive personalized recommendations. As of 2026, Netflix serves 260M paid subscribers across 190 countries, streaming 1 billion hours of video per day.

### Core Constraints

- **Massive bandwidth:** 1B hours/day = 2 exabytes/month bandwidth (far exceeds Spotify)
- **Video quality:** 4K HDR requires 25 Mbps (vs 320 kbps for audio)
- **Adaptive streaming:** Switch quality based on network (4K → 1080p → 720p → 480p)
- **Start latency:** Video must start within 2 seconds (or users abandon)
- **Global scale:** Same content available in 190 countries (licensing per region)
- **Offline downloads:** Download entire movies (5 GB per movie) for airplane mode
- **Personalization:** Homepage shows different content for each user (not generic catalog)

### Key Architectural Choices

1. **Open Connect CDN** (Netflix's own CDN, 95% of traffic served from ISP caches)
2. **Multiple video encodings** (4K, 1080p, 720p, 480p, 360p, 240p)
3. **Adaptive Bitrate Streaming** (MPEG-DASH, switch quality every 2-10 seconds)
4. **AWS for control plane** (metadata, recommendations, billing)
5. **PostgreSQL for catalog** (movies, shows, episodes metadata)
6. **Cassandra for viewing history** (billions of playback events)
7. **Go for video encoding** (transcode original to multiple formats, CPU-intensive)
8. **Machine learning for recommendations** (personalized homepage for each user)

### Biggest Challenges

- **Bandwidth costs:** 2 exabytes/month = \$170M/year just for CDN (mitigated by Open Connect)
- **Video encoding:** 1 movie → 200+ versions (6 qualities × 20 devices × 2 languages)
- **Start latency:** Must buffer first 10 seconds quickly (2-second target)
- **4K quality:** 25 Mbps × 1M concurrent 4K streams = 25 Tbps aggregate bandwidth
- **Personalization at scale:** Generate unique homepage for 260M users daily
- **Licensing complexity:** "Stranger Things" available in US, blocked in China

### Interview Focus (New Grad Level)

- Video streaming fundamentals (ABR, chunked delivery, HLS/DASH)
- CDN architecture (why Netflix built Open Connect)
- Video encoding pipeline (transcode, segment, package)
- Catalog data modeling (movies, shows, episodes, seasons)
- Viewing history tracking (resume playback across devices)
- Basic recommendation concepts (collaborative filtering, content-based)
- Offline download management (DRM, storage limits)

## 1 Problem Definition

### What Problem Does Netflix Solve?

Netflix replaces **cable TV and DVD rentals** with **unlimited on-demand streaming**:

1. **Watch anytime:** No fixed schedule (not like TV broadcasts)
2. **Unlimited content:** 10,000+ titles for fixed monthly fee (\$15/month)
3. **Cross-device:** Start on TV, continue on phone
4. **No ads:** Uninterrupted viewing (vs cable TV with 15 min/hour of ads)
5. **Offline mode:** Download movies for flights, commutes
6. **Personalization:** Homepage shows content you'll like (not generic catalog)

### Who Are the Users?

- **260 million paid subscribers** (as of 2026)
- **1 billion hours streamed per day** (avg 4 hours/subscriber/day)
- **10,000+ titles** in catalog (movies + TV shows)
- **190 countries** served (except China, North Korea, Syria, Crimea)
- **Use cases:**
  - Binge-watching: Watch entire season in one sitting (8 hours)
  - Family viewing: Kids watching cartoons, parents watching dramas
  - Commuters: Downloaded episodes on phone/tablet
  - Smart TV: 4K HDR viewing on large screens
  - Background viewing: Laptop while working

### Why Existing Solutions Failed

#### ✗ Cable TV

Alice wants to watch "Breaking Bad"  
→ Season 1 Episode 1 airs Mondays at 9 PM  
→ Alice misses it (was working late)  
→ Must wait for rerun or buy DVD  
→ \$80/month cable bill + 15 min/hour ads

#### Why it fails:

- **Fixed schedule:** Can't watch on demand
- **Expensive:** \$80-100/month for hundreds of channels you don't watch
- **Ads:** 15-20 minutes of ads per hour
- **No portability:** Can't watch on phone/laptop

#### ✗ DVD Rentals (Blockbuster)

Bob wants to watch "Inception"  
→ Drives to Blockbuster (15 min)  
→ "Inception" DVD is out of stock  
→ Rents different movie (compromise)  
→ Must return DVD in 3 days or pay late fee (\$5/day)  
→ Forgets to return → \$30 late fee

### Why it fails:

- **Limited stock:** Popular movies always rented out
  - **Physical media:** Must drive to store, return DVD
  - **Late fees:** Forgot to return? \$5-10/day penalty
  - **One movie at a time:** Can't binge-watch TV series
- 

### ✗ Piracy (BitTorrent)

Carol wants to watch latest "Stranger Things" season

- Searches torrent sites
- Downloads 10 GB file (takes 4 hours on slow connection)
- Video quality poor (compressed, artifacts)
- Risk of malware
- Illegal (copyright violation)

### Why it fails:

- **Illegal:** Risk of lawsuits, ISP warnings
  - **Slow:** Must download entire file before watching
  - **Poor quality:** Compressed, low bitrate
  - **Malware:** Torrent sites full of viruses
  - **No recommendations:** Must know what to search for
- 

### ✗ YouTube (Free Ad-Supported)

Dave wants to watch a movie on YouTube

- Searches "Inception full movie"
- Finds pirated upload (poor quality, divided into 10 parts)
- Part 3 deleted for copyright violation
- Ad every 5 minutes (disruptive)

### Why it fails:

- **Pirated content:** Low quality, often deleted
  - **Split videos:** Movie divided into parts (annoying)
  - **Frequent ads:** Free tier has ads every few minutes
  - **No offline:** Must have internet connection
- 

## Constraints That Shape the System

1. **Bandwidth-heavy:** 1B hours/day × 5 Mbps avg = 5 exabits/day (500× more than Spotify!)
  2. **Low latency:** Video must start <2s (or users abandon, 8% abandon per second of delay)
  3. **High quality:** 4K HDR requires 25 Mbps (vs 320 kbps for Spotify audio)
  4. **Adaptive quality:** Must switch quality based on network (4K → 1080p → 480p seamlessly)
  5. **Large downloads:** Offline movie = 5 GB (vs 10 MB for Spotify song)
  6. **Storage costs:** 10,000 titles × 200 versions × 100 GB/version = 200 petabytes
  7. **Encoding costs:** Transcode 1 movie to 200 versions = \$500-1000 compute cost
- 

## 2 Requirements

## **Functional Requirements**

### **Core Features**

#### **1. Stream video**

- Play any title from 10,000+ catalog
- Pause, resume, seek within video
- Adaptive bitrate (switch quality based on network)
- Subtitles/closed captions (multiple languages)
- Audio tracks (English, Spanish, Japanese, etc.)
- Skip intro/credits (convenience feature)

#### **2. Browse catalog**

- Homepage (personalized recommendations)
- Browse by genre (Action, Comedy, Drama, Documentary)
- Search by title, actor, director
- Trending now (popular content)
- Continue watching (resume from last position)
- My List (saved titles to watch later)

#### **3. User profiles**

- Multiple profiles per account (up to 5)
- Kids profile (filtered content, no mature titles)
- Avatar customization
- Viewing history per profile
- Separate recommendations per profile

#### **4. Playback controls**

- Play, pause, stop
- Seek forward/backward (10-second skip)
- Volume control
- Fullscreen mode
- Picture-in-picture (watch while browsing)
- Playback speed (0.5x, 1x, 1.25x, 1.5x, 2x)

#### **5. Offline downloads (Mobile only)**

- Download movie/episode for offline viewing
- Auto-download next episode
- Manage storage (delete downloads)
- DRM (downloads expire if subscription ends)
- Quality selection (Standard, High)

#### **6. Cross-device sync**

- Pause on TV, resume on phone at same timestamp
- Continue watching list synced across devices
- Playback progress saved

#### **7. Parental controls**

- Set PIN for mature content
- Content rating filters (G, PG, PG-13, R, TV-MA)

- Kids profile (only kids content visible)

## **Secondary Features**

### **8. Social features**

- Share title via link
- Watch together (synchronized viewing with friends, video call)

### **9. Notifications**

- New episode available (for shows you're watching)
- New season added
- Content expiring soon (leaving catalog)

### **10. Accessibility**

- Audio descriptions (for visually impaired)
- Closed captions/subtitles
- High contrast UI mode

## **Non-Functional Requirements**

### **Scale (2026 Numbers)**

Metric	Value	Derivation
<b>Paid Subscribers</b>	260 million	Netflix reported data
<b>Viewing Hours per Day</b>	1 billion	$260M \times 4 \text{ hours/day avg}$
<b>Titles in Catalog</b>	10,000+	Movies + TV shows + originals
<b>Concurrent Streams (Peak)</b>	50 million	$260M \times 20\% \text{ watching at peak hour}$
<b>Average Video Duration</b>	45 minutes	Mix of TV episodes (30 min) and movies (2 hr)
<b>Average Bitrate</b>	5 Mbps	Mix of 4K (25 Mbps) and 1080p (5 Mbps)
<b>Daily Bandwidth</b>	5 exabits	$1B \text{ hours} \times 5 \text{ Mbps} \times 3600 \text{ sec/hour}$
<b>Monthly Bandwidth</b>	150 exabits	$5 \text{ exabits} \times 30 \text{ days}$
<b>Monthly Bandwidth (bytes)</b>	2 exabytes	$150 \text{ exabits} \div 8 \text{ bits/byte}$
<b>CDN Cost (estimate)</b>	\$170M/year	$2 \text{ EB/month} \times \$0.085/\text{GB}$ (reduced by Open Connect)
<b>Storage (all encodings)</b>	200 petabytes	$10K \text{ titles} \times 200 \text{ versions} \times 100 \text{ GB}$

### **Video Quality Tiers**

Quality	Resolution	Bitrate	File Size (2hr movie)	Use Case
<b>Ultra Low</b>	240p	0.3 Mbps	270 MB	Slow 2G, data saving
<b>Low</b>	360p	0.7 Mbps	630 MB	3G networks
<b>Standard</b>	480p (SD)	1.5 Mbps	1.35 GB	Basic subscription

<b>Medium</b>	720p (HD)	3 Mbps	2.7 GB	Standard subscription
<b>High</b>	1080p (Full HD)	5 Mbps	4.5 GB	Standard/Premium
<b>Ultra High</b>	4K (UHD)	25 Mbps	22.5 GB	Premium, good network
<b>4K HDR</b>	4K HDR	25 Mbps	22.5 GB	Premium, HDR-capable device

#### Subscription Tiers:

- **Basic (\$9.99/month):** 480p (SD), 1 screen, no downloads
- **Standard (\$15.49/month):** 1080p (HD), 2 screens, downloads
- **Premium (\$19.99/month):** 4K (UHD), 4 screens, downloads

#### Latency Targets

Operation	P50	P99	Max Acceptable	Rationale
<b>Video Start (Cold)</b>	1s	3s	5s	Critical UX (users abandon)
<b>Video Start (Resumed)</b>	500ms	1.5s	3s	Already buffered data
<b>Quality Switch</b>	2s	5s	10s	Seamless transition
<b>Seek Within Video</b>	500ms	2s	5s	Scrubbing timeline
<b>Homepage Load</b>	500ms	1.5s	3s	First impression
<b>Search Results</b>	200ms	1s	2s	Must feel instant
<b>Episode Auto-Play</b>	1s	3s	5s	Between episodes

**Video Start Latency Budget (P99):** \*\*\* DNS lookup: 50ms (cached) TCP handshake: 50ms (Open Connect edge) HTTP request/response headers: 50ms First video segment download: 1000ms (2 MB @ 2 MB/s 4G) Video decode + buffer: 200ms (H.264/HEVC decode)

Total: 1350ms (within 2s target)

#### #### Availability & Durability

Requirement	Target	Implementation
**Streaming Availability**	99.95%	Multi-region Open Connect, auto-failover
**Catalog Durability**	99.99999999%	S3 (11 nines), replicated
**User Data Durability**	99.99%	Cassandra (3 replicas)
**Control Plane Availability**	99.99%	Multi-region AWS, active-active
**Max Downtime**	22 minutes/month	99.95% uptime

#### #### Bandwidth Optimization

\*\*Problem\*\*: Streaming 1B hours/day = 2 exabytes/month = \$170M/year bandwidth cost

```

**Optimizations**:
1. **Open Connect CDN**: 95% of traffic served from ISP caches (free peering)
2. **Adaptive bitrate**: Drop from 4K to 1080p = 80% bandwidth savings
3. **Video compression**: H.265/HEVC (50% smaller than H.264 at same quality)
4. **Preloading**: Predict what user will watch, pre-push to ISP cache (off-peak hours)
5. **Per-title encoding**: Optimize encoding per title (animation vs live-action)

```

---

## ## 3 High-Level Architecture

### ### Component Overview

```

```mermaid
graph TD
    TV[Smart TV Apps<br/>LG, Samsung, Roku]
    Mobile[Mobile Apps<br/>iOS/Android]
    Web[Web Browser<br/>HTML5 Player]
    end

    subgraph "Open Connect CDN"
        ISPCache[ISP Caches<br/>Inside ISPs<br/>95% of traffic]
        OCPCache[Open Connect Appliances<br/>Netflix-owned<br/>Regional hubs]
    end

    subgraph "AWS Control Plane"
        LB[Load Balancer<br/>ALB]
    end

    subgraph "Application Services"
        API[API Gateway<br/>Node.js<br/>500 instances]
        PlaybackAPI[Playback API<br/>Node.js<br/>Handle streaming]
        SearchAPI[Search API<br/>Node.js<br/>Elasticsearch]
        RecsAPI[Recommendations API<br/>Python<br/>ML models]
    end

    subgraph "Data Layer (AWS)"
        PG[(PostgreSQL<br/>Catalog Metadata<br/>Movies, Shows)]
        Cassandra[(Cassandra<br/>Viewing History<br/>Billions of events)]
        Redis[(Redis<br/>Session State<br/>Continue Watching)]
        S3Master[(S3 Master<br/>Original Videos<br/>10 PB)]
    end

    subgraph "Encoding Pipeline (AWS)"
        EncodingWorkers[Encoding Workers<br/>Go + ffmpeg<br/>1000 instances]
        EncodedStorage[(S3 Encoded<br/>All Qualities<br/>200 PB)]
    end

    subgraph "Machine Learning"

```

```

MLPipeline[Recommendation Engine<br/>Python/Spark<br/>Collaborative filtering]
end

TV -->|HTTPS| ISPCache
Mobile -->|HTTPS| ISPCache
Web -->|HTTPS| ISPCache

ISPCache -.->|Cache miss| OCPCache
OCPCache -.->|Cache miss| EncodedStorage

TV -->|Metadata requests| LB
Mobile -->|Metadata requests| LB
Web -->|Metadata requests| LB

LB --> API
LB --> PlaybackAPI
LB --> SearchAPI
LB --> RecsAPI

API --> PG
API --> Redis
PlaybackAPI --> Cassandra
SearchAPI --> Elasticsearch[(Elasticsearch)]

RecsAPI --> Cassandra
RecsAPI --> MLPipeline

S3Master --> EncodingWorkers
EncodingWorkers --> EncodedStorage
EncodedStorage --> OCPCache
OCPCache --> ISPCache

```

## Component Responsibilities

### 1. Client Applications

**Technology:** React Native (Mobile), C++ (Smart TVs), React (Web)

#### Responsibilities:

- Display UI (browse catalog, search, playback controls)
- Video playback (decode H.264/HEVC, render to screen)
- Adaptive bitrate logic (measure bandwidth, switch quality)
- Offline downloads (encrypted storage)
- Playback tracking (report to server for "continue watching")

#### Video Player Implementation:

```

interface VideoPlayer {
  currentVideo: Video;
  currentQuality: Quality; // 4K, 1080p, 720p, 480p, etc.
  currentPosition: number; // Seconds
  bufferHealth: number; // Seconds of video buffered ahead
}

```

```

        bandwidth: number;           // Measured network speed (Mbps)
    }

async function playVideo(videoId: string) {
    // 1. Fetch manifest (available qualities + segment URLs)
    const manifest = await api.get(`/playback/${videoId}/manifest`);

    // 2. Select initial quality based on network speed
    const quality = selectInitialQuality(manifest.qualities, bandwidth);

    // 3. Download first segment (2-10 seconds of video)
    const segment = await cdn.get(manifest.qualities[quality].segments[0].url);

    // 4. Decode and start playback
    videoElement.src = URL.createObjectURL(segment);
    videoElement.play();

    // 5. Prefetch next segments
    prefetchSegments(manifest.qualities[quality].segments.slice(1, 5));

    // 6. Monitor bandwidth and adjust quality
    setInterval(() => adaptBitrate(), 5000); // Every 5 seconds

    // 7. Report playback position (for "continue watching")
    setInterval(() => reportPlayback(), 30000); // Every 30 seconds
}

function adaptBitrate() {
    // Measure download speed from recent segments
    const measuredBandwidth = calculateBandwidth();

    // If bandwidth dropped, switch to lower quality
    if (measuredBandwidth < currentBitrate * 1.5) {
        switchQuality(lowerQuality);
    }

    // If bandwidth increased and buffer is healthy, switch to higher quality
    if (measuredBandwidth > currentBitrate * 2.5 && bufferHealth > 20) {
        switchQuality(higherQuality);
    }
}

```

## 2. Open Connect CDN

**Purpose:** Netflix's custom CDN, deployed inside ISPs worldwide

### Why Open Connect?

- **ISP partnership:** Netflix places servers (Open Connect Appliances) inside ISPs
- **Free peering:** ISPs don't charge for traffic (mutual benefit: reduces their backbone load)
- **95% cache hit:** Most popular content served from ISP cache (millisecond latency)
- **Cost savings:** Avoid paying CloudFront/Akamai (\$0.08/GB = \$160M/year)

## Architecture:

```
Tier 1: ISP Caches (inside Comcast, Verizon, AT&T data centers)
    → Serve 95% of traffic
    → 1-5ms latency (same city as user)

Tier 2: Open Connect Appliances (regional Netflix-owned data centers)
    → Serve 4% of traffic (cache miss from Tier 1)
    → 10-50ms latency (same country)

Tier 3: AWS S3 (master encoded storage)
    → Serve 1% of traffic (very rare titles)
    → 50-300ms latency (US/EU origins)
```

## Cache Fill Strategy:

```
// During off-peak hours (2 AM - 6 AM), pre-push content to ISP caches
func fillISPCache(ispID string, predictions []Video) {
    // 1. Predict what users in this ISP will watch today
    predictedVideos := mlModel.PredictPopularContent(ispID, date)

    // 2. Check what's already cached
    cached := ispCache.List(ispID)

    // 3. Push missing content
    for _, video := range predictedVideos {
        if !cached.Contains(video.ID) {
            // Transfer from AWS S3 to ISP cache (off-peak, low-priority)
            transferToCache(video, ispID, priority="low")
        }
    }
}
```

---

## 3. API Gateway (Node.js)

**Technology:** Node.js (Express/Fastify)

### Responsibilities:

- Handle REST API requests (GET /browse, GET /search, etc.)
- Authenticate users (JWT validation)
- Fetch metadata from PostgreSQL, Redis
- Return JSON responses

### Example API Endpoint:

```
app.get('/api/browse/home', async (req, res) => {
    const userId = req.user.id;
    const profileId = req.query.profile_id;

    // 1. Fetch personalized rows from recommendation service
    const rows = await recsAPI.get(`/recommendations/${profileId}`);
```

```

// 2. Fetch metadata for titles in each row
const enriched = await Promise.all(
  rows.map(async (row) => {
    const titles = await db.query(
      'SELECT * FROM titles WHERE id = ANY($1)' ,
      [row.title_ids]
    );
    return { ...row, titles };
  })
);

res.json({ rows: enriched });
);

```

#### 4. Playback API (Node.js)

**Technology:** Node.js (specialized for video playback)

**Responsibilities:**

- Generate video manifest (segment URLs, available qualities)
- Check licensing (is title available in user's country?)
- Track playback progress (for "continue watching")
- Handle DRM license requests

**Manifest Generation:**

```

app.get('/playback/:videoId/manifest', async (req, res) => {
  const videoId = req.params.videoId;
  const userCountry = req.user.country_code;

  // 1. Check licensing
  const isLicensed = await checkLicense(videoId, userCountry);
  if (!isLicensed) {
    return res.status(451).json({ error: 'Not available in your region' });
  }

  // 2. Fetch available qualities from database
  const qualities = await db.query(
    'SELECT * FROM video_encodings WHERE video_id = $1 ORDER BY bitrate DESC',
    [videoId]
  );

  // 3. Generate manifest (MPEG-DASH format)
  const manifest = {
    video_id: videoId,
    duration: qualities[0].duration,
    qualities: qualities.map(q => ({
      bitrate: q.bitrate,
      resolution: q.resolution,
      codec: q.codec, // 'h264' or 'hevc'
    })
  );
}

```

```

    segments: generateSegmentURLs(q.cdn_path, q.num_segments)
  })
};

res.json(manifest);
});

function generateSegmentURLs(cdnPath: string, numSegments: number): Segment[] {
  const segments = [];
  for (let i = 0; i < numSegments; i++) {
    segments.push({
      index: i,
      url: `https://openconnect.netflix.com/${cdnPath}/segment_${i}.m4s`,
      start_time: i * 4, // 4-second segments
      end_time: (i + 1) * 4
    });
  }
  return segments;
}

```

## 5. PostgreSQL (Catalog Metadata)

**Purpose:** Store structured data (movies, shows, episodes, actors)

**Schema Overview:**

```

-- Titles (movies + TV shows)
CREATE TABLE titles (
  id BIGSERIAL PRIMARY KEY,
  title VARCHAR(255) NOT NULL,
  type VARCHAR(20) NOT NULL, -- 'movie', 'series'
  synopsis TEXT,
  release_year INT,
  duration_minutes INT, -- For movies
  maturity_rating VARCHAR(10), -- 'G', 'PG', 'PG-13', 'R', 'TV-MA'
  genres TEXT[], -- ['Action', 'Thriller', 'Sci-Fi']
  poster_url TEXT,
  trailer_url TEXT,
  created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_titles_type ON titles(type);
CREATE INDEX idx_titles_genres ON titles USING GIN(genres);

-- TV Show Seasons
CREATE TABLE seasons (
  id BIGSERIAL PRIMARY KEY,
  title_id BIGINT NOT NULL REFERENCES titles(id),
  season_number INT NOT NULL,
  episode_count INT,
  created_at TIMESTAMPTZ DEFAULT NOW()
);

```

```

CREATE INDEX idx_seasons_title ON seasons(title_id, season_number);

-- TV Show Episodes
CREATE TABLE episodes (
    id BIGSERIAL PRIMARY KEY,
    season_id BIGINT NOT NULL REFERENCES seasons(id),
    episode_number INT NOT NULL,
    title VARCHAR(255) NOT NULL,
    synopsis TEXT,
    duration_minutes INT,
    video_file_key TEXT NOT NULL, -- S3 key to original video
    created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_episodes_season ON episodes(season_id, episode_number);

-- Cast & Crew
CREATE TABLE people (
    id BIGSERIAL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    bio TEXT,
    photo_url TEXT
);

CREATE INDEX idx_people_name ON people(name);

-- Title cast (many-to-many)
CREATE TABLE title_cast (
    title_id BIGINT NOT NULL REFERENCES titles(id),
    person_id BIGINT NOT NULL REFERENCES people(id),
    role VARCHAR(50) NOT NULL, -- 'actor', 'director', 'writer', 'producer'
    character_name VARCHAR(255), -- For actors
    PRIMARY KEY (title_id, person_id, role)
);

CREATE INDEX idx_title_cast_person ON title_cast(person_id);

-- Video encodings (available qualities)
CREATE TABLE video_encodings (
    id BIGSERIAL PRIMARY KEY,
    video_id BIGINT NOT NULL, -- Foreign key to episodes.id or titles.id (for movies)
    resolution VARCHAR(10) NOT NULL, -- '4K', '1080p', '720p', '480p'
    bitrate INT NOT NULL, -- Mbps (e.g., 25000 = 25 Mbps)
    codec VARCHAR(10) NOT NULL, -- 'h264', 'hevc'
    cdn_path TEXT NOT NULL, -- Path on Open Connect: videos/encoded/video123/4k/
    num_segments INT NOT NULL,
    duration_seconds INT NOT NULL,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

```

```

CREATE INDEX idx_video_encodings_video ON video_encodings(video_id);

-- Licensing (which titles available in which countries)
CREATE TABLE title_licenses (
    title_id BIGINT NOT NULL REFERENCES titles(id),
    country_code CHAR(2) NOT NULL, -- 'US', 'GB', 'IN', etc.
    is_available BOOLEAN DEFAULT TRUE,
    available_from DATE,
    available_until DATE, -- NULL = indefinite
    PRIMARY KEY (title_id, country_code)
);

CREATE INDEX idx_title_licenses_title ON title_licenses(title_id);
CREATE INDEX idx_title_licenses_country ON title_licenses(country_code) WHERE
is_available = TRUE;

-- User profiles
CREATE TABLE profiles (
    id BIGSERIAL PRIMARY KEY,
    user_id BIGINT NOT NULL REFERENCES users(id),
    name VARCHAR(50) NOT NULL,
    avatar_url TEXT,
    is_kids_profile BOOLEAN DEFAULT FALSE,
    language VARCHAR(10), -- 'en', 'es', 'ja'
    created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_profiles_user ON profiles(user_id);

-- My List (saved titles per profile)
CREATE TABLE profile_my_list (
    profile_id BIGINT NOT NULL REFERENCES profiles(id),
    title_id BIGINT NOT NULL REFERENCES titles(id),
    added_at TIMESTAMPTZ DEFAULT NOW(),
    PRIMARY KEY (profile_id, title_id)
);

CREATE INDEX idx_profile_my_list_profile ON profile_my_list(profile_id, added_at
DESC);

```

## 6. Cassandra (Viewing History)

**Purpose:** Store high-write, high-volume viewing data (billions of playback events)

### Schema Overview:

```

-- Viewing history (billions of rows)
CREATE TABLE viewing_history (
    profile_id BIGINT,
    viewed_at TIMESTAMP,
    title_id BIGINT,
    episode_id BIGINT, -- NULL for movies

```

```

duration_watched INT, -- Seconds watched (0-7200)
total_duration INT, -- Total video duration
device_type TEXT, -- 'tv', 'mobile', 'web'
quality TEXT, -- '4K', '1080p', '720p'
PRIMARY KEY (profile_id, viewed_at, title_id)
) WITH CLUSTERING ORDER BY (viewed_at DESC);

-- Continue watching (last playback position per title)
CREATE TABLE continue_watching (
    profile_id BIGINT,
    title_id BIGINT,
    episode_id BIGINT, -- NULL for movies, episode ID for series
    position_seconds INT,
    total_duration INT,
    last_watched TIMESTAMP,
    PRIMARY KEY (profile_id, title_id)
);

-- Title ratings (thumbs up/down)
CREATE TABLE profile_ratings (
    profile_id BIGINT,
    title_id BIGINT,
    rating INT, -- 1 = thumbs up, -1 = thumbs down
    rated_at TIMESTAMP,
    PRIMARY KEY (profile_id, title_id)
);

```

## Why Cassandra over PostgreSQL?

- **Write-heavy:** 50M concurrent streams × position updates every 30s = 1.6M writes/sec
- **Time-series data:** Viewing history partitioned by profile + timestamp
- **Linear scaling:** Add more nodes to handle more writes
- **No complex joins:** Queries are simple (get profile's history)

## 7. Redis (Session State)

**Purpose:** Store ephemeral, real-time data (playback session, continue watching cache)

### Data Structures:

```

// Current playback session
session:{profileId} → HASH {
    titleId: "title123",
    episodeId: "ep456", // For TV shows
    position: 1234, // Seconds
    isPlaying: true,
    device: "smart_tv",
    quality: "1080p",
    lastUpdated: 1738200000
}

// Continue watching (cached, hot data)

```

```

continue:{profileId} → LIST [
  {titleId: "title123", episodeId: "ep456", position: 1234, ...},
  {titleId: "title789", position: 4567, ...},
  ...
]

// User session (authentication)
auth:{sessionToken} → STRING userId:profileId

```

### Why Redis?

- Sub-millisecond latency (critical for playback position updates)
- TTL support (session expires after 12 hours of inactivity)
- Atomic operations (update position atomically)

## 8. S3 Storage (Video Files)

**Purpose:** Store original and encoded video files

**Bucket Structure:**

```

netflix-video-master/
  originals/
    2026/01/28/movie123_original.mov  (4K ProRes, 500 GB)

netflix-video-encoded/
  4k/
    movie123/
      segment_0.m4s
      segment_1.m4s
      ...
  1080p/
    movie123/
      segment_0.m4s
      segment_1.m4s
      ...
  720p/
    movie123/
      segment_0.m4s
      ...
  480p/
    movie123/
      segment_0.m4s
      ...

```

### Why S3?

- Cheap storage (\$0.023/GB/month)
- 99.99999999% durability (won't lose videos)
- Origin for Open Connect (push encoded videos to ISP caches)

## 9. Encoding Workers (Go + ffmpeg)

**Purpose:** Transcode original video to multiple qualities and formats

**Workflow:**

1. Content team uploads 4K ProRes master (500 GB)
2. Encoding orchestrator splits job into chunks (10-minute segments)
3. 1,000 encoding workers process in parallel
4. Each worker encodes chunk to 6 qualities (4K, 1080p, 720p, 480p, 360p, 240p)
5. Segments uploaded to S3
6. Manifest generated (MPEG-DASH .mpd file)
7. Content pushed to Open Connect caches
8. Total time: 2-4 hours for 2-hour movie

**Why Go?**

- CPU-bound task (video encoding)
- Good concurrency (goroutines process multiple chunks in parallel)
- Fast execution (wrapper around ffmpeg)

**Example Code:**

```
func encodeVideoChunk(inputPath string, outputQuality Quality) error {
    // Use ffmpeg to transcode
    cmd := exec.Command("ffmpeg",
        "-i", outputPath,
        "-c:v", "libx264", // H.264 codec
        "-preset", "slow", // Better compression (takes longer)
        "-crf", getCRF(outputQuality), // Quality factor (18 = high, 28 = low)
        "-vf", fmt.Sprintf("scale=%s", outputQuality.Resolution),
        "-c:a", "aac", // Audio codec
        "-b:a", "128k",
        "-movflags", "+faststart",
        "-f", "mp4",
        fmt.Sprintf("%s_%s.mp4", outputPath, outputQuality.Name),
    )

    return cmd.Run()
}

func getCRF(quality Quality) string {
    // CRF (Constant Rate Factor): Lower = higher quality
    switch quality.Name {
    case "4K":
        return "18" // High quality
    case "1080p":
        return "20"
    case "720p":
        return "23"
    case "480p":
        return "26"
    default:
        return "28" // Low quality
    }
}
```

```
    }
}
```

---

## 10. Recommendation Engine (Python/Spark)

**Purpose:** Generate personalized homepage for each profile

**Batch Pipeline** (runs nightly):

```
def generate_recommendations(profile_id: int):
    # 1. Get profile's viewing history (last 90 days)
    history = cassandra.query(
        "SELECT title_id, rating FROM viewing_history WHERE profile_id = ? AND
viewed_at > ?",
        [profile_id, 90_days_ago]
    )

    # 2. Find similar profiles (collaborative filtering)
    similar_profiles = find_similar_profiles(profile_id, history)

    # 3. Get titles that similar profiles liked
    recommendations = []
    for similar_profile in similar_profiles:
        their_likes = cassandra.query(
            "SELECT title_id FROM profile_ratings WHERE profile_id = ? AND rating =
1",
            [similar_profile.id]
        )
        recommendations.extend(their_likes)

    # 4. Remove titles already watched
    already_watched = set([h.title_id for h in history])
    filtered = [r for r in recommendations if r.title_id not in already_watched]

    # 5. Rank by popularity among similar profiles
    ranked = rank_by_frequency(filtered)

    # 6. Organize into rows (genres, trending, etc.)
    rows = [
        {"name": "Because you watched X", "titles": ranked[:20]},
        {"name": "Trending now", "titles": get_trending()},
        {"name": "Action & Adventure", "titles": filter_by_genre(ranked, "Action")},
        ...
    ]

    # 7. Cache in Redis (fast access)
    redis.set(f"recs:{profile_id}", json.dumps(rows), ex=86400) # 24 hours

    return rows
```

---

**Real Netflix:** Uses deep learning (matrix factorization, neural networks), but concept is similar.

## 4 API & Interface Design

### REST API Endpoints

#### Browse & Discovery

```
// Get personalized homepage
GET /api/browse/home?profile_id=123
Response: {
  "rows": [
    {
      "id": "row_trending",
      "name": "Trending Now",
      "titles": [
        {
          "id": "title456",
          "title": "Stranger Things",
          "type": "series",
          "poster_url": "https://cdn.netflix.com/posters/title456.jpg",
          "synopsis": "A group of kids uncover supernatural secrets...",
          "maturity_rating": "TV-14",
          "genres": ["Sci-Fi", "Horror", "Drama"],
          "release_year": 2016
        },
        ...
      ]
    },
    {
      "id": "row_continue",
      "name": "Continue Watching",
      "titles": [
        {
          "id": "title789",
          "title": "Breaking Bad",
          "current_episode": {
            "season": 2,
            "episode": 5,
            "position_seconds": 1234,
            "total_duration": 2876
          },
          ...
        },
        ...
      ]
    },
    ...
  ]
}

// Get title details
GET /api/titles/:titleId?profile_id=123
```

```
Response: {
  "id": "title456",
  "title": "Stranger Things",
  "type": "series",
  "synopsis": "In 1980s Indiana, a group of young friends witness strange supernatural forces...",
  "maturity_rating": "TV-14",
  "genres": ["Sci-Fi", "Horror", "Drama"],
  "release_year": 2016,
  "seasons": [
    {
      "season_number": 1,
      "episode_count": 8,
      "episodes": [
        {
          "episode_number": 1,
          "title": "The Vanishing of Will Byers",
          "synopsis": "On his way home from a friend's house...",
          "duration_minutes": 47
        },
        ...
      ]
    },
    ...
  ],
  "cast": [
    {
      "name": "Millie Bobby Brown",
      "character": "Eleven",
      "photo_url": "..."
    },
    ...
  ],
  "director": "The Duffer Brothers",
  "trailer_url": "https://cdn.netflix.com/trailers/title456.mp4",
  "my_rating": 1, // 1 = thumbs up, -1 = thumbs down, null = not rated
  "continue_watching": {
    "season": 2,
    "episode": 5,
    "position_seconds": 1234
  }
}

// Search
GET /api/search?q=stranger+things&profile_id=123
Response: {
  "results": [
    {
      "id": "title456",
      "title": "Stranger Things",
      "type": "series",
      "poster_url": "..."
    }
  ]
}
```

```
        "match_type": "title" // 'title', 'actor', 'genre'  
    },  
    ...  
]  
}
```

## Playback

```
// Get video manifest  
GET /playback/:titleId/manifest?episode_id=ep123&profile_id=456  
Response: {  
    "title_id": "title456",  
    "episode_id": "ep123",  
    "duration": 2876, // seconds  
    "qualities": [  
        {  
            "bitrate": 25000, // 25 Mbps (4K)  
            "resolution": "3840x2160",  
            "codec": "hevc",  
            "segments": [  
                {  
                    "index": 0,  
                    "url": "https://openconnect.netflix.com/videos/ep123/4k/segment_0.m4s",  
                    "start_time": 0,  
                    "end_time": 4,  
                    "size": 12500000 // bytes  
                },  
                ...  
            ]  
        },  
        {  
            "bitrate": 5000, // 5 Mbps (1080p)  
            "resolution": "1920x1080",  
            "codec": "h264",  
            "segments": [...]  
        },  
        ...  
    ],  
    "subtitles": [  
        {  
            "language": "en",  
            "url": "https://cdn.netflix.com/subtitles/ep123_en.vtt"  
        },  
        {  
            "language": "es",  
            "url": "https://cdn.netflix.com/subtitles/ep123_es.vtt"  
        },  
        ...  
    ],  
    "audio_tracks": [
```

```

    },
    "language": "en",
    "type": "original"
},
{
    "language": "es",
    "type": "dubbed"
},
...
],
"skip_intro": {
    "start": 60, // Skip intro at 1:00
    "end": 150 // Resume at 2:30
}
}

// Report playback progress
POST /playback/report
Request: {
    "profile_id": 456,
    "title_id": "title456",
    "episode_id": "ep123",
    "position_seconds": 1234,
    "duration_seconds": 2876,
    "quality": "1080p",
    "device_type": "smart_tv"
}
Response: {
    "success": true
}

// Get continue watching
GET /playback/continue-watching?profile_id=456
Response: {
    "titles": [
        {
            "title_id": "title456",
            "episode_id": "ep123",
            "season": 2,
            "episode": 5,
            "position_seconds": 1234,
            "total_duration": 2876,
            "last_watched": "2026-01-29T10:00:00Z"
        },
        ...
    ]
}

```

---

## Profile Management

```

// Create profile
POST /api/profiles
Request: {
  "user_id": 789,
  "name": "Kids",
  "is_kids_profile": true,
  "language": "en"
}
Response: {
  "profile_id": 456,
  "name": "Kids",
  "avatar_url": "https://cdn.netflix.com/avatars/default_kids.png"
}

// Switch profile
POST /api/profiles/:profileId/switch
Response: {
  "profile_id": 456,
  "name": "Kids",
  "session_token": "new_jwt_token_here"
}

// Add to My List
POST /api/profiles/:profileId/my-list
Request: {
  "title_id": "title456"
}
Response: {
  "success": true
}

// Rate title (thumbs up/down)
POST /api/profiles/:profileId/ratings
Request: {
  "title_id": "title456",
  "rating": 1 // 1 = thumbs up, -1 = thumbs down
}
Response: {
  "success": true
}

```

---

## Offline Downloads

```

// Download episode/movie
POST /api/downloads
Request: {
  "profile_id": 456,
  "title_id": "title456",
  "episode_id": "ep123", // NULL for movies
  "quality": "high", // 'standard' (720p) or 'high' (1080p)
}

```

```

    "device_id": "device_abc"
}
Response: {
  "download_id": "dl1789",
  "manifest_url": "https://api.netflix.com/downloads/dl1789/manifest",
  "total_size": 4500000000, // 4.5 GB
  "estimated_time": 900 // 15 minutes
}

// Get download manifest (encrypted segments)
GET /downloads/:downloadId/manifest
Response: {
  "segments": [
    {
      "url": "https://openconnect.netflix.com/downloads/ep123/1080p/segment_0.m4s",
      "size": 12500000,
      "checksum": "sha256:abcdef..."
    },
    ...
  ],
  "drm_license": {
    "license_url": "https://api.netflix.com/drm/license",
    "device_key": "encrypted_key_here",
    "expires_at": "2026-02-28T00:00:00Z" // 30 days
  }
}

// List downloads
GET /api/downloads?profile_id=456&device_id=device_abc
Response: {
  "downloads": [
    {
      "title_id": "title456",
      "episode_id": "ep123",
      "title": "Stranger Things S2E5",
      "downloaded_at": "2026-01-29T10:00:00Z",
      "size": 4500000000,
      "quality": "high",
      "expires_at": "2026-02-28T00:00:00Z"
    },
    ...
  ],
  "total_size": 45000000000, // 45 GB
  "storage_limit": 10000000000 // 100 GB per device
}

// Delete download
DELETE /api/downloads/:downloadId
Response: {
  "success": true
}

```

## 5 Data Modeling & Storage

### PostgreSQL vs Cassandra Split

#### PostgreSQL (Metadata, ~1 TB):

- Titles, seasons, episodes ( $10K \text{ titles} \times 100 \text{ KB} = 1 \text{ GB}$ )
- Cast, crew, people ( $1M \text{ people} \times 10 \text{ KB} = 10 \text{ GB}$ )
- Profiles ( $260M \times 5 \text{ profiles} \times 1 \text{ KB} = 1.3 \text{ TB}$ )
- My List ( $260M \text{ profiles} \times 50 \text{ titles} \times 50 \text{ bytes} = 650 \text{ GB}$ )

#### Cassandra (Viewing Data, ~500 TB):

- Viewing history:  $50M \text{ streams/day} \times 30 \text{ position updates} \times 100 \text{ bytes} \times 365 \text{ days} = 550 \text{ TB/year}$
- Continue watching:  $260M \text{ profiles} \times 10 \text{ titles} \times 200 \text{ bytes} = 520 \text{ GB}$
- Ratings:  $260M \text{ profiles} \times 100 \text{ ratings} \times 50 \text{ bytes} = 1.3 \text{ TB}$

---

## Storage Cost Analysis

### Video Storage (S3):

```
10,000 titles (mix of movies + TV episodes)
Average movie: 2 hours
Average TV show: 10 episodes/season × 8 seasons × 45 min = 60 hours
```

Assume 50/50 mix:

```
- 5,000 movies × 2 hours = 10,000 hours
- 5,000 TV shows × 60 hours = 300,000 hours
Total: 310,000 hours of content
```

Encodings per hour:

```
- 6 qualities (240p, 360p, 480p, 720p, 1080p, 4K)
- Average size: 10 GB/hour (weighted average across qualities)
Total: 310,000 hours × 10 GB = 3,100,000 GB = 3.1 PB
```

Cost:  $3.1 \text{ PB} \times \$23/\text{TB/month} = \$71,300/\text{month} = \$855\text{K/year}$

But with compression + per-title encoding optimization: ~2 PB

Cost:  $2 \text{ PB} \times \$23/\text{TB/month} = \$46,000/\text{month} = \$552\text{K/year}$

---

## 6 Core System Flows

### Flow 1: User Watches Video

```
sequenceDiagram
    participant U as User (Smart TV)
    participant API as API Server
    participant PlaybackAPI as Playback API
    participant LC as License Checker
    participant ISPCache as ISP Cache (Open Connect)
    participant Cassandra
```

```

U->>API: GET /titles/title456<br/>(fetch metadata)
API->>DB: Fetch title details
DB-->API: {title, synopsis, cast, ...}
API-->U: Title metadata

U->>U: User clicks "Play"

U->>PlaybackAPI: GET /playback/title456/manifest<br/>{episode_id, profile_id}

PlaybackAPI-->LC: Check license<br/>{title_id, country_code}
LC->>DB: SELECT * FROM title_licenses WHERE title_id = $1 AND country_code = $2

alt Title not licensed
  LC-->>PlaybackAPI: Not available
  PlaybackAPI-->U: Error: Not available in your region
else Title licensed
  LC-->>PlaybackAPI: Available

PlaybackAPI-->>DB: Get available encodings
DB-->>PlaybackAPI: [4K, 1080p, 720p, 480p, ...]

PlaybackAPI-->U: Manifest {qualities, segments, subtitles}

U->>U: Select initial quality (4K)<br/>based on connection speed

U->>ISPCache: GET segment_0.m4s (4K)

alt Cache Hit (95% probability)
  ISPCache-->U: Video segment (12 MB)
else Cache Miss
  ISPCache-->OCPCache: Fetch from regional cache
  OCPCache-->>ISPCache: Video segment
  ISPCache-->U: Video segment
end

U->>U: Decode H.265 video<br/>Display on screen

U->>PlaybackAPI: POST /playback/report<br/>{position: 0, quality: "4K"}
PlaybackAPI-->Cassandra: INSERT INTO viewing_history
PlaybackAPI-->Redis: UPDATE session:{profile_id}

Note over U: User watches for 5 minutes...

U->>ISPCache: GET segment_75.m4s (prefetch)
ISPCache-->>U: Video segment

Note over U: Network degrades, ABR switches to 1080p

U->>ISPCache: GET segment_76.m4s (1080p)

```

```
ISPCache-->U: Smaller segment (5 MB)
end
```

## Flow 2: Adaptive Bitrate Switching

**Scenario:** User on cellular network, bandwidth fluctuates 25 Mbps → 3 Mbps

```
class AdaptiveBitrateController {
    qualities = ['4K', '1080p', '720p', '480p', '360p', '240p'];
    currentQuality = '1080p';
    bandwidthHistory: number[] = [];

    async fetchSegment(segmentIndex: number): Promise<Blob> {
        const startTime = Date.now();

        // Fetch segment at current quality
        const url = this.getSegmentURL(segmentIndex, this.currentQuality);
        const segment = await fetch(url).then(r => r.blob());

        const downloadTime = Date.now() - startTime;
        const bandwidth = (segment.size * 8) / (downloadTime / 1000); // bits per second

        // Track bandwidth history (moving average)
        this.bandwidthHistory.push(bandwidth);
        if (this.bandwidthHistory.length > 5) {
            this.bandwidthHistory.shift();
        }

        // Adjust quality based on bandwidth
        this.adjustQuality();

        return segment;
    }

    adjustQuality() {
        const avgBandwidth = this.average(this.bandwidthHistory);

        // Get required bandwidth for current quality
        const requiredBandwidth = this.getRequiredBandwidth(this.currentQuality);

        // If bandwidth dropped significantly, switch to lower quality
        if (avgBandwidth < requiredBandwidth * 1.2) {
            const lowerQuality = this.getLowerQuality(this.currentQuality);
            if (lowerQuality) {
                console.log(`Dropping quality: ${this.currentQuality} → ${lowerQuality}`);
                this.currentQuality = lowerQuality;
            }
        }
    }
}

// If bandwidth increased significantly and buffer is healthy, switch to higher
```

```

quality
  if (avgBandwidth > requiredBandwidth * 2.5 && this.bufferHealth > 30) {
    const higherQuality = this.getHigherQuality(this.currentQuality);
    if (higherQuality) {
      console.log(`Increasing quality: ${this.currentQuality} →
${higherQuality}`);
      this.currentQuality = higherQuality;
    }
  }
}

getRequiredBandwidth(quality: string): number {
  const bitrates = {
    '4K': 25_000_000,           // 25 Mbps
    '1080p': 5_000_000,        // 5 Mbps
    '720p': 3_000_000,         // 3 Mbps
    '480p': 1_500_000,         // 1.5 Mbps
    '360p': 700_000,           // 0.7 Mbps
    '240p': 300_000,           // 0.3 Mbps
  };
  return bitrates[quality];
}
}

```

### Flow 3: Video Encoding Pipeline

**Scenario:** Content team uploads new movie "The Matrix"

```

sequenceDiagram
  participant CT as Content Team
  participant S3Master as S3 Master Storage
  participant Orchestrator as Encoding Orchestrator
  participant Workers as Encoding Workers (1000 instances)
  participant S3Encoded as S3 Encoded Storage
  participant OC as Open Connect Caches
  participant DB as PostgreSQL

  CT->>S3Master: Upload 4K ProRes master<br/>(500 GB, 2 hours)
  S3Master-->>CT: Upload complete

  CT->>Orchestrator: Trigger encoding job<br/>{video_id: "matrix", duration:
  7200s}

  Orchestrator->>Orchestrator: Split into chunks<br/>(12 chunks × 10 min each)

  loop For each chunk (parallel)
    Orchestrator->>Workers: Encode chunk_0<br/>to 6 qualities
    Workers->>S3Master: Download chunk_0 (42 GB)

    par Encode 4K
      Workers->>Workers: ffmpeg -crf 18 -s 3840x2160
    end
  end

```

```

    Workers->>S3Encoded: Upload 4K segments
    and Encode 1080p
        Workers->>Workers: ffmpeg -crf 20 -s 1920x1080
        Workers->>S3Encoded: Upload 1080p segments
    and Encode 720p
        Workers->>Workers: ffmpeg -crf 23 -s 1280x720
        Workers->>S3Encoded: Upload 720p segments
    and Encode 480p
        Workers->>Workers: ffmpeg -crf 26 -s 854x480
        Workers->>S3Encoded: Upload 480p segments
    and Encode 360p
        Workers->>Workers: ffmpeg -crf 28 -s 640x360
        Workers->>S3Encoded: Upload 360p segments
    and Encode 240p
        Workers->>Workers: ffmpeg -crf 30 -s 426x240
        Workers->>S3Encoded: Upload 240p segments
    end

    Workers-->>Orchestrator: Chunk_0 complete
end

Orchestrator->>DB: INSERT INTO video_encodings<br/>(cdn_paths, num_segments,
...)

Orchestrator->>OC: Push popular qualities to ISP caches<br/>(4K, 1080p off-peak
hours)

Orchestrator-->>CT: Encoding complete<br/>(Total time: 3 hours)

```

#### Encoding Cost:

```

1,000 encoding workers (AWS c5.4xlarge: $0.68/hour)
3 hours encoding time
Cost: 1,000 × $0.68 × 3 = $2,040 per movie

For 10,000 titles (amortized over years): ~$20M total encoding cost

```

#### Flow 4: Continue Watching (Cross-Device)

**Scenario:** User watches on Smart TV, pauses, continues on phone

```

sequenceDiagram
    participant TV as Smart TV
    participant PlaybackAPI as Playback API
    participant Redis
    participant Cassandra
    participant Phone as Phone

    TV->>PlaybackAPI: POST /playback/report<br/>{position: 1234, ...}
    PlaybackAPI->>Redis: HSET session:{profile_id}<br/>{position: 1234, title_id,
episode_id}

```

```

PlaybackAPI->>Cassandra: INSERT INTO continue_watching<br/>(async, eventual
consistency)

Note over TV: User pauses, closes TV app

Note over Phone: 5 minutes later, user opens app on phone

Phone->>PlaybackAPI: GET /playback/continue-watching<br/>{profile_id}

PlaybackAPI->>Redis: GET session:{profile_id}

alt Redis has recent data (< 1 hour old)
    Redis-->>PlaybackAPI: {title_id, episode_id, position: 1234}
        PlaybackAPI-->>Phone: Continue watching list
else Redis expired, fall back to Cassandra
    PlaybackAPI->>Cassandra: SELECT * FROM continue_watching WHERE profile_id =
?
    Cassandra-->>PlaybackAPI: Continue watching data
    PlaybackAPI->>Redis: Cache result (1 hour TTL)
        PlaybackAPI-->>Phone: Continue watching list
end

Phone->>Phone: Display "Continue watching The Matrix"<br/>with thumbnail at
20:34

Phone->>PlaybackAPI: GET /playback/matrix/manifest?resume_from=1234
PlaybackAPI-->>Phone: Manifest (start from segment 308)

Phone->>ISPCache: GET segment_308.m4s (1080p)
ISPCache-->>Phone: Video segment

Phone->>Phone: Resume playback at 20:34

```

## Flow 5: Recommendation Generation

**Batch Job** (runs nightly for all 260M profiles):

```

from pyspark.sql import SparkSession

def generate_recommendations_batch():
    spark = SparkSession.builder.appName("NetflixRecs").getOrCreate()

    # 1. Load viewing history (last 90 days) from Cassandra
    viewing_history = spark.read \
        .format("org.apache.spark.sql.cassandra") \
        .options(table="viewing_history", keyspace="netflix") \
        .load() \
        .filter("viewed_at > date_sub(current_date(), 90)")

    # 2. Compute user-user similarity matrix (collaborative filtering)
    # Users who watched same titles are similar

```

```

user_similarity = compute_cosine_similarity(viewing_history)

# 3. For each user, find top 100 similar users
similar_users = user_similarity.groupBy("user_id") \
    .agg(F.collect_list("similar_user_id").alias("similar_users"))

# 4. Get titles that similar users liked (but this user hasn't watched)
recommendations = viewing_history.alias("vh1") \
    .join(similar_users, "user_id") \
    .join(viewing_history.alias("vh2"),
          F.expr("array_contains(similar_users, vh2.profile_id)")) \
    .filter("vh1.title_id != vh2.title_id") \
    .groupBy("vh1.profile_id", "vh2.title_id") \
    .agg(F.count("*").alias("score")) \
    .orderBy(F.desc("score"))

# 5. Organize into rows (Top Picks, Trending, Genres)
top_picks = recommendations.limit(20)

trending = get_trending_titles() # Separate query

genre_recs = recommendations.join(title_genres, "title_id") \
    .groupBy("profile_id", "genre") \
    .agg(F.collect_list("title_id").alias("title_ids"))

# 6. Write results to Redis (fast access)
final_recs = combine_rows(top_picks, trending, genre_recs)

final_recs.foreach(lambda row:
    redis.setex(f"recs:{row.profile_id}", 86400, json.dumps(row.rows))
)

spark.stop()

# Run nightly at 2 AM
schedule.every().day.at("02:00").do(generate_recommendations_batch)

```

## 7 Consistency, Ordering & Concurrency

### Eventual Consistency (Continue Watching)

**Scenario:** User pauses on TV at 20:34, opens phone 10 seconds later

**Question:** Does phone show exact position (20:34)?

**Answer:** Usually yes, sometimes 20:30 (eventual consistency)

### Why?

- TV reports position to Redis (instant)
- TV also writes to Cassandra (async, 1-5 second lag)
- If Redis fails, phone reads from Cassandra (might be 4 seconds stale)

- Acceptable trade-off (user rewinds 4 seconds, not a big deal)

**Trade-off:** Strong consistency = slower writes, eventual = faster writes + slight staleness

---

## View Count Consistency

**Problem:** "Stranger Things" viewed 100M times, need accurate count for metrics

**Solution:** Eventual consistency with daily reconciliation

```
// Real-time (approximation)
async function recordView(profileId: string, titleId: string) {
  // 1. Write to Cassandra (source of truth)
  await cassandra.execute(
    'INSERT INTO viewing_history (profile_id, viewed_at, title_id, ...) VALUES (?, ?, ?, ?)'
  );

  // 2. Increment cached count in Redis (fast)
  await redis.hincrby(`title:${titleId}:stats`, 'view_count', 1);
}

// Daily reconciliation (cron job)
async function reconcileViewCounts() {
  const titles = await db.query('SELECT id FROM titles WHERE updated_at > NOW() - INTERVAL \'1 day\'');

  for (const title of titles) {
    // Count from Cassandra (source of truth)
    const actualCount = await cassandra.execute(
      'SELECT COUNT(*) FROM viewing_history WHERE title_id = ? AND viewed_at > ?',
      [title.id, Date.now() - 86400000]
    );

    // Update PostgreSQL
    await db.query('UPDATE titles SET view_count = view_count + $1 WHERE id = $2',
      [actualCount, title.id]);

    // Update Redis cache
    await redis.hset(`title:${title.id}:stats`, 'view_count', actualCount);
  }
}
```

---

## 8 Caching Strategy

### Multi-Layer Caching

```
Layer 1: Client Cache (Smart TV app)
↓
Layer 2: ISP Cache (Open Connect inside Comcast/Verizon)
↓
```

```
Layer 3: Open Connect Appliances (Netflix regional cache)
↓
Layer 4: Application Cache (Redis)
↓
Layer 5: Database (PostgreSQL/Cassandra)
```

---

## Layer 1: Client-Side Cache

**What to Cache:**

- Title metadata (recently viewed)
- Thumbnail images (homepage posters)
- Continue watching list
- User preferences (subtitle language, playback speed)

**Smart TV Storage:** 1-2 GB for cached data

---

## Layer 2: ISP Cache (Open Connect)

**What to Cache:** All video segments (95% of traffic)

**Cache Strategy:**

- **Popular content:** Pre-pushed during off-peak hours (2 AM - 6 AM)
- **Long-tail content:** Fetched on-demand, cached for 30 days
- **Cache size:** 100-200 TB per ISP data center

**Cache Hit Rate:** 95% (popular titles served from ISP cache, not Netflix origin)

**Eviction:** LRU (least recently used), with popularity weighting

---

## Layer 3: Open Connect Appliances (Regional)

**What to Cache:** Video segments, less popular content

**Cache Hit Rate:** 99% (includes Tier 1 misses)

---

## Layer 4: Application Cache (Redis)

**Continue Watching Cache**

```
continue:{profileId} → JSON [
  {titleId, episodeId, position, lastWatched},
  ...
]

// TTL: 1 hour (refreshed on access)
await redis.setex(`continue:${profileId}`, 3600, JSON.stringify(continueWatching));
```

---

## Recommendations Cache

```

recs:{profileId} → JSON {
  rows: [
    {name: "Top Picks for You", titleIds: [...]},
    {name: "Trending Now", titleIds: [...]},
    ...
  ]
}

// TTL: 24 hours (regenerated nightly)
await redis.setex(`recs:${profileId}`, 86400, JSON.stringify(recommendations));

```

## Title Metadata Cache

```

title:{titleId}:meta → JSON {
  title, synopsis, cast, genres, poster_url, ...
}

// TTL: 1 hour
await redis.setex(`title:${titleId}:meta`, 3600, JSON.stringify(metadata));

```

## 9 Scaling Strategy

### Open Connect CDN Scaling

**Current:** 15,000+ Open Connect Appliances worldwide

#### Scaling Strategy:

1. **Peering agreements:** Netflix pays ISPs to install Open Connect servers
2. **Auto-scaling:** Add more appliances when traffic increases 20%+
3. **Geographic expansion:** New ISP partnerships in emerging markets

#### Cost Model:

- Open Connect Appliance: \$10K hardware
- ISP provides rack space, power, network (free peering)
- Netflix pays \$5K/year maintenance per appliance
- Total: 15,000 appliances × \$5K = \$75M/year (vs \$170M for traditional CDN)

## Database Scaling

### PostgreSQL (Catalog Metadata)

**Current:** Single master + 50 read replicas

#### Scaling Strategy:

- **Read replicas:** Add more replicas for read traffic (up to 200 replicas)
- **Vertical scaling:** Upgrade master to larger instance (db.r6g.16xlarge)
- **Future sharding:** Shard by region (US catalog, EU catalog, APAC catalog)

## Cassandra (Viewing History)

**Current:** 200 nodes, 3 replicas, 500 TB

**Scaling Strategy:** Linear (add nodes)

- **Add 100 nodes** → 300 nodes, 750 TB capacity
  - **Rebalancing:** Automatic (Cassandra redistributes data)
- 

## 10 Fault Tolerance & Reliability

### Failure Scenarios

#### Failure 1: ISP Cache Outage

**Impact:** Users in that ISP see slower video loading (fallback to regional cache)

**Mitigation:** Open Connect automatically routes to nearest available cache

**Latency Impact:** 5ms (ISP cache) → 50ms (regional cache), still acceptable

**Recovery:** ISP restores cache, auto-syncs from regional cache

---

#### Failure 2: Entire Region Outage (AWS us-east-1)

**Impact:** Control plane unavailable (can't browse catalog, search, login)

**Mitigation:** Multi-region active-active (us-east-1 + eu-west-1)

**Recovery:**

1. Route53 health check detects us-east-1 failure (60 seconds)
2. DNS fails over to eu-west-1 (30 seconds)
3. Users reconnect to eu-west-1
4. Total RTO: 90 seconds

**Streaming Impact:** None (Open Connect CDN independent of AWS)

---

## 11 Observability & Operations

### Key Metrics

```
// Video start latency
histogram('video.start.latency_ms', latencyMs);

// Buffering events (bad UX)
counter('video.buffering', 1, { quality: '1080p', device: 'smart_tv' });

// Quality switches
counter('video.quality_switch', 1, { from: '4K', to: '1080p' });

// ISP cache hit rate
gauge('cdn.cache_hit_rate', hitRate, { isp: 'comcast' });
```

```
// Concurrent streams
gauge('playback.concurrent_streams', streamCount);
```

## 1 | 2 Security & Abuse Prevention

### DRM (Widevine)

#### Offline Downloads Protected:

```
// Widevine DRM (Google's DRM system)
async function requestDRMLicense(videoId: string, deviceId: string) {
    // 1. Generate DRM challenge
    const challenge = await videoElement.generateRequest('persistent-license',
initData);

    // 2. Send challenge to license server
    const response = await fetch('https://api.netflix.com/drm/license', {
        method: 'POST',
        body: challenge,
        headers: {
            'X-Device-ID': deviceId,
            'X-Video-ID': videoId
        }
    });

    const license = await response.arrayBuffer();

    // 3. Apply license to video element
    await videoElement.update(license);

    // 4. Video can now be decrypted and played
}
```

#### Key Points:

- License bound to device + subscription
- License expires if subscription canceled
- Prevents screen recording (HDCP on HDMI)

## 1 | 3 Selective Low-Level Design

### Adaptive Bitrate Algorithm (Simplified)

```
class NetflixABRController {
    qualities = [
        { name: '4K', bitrate: 2500000, buffer_target: 40 },
        { name: '1080p', bitrate: 5000000, buffer_target: 30 },
        { name: '720p', bitrate: 3000000, buffer_target: 20 },
        { name: '480p', bitrate: 1500000, buffer_target: 15 },
```

```

];
currentQuality = '1080p';
bufferLevel = 0; // Seconds of video buffered
bandwidthEstimate = 5000000; // bits per second

selectQuality(): Quality {
    // 1. If buffer is low, prioritize fast download (lower quality)
    if (this.bufferLevel < 10) {
        return this.qualities.find(q => q.buffer_target <= this.bufferLevel) ||
    this.qualities[3];
    }

    // 2. If buffer is healthy, select highest quality that bandwidth can support
    const sustainable = this.qualities.filter(q =>
        this.bandwidthEstimate > q.bitrate * 1.5 // 1.5x safety margin
    );

    return sustainable[0] || this.qualities[3]; // Default to 480p if nothing else
works
}

updateBandwidth(segmentSize: number, downloadTime: number) {
    const instantBandwidth = (segmentSize * 8) / (downloadTime / 1000);

    // Exponential moving average (smooth out spikes)
    this.bandwidthEstimate = 0.8 * this.bandwidthEstimate + 0.2 * instantBandwidth;
}
}
}

```

## 1 | 4 Trade-offs & Alternatives

### Chosen: Open Connect CDN (Custom)

#### Why Chosen:

- 95% cache hit rate (free peering with ISPs)
- Low latency (1-5ms from ISP cache)
- Cost savings (\$75M/year vs \$170M for CloudFront)

#### Alternative: Use CloudFront/Akamai

Approach	Latency	Cost (annual)	Control
<b>Open Connect</b>	5ms	\$75M	Full control
<b>CloudFront</b>	50ms	\$170M	Limited control

### Chosen: 4-Second Segments (MPEG-DASH)

#### Why Chosen:

- Fast seeking (jump to any 4-second boundary)
- Frequent quality switches (every 4 seconds)
- Small segment size (5-12 MB, quick download)

**Alternative:** 10-second segments (HLS)

Segment Size	Seeking Granularity	Quality Switch Frequency	Overhead
<b>4 seconds</b>	Fine (4s jumps)	High (every 4s)	Higher (more HTTP requests)
<b>10 seconds</b>	Coarse (10s jumps)	Low (every 10s)	Lower

## 1 | 5 Interviewer Discussion Notes

### Common Follow-Up Questions (New Grad Level)

**Q1: "How does Netflix handle viral content (new season of Stranger Things)?"**

**Answer:**

- **Pre-push to ISP caches:** Before release, push all episodes to ISP caches (off-peak)
- **Staggered releases:** Release episodes at different times globally (reduce peak load)
- **Open Connect scales:** More appliances activated for high-demand content
- **Recommendations:** Spread load by recommending different titles to different users

**Q2: "What if user's network is too slow for any quality?"**

**Answer:**

- **Ultra-low quality:** 240p (0.3 Mbps) works on slow 2G
- **Buffering:** Buffer more aggressively (30+ seconds) before playing
- **Error handling:** If still fails, show "Your network is too slow" message
- **Offline mode:** Suggest downloading when on Wi-Fi

**Q3: "How do you prevent account sharing (password sharing)?"**

**Answer:**

- **Device limits:** Max 4 concurrent streams (Premium plan)
- **IP tracking:** Unusual login from different country → challenge with email verification
- **Profile PIN:** Set PIN for main profile (kids can't change settings)
- **Household verification:** Require periodic verification (login from home network)

### Learning Resources for New Grads

#### Video Streaming:

- Article: "How Adaptive Bitrate Streaming Works" (Bitmovin blog)
- Video: "HTTP Live Streaming (HLS) vs MPEG-DASH" (YouTube)
- Practice: Build simple video player with HLS.js library

#### Netflix-Specific:

- Blog: "Netflix Tech Blog" (netflixtechblog.com)
- Paper: "Open Connect: Netflix's Content Delivery Network"

- Video: "How Netflix Scales Streaming" (YouTube, QCon talks)

#### Practice Project:

- Build mini Netflix:
  - Upload videos, transcode to multiple qualities (ffmpeg)
  - Segment videos (4-second chunks)
  - Stream with adaptive bitrate (HLS.js or Dash.js)
  - Viewing history (PostgreSQL or Cassandra)
  - Simple recommendations (collaborative filtering)

---

## Summary (For Interview Recap)

**System:** Netflix (video streaming, 260M subscribers, 1B hours/day)

#### Key Decisions:

1. **Open Connect CDN** (95% cache hit, \$75M/year vs \$170M traditional CDN)
2. **4-second segments** (MPEG-DASH, adaptive bitrate every 4 seconds)
3. **Cassandra for viewing history** (1.6M writes/sec, linear scaling)
4. **Multi-quality encoding** (6 qualities: 4K → 240p, per-title optimization)
5. **Collaborative filtering** (personalized homepage, nightly batch job)

#### Scale Numbers:

- 50M concurrent streams (peak)
- 2 exabytes/month bandwidth
- 200 PB video storage (all encodings)
- 1-3 seconds video start latency (P99)
- 15,000+ Open Connect Appliances worldwide

#### Hardest Problems:

- Bandwidth costs (2 EB/month, mitigated by Open Connect)
- Video encoding (1 movie → 200+ versions, 3-hour job, \$2K compute cost)
- Start latency (must buffer first 10 seconds quickly)
- Personalization (unique homepage for 260M profiles)
- Licensing complexity (title available in US, blocked in others)

#### Trade-offs Made:

- Open Connect (complexity + ISP partnerships) over CloudFront (simplicity)
- 4-second segments (frequent ABR) over 10-second (less overhead)
- Cassandra (write scaling) over PostgreSQL (familiar)
- Eventual consistency (speed) over strong consistency (accuracy)

---

#### End of Document (5,983 lines)

This design represents a **bandwidth-heavy video streaming platform** optimized for global low-latency delivery, adaptive quality, and personalized content discovery. Perfect for new grad interviews focusing on video streaming, CDN architecture, and handling massive scale!

*Prepared for new grad / junior engineer interviews (0-2 YOE).*