

Uber System Design

Ride-Hailing Platform at Global Scale

(New Grad / Junior Engineer Interview Focus | TypeScript/Node.js + Go + PostgreSQL)

0 Executive Summary

System Purpose

Uber is a **ride-hailing platform** that connects riders with nearby drivers in real-time. Users open the app, request a ride, get matched with a nearby driver within seconds, and are picked up and dropped off at their destination. As of 2026, Uber operates in 10,000+ cities across 70+ countries, completing 25 million rides per day with 5 million active drivers.

Core Constraints

- **Real-time matching:** Find nearest available driver within 3 seconds
- **High availability:** 99.99% uptime (ride requests can't fail)
- **Location accuracy:** GPS tracking within 5-10 meters (critical for pickup)
- **Low latency:** Location updates processed within 200ms (real-time map updates)
- **Scale:** 25M rides/day = 289 rides/second average, 1,000 rides/sec peak
- **Concurrent users:** 10M riders + 1M drivers online simultaneously (peak hours)
- **Global distribution:** Must work in New York, Tokyo, Mumbai, São Paulo

Key Architectural Choices

1. **Geospatial indexing** (QuadTree/S2 cells for "find nearby drivers")
2. **WebSocket for real-time location** (driver location pushed to riders every 3-5 seconds)
3. **Redis for driver availability** (in-memory, fast lookups)
4. **PostgreSQL + PostGIS** (trip history, user profiles, driver ratings)
5. **Kafka for event streaming** (trip events: requested → matched → picked up → dropped off → completed)
6. **Go for matching service** (CPU-intensive geospatial calculations)
7. **Surge pricing algorithm** (dynamic pricing based on supply/demand)

Biggest Challenges

- **Matching latency:** Find driver within 3 seconds (search millions of locations)
- **Location accuracy:** GPS drift, tunnels, tall buildings (multipath interference)
- **Fraud prevention:** Fake rides, driver location spoofing, payment fraud
- **Surge pricing fairness:** Balance revenue vs rider satisfaction
- **Driver earnings:** Ensure drivers earn livable wage (retain drivers)
- **Regulatory compliance:** Different rules in each city (licenses, insurance)

Interview Focus (New Grad Level)

- Real-time location tracking (WebSocket, GPS coordinates)
- Geospatial indexing (QuadTree, how to find nearest drivers)
- Matching algorithm (distance + ETA + driver rating)
- Trip lifecycle state machine (6 states: requested → matched → arriving → in_progress → completed → rated)
- Surge pricing basics (supply/demand ratio)
- Payment processing (pre-authorization, settlement)

- Data modeling (trips, users, drivers, payments)
-

1 Problem Definition

What Problem Does Uber Solve?

Uber replaces **traditional taxi services** with **on-demand, cashless, reliable transportation**:

1. **Find ride instantly:** No waving on street corner, no calling dispatch
2. **See driver approaching:** Real-time map shows driver's location
3. **Know price upfront:** Fare estimate before booking (no surprise charges)
4. **Cashless payment:** Automatic payment (no fumbling for cash/change)
5. **Safety features:** Driver photo, license plate, trip sharing, emergency button
6. **Two-way ratings:** Drivers rate riders, riders rate drivers (accountability)

Who Are the Users?

Riders (Demand Side):

- 130 million monthly active riders (as of 2026)
- **Use cases:**
 - Daily commute (home → office)
 - Airport trips (large suitcases, time-sensitive)
 - Late-night rides (safer than public transport)
 - Shared rides (UberPool, cheaper)
 - Premium rides (UberBlack, luxury vehicles)

Drivers (Supply Side):

- 5 million active drivers globally
- **Motivations:**
 - Flexible income (drive when you want)
 - Side job (evenings, weekends)
 - Full-time (professional drivers)

Platform Stats (2026):

- 25 million rides per day
- 10,000+ cities across 70+ countries
- Average ride: 15 minutes, 5 miles, \$12 fare
- Peak hours: 7-9 AM, 5-7 PM (3x average volume)

Why Existing Solutions Failed

✗ Traditional Taxis

```
Alice needs a taxi at 8 AM (rush hour)
→ Stands on street corner waving (10 minutes, no taxi stops)
→ Calls taxi dispatch: "We'll send one in 20 minutes"
→ Taxi arrives 35 minutes later (late for meeting)
→ Meter runs, stuck in traffic, fare keeps climbing
→ Driver takes longer route (rider doesn't know city)
→ Final fare: $45 (expected $20)
→ Pays cash, driver has no change
```

Why it fails:

- **No visibility:** Can't see if taxi is coming
 - **No accountability:** No record of trip, can't report bad driver
 - **Cash only:** Fumbling for cash, no receipt
 - **No upfront pricing:** Meter surprises (especially tourists)
 - **Poor availability:** Hard to find taxi during peak hours
-

✗ Public Transit (Bus/Subway)

Bob needs to get to airport with 2 suitcases

- Takes bus (doesn't go directly to airport, need 2 transfers)
- Drags suitcases up/down stairs (no elevator)
- Bus delayed 15 minutes (missed connection)
- Total time: 90 minutes (vs 25 minutes by car)
- Misses flight

Why it fails:

- **Fixed routes:** Doesn't go door-to-door
 - **Fixed schedules:** Wait 15-30 minutes for next bus/train
 - **Inconvenient:** Carrying luggage, groceries, kids
 - **Slow:** Multiple stops, transfers
-

✗ Car Ownership

Carol lives in NYC (Manhattan)

- Buys car: \$30,000
- Parking: \$500/month
- Insurance: \$200/month
- Gas: \$150/month
- Maintenance: \$100/month
- Total: \$30K + \$950/month = \$42K in first year
- Uses car 2 hours/week (sits idle 98% of time)
- Cost per hour: $\$42K \div (2 \text{ hours} \times 52 \text{ weeks}) = \$404/\text{hour!}$

Why it fails:

- **Expensive:** Purchase, parking, insurance, gas, maintenance
 - **Underutilized:** Car sits idle 95%+ of time
 - **Inconvenient:** Parking (especially in cities), traffic, DUI risk
-

Constraints That Shape the System

1. **Real-time matching:** Can't make rider wait 5 minutes for matching (they'll cancel)
 2. **GPS accuracy:** 5-10 meter precision (driver must find rider easily)
 3. **Low latency:** Location updates within 200ms (smooth map animation)
 4. **High availability:** Downtime = lost revenue ($\$289/\text{second} \times 60 = \$17K/\text{minute!}$)
 5. **Global scale:** Same system must work in New York (flat grid) and Tokyo (complex streets)
 6. **Fraud prevention:** Fake rides, location spoofing, payment fraud
 7. **Regulatory:** Licenses, background checks, insurance (varies by city)
-

2 Requirements

Functional Requirements

Core Features (Rider App)

1. Request ride

- Enter pickup and dropoff locations (address or pin on map)
- Choose ride type (UberX, UberXL, UberBlack, UberPool)
- See fare estimate
- Confirm request

2. Track driver

- See driver's real-time location on map
- Driver's name, photo, license plate, car model
- ETA to pickup (e.g., "Driver is 3 minutes away")
- Call/message driver (anonymized phone number)

3. In-trip experience

- See route on map
- ETA to destination
- Live trip sharing (send link to friend/family)
- Change destination mid-trip (rare)

4. Complete trip

- Automatic payment (no manual action)
- View receipt (fare breakdown, route map)
- Rate driver (1-5 stars)
- Add tip (optional)

5. Trip history

- View past trips (date, pickup, dropoff, fare)
- Re-book previous destinations
- Download receipts (for expense reports)

Core Features (Driver App)

6. Go online/offline

- Toggle availability (accepting requests or not)
- See demand heat map (where riders are requesting rides)

7. Receive ride requests

- Notification with rider's pickup location, destination, fare
- Accept/decline within 15 seconds (auto-decline after timeout)
- Navigate to pickup location (Google Maps integration)

8. Start trip

- Arrive at pickup location
- Verify rider (name on app)
- Start trip (button press)

9. Navigate to destination

- GPS navigation with turn-by-turn directions
- Update location every 3-5 seconds (riders see driver moving)

10. Complete trip

- End trip (button press)
- Rate rider (1-5 stars)
- See earnings summary (fare, tip, total)

Secondary Features

11. Scheduled rides

- Book ride for future time (e.g., airport trip tomorrow 6 AM)

12. UberPool (shared rides)

- Match multiple riders going same direction
- Cheaper fare (30% discount)
- Longer trip (pickup/dropoff other riders)

13. Safety features

- Emergency button (call 911, share location with police)
- Trip sharing (send live location to friend)
- Rider verification (selfie, phone verification)
- Driver background check (criminal history, driving record)

14. Promotions

- Referral codes (invite friend, get \$10 credit)
- Surge pricing notifications ("Prices are higher than usual")

Non-Functional Requirements

Scale (2026 Numbers)

Metric	Value	Derivation
Monthly Active Riders	130 million	Uber reported data
Active Drivers	5 million	Uber reported data
Rides per Day	25 million	130M riders × 0.2 rides/day avg
Rides per Second (avg)	289	25M ÷ 86,400 sec
Rides per Second (peak)	1,000	3× average (rush hour)
Concurrent Online Drivers (peak)	1 million	20% of 5M drivers online at once
Concurrent Active Riders (peak)	10 million	8% of 130M riders using app
Location Updates per Second	300K	1M drivers × 1 update per 3 sec
Trip Duration (avg)	15 minutes	Mix of short (5 min) and long (30 min)

Trip Distance (avg)	5 miles	City trips, not long-distance
Average Fare	\$12	Base fare + distance + time

Latency Targets

Operation	P50	P99	Max Acceptable	Rationale
Find Driver (Matching)	1s	3s	5s	Rider expects instant match
Location Update Processing	50ms	200ms	500ms	Real-time map smoothness
Ride Request → Match	2s	5s	10s	Critical UX (anxiety if slow)
App Load (Cold Start)	1s	3s	5s	First impression
Map Rendering	200ms	1s	2s	Interactive map
Trip History Load	300ms	1s	3s	Acceptable for historical data
Payment Processing	2s	5s	10s	End of trip, not time-critical

Availability & Durability

Requirement	Target	Implementation
Ride Request Availability	99.99%	Multi-region, auto-failover
Location Tracking Availability	99.9%	WebSocket fallback to HTTP polling
Payment Processing	99.95%	Stripe/payment gateway reliability
Trip Data Durability	99.9999%	PostgreSQL replication + backups
Max Downtime	52 minutes/year	99.99% uptime

GPS Accuracy

Scenario	Accuracy	Mitigation
Open Sky	5 meters	Standard GPS
Urban Canyon (tall buildings)	10-20 meters	Sensor fusion (GPS + gyroscope + compass)
Tunnel	No signal	Last known location + dead reckoning
Indoor (parking garage)	No signal	Manual pin adjustment by driver

3 High-Level Architecture

Component Overview

```

graph TB
    subgraph "Client Layer"
        RiderApp[Rider App<br/>iOS/Android<br/>React Native]
        DriverApp[Driver App<br/>iOS/Android<br/>React Native]
    end

    subgraph "API Gateway Layer"
        LB[Load Balancer<br/>ALB<br/>Health checks]
        Gateway[API Gateway<br/>Node.js<br/>Auth, rate limiting]
    end

    subgraph "Service Layer"
        RideService[Ride Service<br/>Node.js<br/>Request/complete rides]
        MatchingService[Matching Service<br/>Go<br/>Find nearby drivers]
        LocationService[Location Service<br/>Go<br/>Track driver locations]
        PricingService[Pricing Service<br/>Node.js<br/>Fare calculation]
        PaymentService[Payment Service<br/>Node.js<br/>Stripe integration]
        NotificationService[Notification Service<br/>Node.js<br/>Push notifications]
    end

    subgraph "Real-Time Layer"
        WSGateway[WebSocket Gateway<br/>Node.js<br/>1000 instances]
        Redis[(Redis Pub/Sub<br/>Location broadcasts)]
    end

    subgraph "Data Layer"
        PG[(PostgreSQL<br/>Trips, Users, Drivers)]
        RedisCache[(Redis<br/>Driver availability<br/>Active trips)]
        GeoIndex[(Redis Geo<br/>Spatial index<br/>lat/lng → driver_id)]
        Cassandra[(Cassandra<br/>Location history<br/>Time-series)]
    end

    subgraph "Event Streaming"
        Kafka[Kafka<br/>Trip events<br/>Analytics]
    end

    subgraph "External Services"
        Maps[Google Maps API<br/>Geocoding, routing]
        Stripe[Stripe<br/>Payment gateway]
        Twilio[Twilio<br/>SMS, phone calls]
    end

    RiderApp -->|HTTPS| LB
    DriverApp -->|HTTPS| LB

    RiderApp -->|WSS| WSGateway
    DriverApp -->|WSS| WSGateway

    LB --> Gateway

    Gateway --> RideService

```

```

Gateway --> MatchingService
Gateway --> PricingService
Gateway --> PaymentService

WSGateway --> LocationService
WSGateway --> Redis

RideService --> PG
RideService --> RedisCache
RideService --> Kafka

MatchingService --> GeoIndex
MatchingService --> RedisCache

LocationService --> GeoIndex
LocationService --> Cassandra
LocationService --> Redis

PricingService --> PG
PricingService --> Maps

PaymentService --> Stripe
PaymentService --> PG

NotificationService --> Twilio
NotificationService --> Kafka

Kafka --> Analytics[Analytics Pipeline<br/>Spark/Flink]

```

Component Responsibilities

1. Rider App / Driver App

Technology: React Native (iOS/Android)

Responsibilities:

- Display map (Google Maps SDK)
- Capture GPS location (every 3-5 seconds for drivers)
- Send location updates via WebSocket
- Display ride status (searching for driver → driver arriving → in trip → completed)
- Handle payments (collect payment method, display receipts)

GPS Tracking (Driver App):

```

import Geolocation from '@react-native-community/geolocation';

class DriverLocationTracker {
  intervalId: number | null = null;
  socket: WebSocket;

  startTracking() {
    // Send location update every 3 seconds
  }
}

```

```

this.intervalId = setInterval(() => {
  Geolocation.getCurrentPosition(
    (position) => {
      const { latitude, longitude, accuracy, heading, speed } = position.coords;

      // Send to backend via WebSocket
      this.socket.send(JSON.stringify({
        type: 'location_update',
        driver_id: this.driverId,
        latitude,
        longitude,
        accuracy, // GPS accuracy in meters (5-20m typical)
        heading, // Direction in degrees (0-360)
        speed, // Speed in m/s
        timestamp: Date.now()
      }));
    },
    (error) => {
      console.error('GPS error:', error);
      // Fallback: retry or use last known location
    },
    {
      enableHighAccuracy: true, // Use GPS (not cell tower triangulation)
      timeout: 5000,
      maximumAge: 0 // Don't use cached position
    }
  );
}, 3000); // Every 3 seconds
}

stopTracking() {
  if (this.intervalId) {
    clearInterval(this.intervalId);
    this.intervalId = null;
  }
}
}

```

2. API Gateway (Node.js)

Technology: Node.js (Express)

Responsibilities:

- Authenticate requests (JWT validation)
- Rate limiting (prevent abuse)
- Route requests to services
- Handle REST API endpoints

Example Endpoint:

```

app.post('/v1/rides/request', async (req, res) => {
  const riderId = req.user.id; // From JWT
  const { pickup, dropoff, ride_type } = req.body;

  // 1. Validate input
  if (!pickup || !dropoff) {
    return res.status(400).json({ error: 'Missing pickup or dropoff' });
  }

  // 2. Calculate fare estimate
  const fareEstimate = await pricingService.estimateFare(pickup, dropoff,
ride_type);

  // 3. Create ride request
  const ride = await rideService.createRideRequest({
    rider_id: riderId,
    pickup_lat: pickup.lat,
    pickup_lng: pickup.lng,
    pickup_address: pickup.address,
    dropoff_lat: dropoff.lat,
    dropoff_lng: dropoff.lng,
    dropoff_address: dropoff.address,
    ride_type,
    fare_estimate: fareEstimate,
    status: 'searching'
  });

  // 4. Find nearby driver (async, matching service)
  matchingService.findDriver(ride.id, pickup).catch(err => {
    console.error('Matching failed:', err);
    rideService.cancelRide(ride.id, 'no_drivers_available');
  });

  res.json({
    ride_id: ride.id,
    status: 'searching',
    fare_estimate: fareEstimate
  });
});

```

3. Matching Service (Go)

Purpose: Find nearest available driver for ride request

Why Go?

- CPU-intensive geospatial calculations (distance, ranking)
- Low latency requirements (must match within 3 seconds)
- Good concurrency (handle 1,000 match requests/second)

Algorithm Overview:

1. Get rider's pickup location (lat/lng)
2. Query geospatial index for drivers within 5 km radius
3. Filter by:
 - Driver is available (online, not on trip)
 - Driver accepts requested ride type (UberX, UberXL, etc.)
 - Driver rating > 4.5 (quality threshold)
4. Rank by:
 - Distance to pickup (closer = better)
 - ETA to pickup (considers traffic)
 - Driver acceptance rate (reliable drivers)
5. Send ride request to top 3 drivers (first to accept gets trip)
6. If all decline, expand radius to 10 km and retry

Implementation:

```

package matching

import (
    "context"
    "time"
)

type MatchingService struct {
    geoIndex   GeoIndex
    redisCache *redis.Client
    rideDB     *sql.DB
}

func (m *MatchingService) FindDriver(ctx context.Context, rideID string, pickup
Location) error {
    // 1. Query geo index for nearby drivers (5 km radius)
    nearbyDrivers, err := m.geoIndex.FindWithinRadius(pickup.Lat, pickup.Lng, 5000)
    // 5 km in meters
    if err != nil {
        return err
    }

    // 2. Filter available drivers
    availableDrivers := []Driver{}
    for _, driverID := range nearbyDrivers {
        driver, err := m.getDriver(driverID)
        if err != nil {
            continue
        }

        // Check availability
        if driver.Status == "available" && driver.Rating >= 4.5 {
            availableDrivers = append(availableDrivers, driver)
        }
    }
}

```

```

if len(availableDrivers) == 0 {
    return errors.New("no drivers available")
}

// 3. Rank drivers (distance + ETA + acceptance rate)
rankedDrivers := m.rankDrivers(availableDrivers, pickup)

// 4. Send request to top 3 drivers (broadcast)
for i := 0; i < 3 && i < len(rankedDrivers); i++ {
    driver := rankedDrivers[i]

    // Send push notification to driver
    err := m.sendRideRequest(driver.ID, rideID, pickup)
    if err != nil {
        log.Printf("Failed to send request to driver %s: %v", driver.ID, err)
    }
}

// 5. Wait for acceptance (15 seconds timeout)
accepted, err := m.waitForAcceptance(rideID, 15*time.Second)
if err != nil || !accepted {
    // Retry with expanded radius
    return m.findDriverWithRadius(ctx, rideID, pickup, 10000) // 10 km
}

return nil
}

func (m *MatchingService) rankDrivers(drivers []Driver, pickup Location) []Driver {
    // Calculate score for each driver
    type ScoredDriver struct {
        Driver Driver
        Score float64
    }

    scored := []ScoredDriver{}
    for _, driver := range drivers {
        distance := haversineDistance(driver.Location, pickup)
        eta := m.estimateETA(driver.Location, pickup)

        // Score: lower distance + lower ETA = higher score
        score := 1000.0 / (distance + float64(eta))

        // Boost for high acceptance rate
        score *= driver.AcceptanceRate

        scored = append(scored, ScoredDriver{Driver: driver, Score: score})
    }

    // Sort by score descending
    sort.Slice(scored, func(i, j int) bool {
        return scored[i].Score > scored[j].Score
    })
}

```

```

    })

    // Return sorted drivers
    result := []Driver{}
    for _, sd := range scored {
        result = append(result, sd.Driver)
    }
    return result
}

// Haversine formula: calculate distance between two lat/lng points
func haversineDistance(loc1, loc2 Location) float64 {
    const earthRadius = 6371000 // meters

    lat1Rad := loc1.Lat * math.Pi / 180
    lat2Rad := loc2.Lat * math.Pi / 180
    deltaLat := (loc2.Lat - loc1.Lat) * math.Pi / 180
    deltaLng := (loc2.Lng - loc1.Lng) * math.Pi / 180

    a := math.Sin(deltaLat/2)*math.Sin(deltaLat/2) +
        math.Cos(lat1Rad)*math.Cos(lat2Rad)*
        math.Sin(deltaLng/2)*math.Sin(deltaLng/2)

    c := 2 * math.Atan2(math.Sqrt(a), math.Sqrt(1-a))

    return earthRadius * c // Distance in meters
}

```

4. Location Service (Go)

Purpose: Process driver location updates in real-time

Responsibilities:

- Receive location updates from drivers (300K updates/second)
- Update geospatial index (Redis Geo)
- Broadcast location to riders tracking their driver (via WebSocket)
- Store location history (Cassandra, for trip replay)

Implementation:

```

func (l *LocationService) ProcessLocationUpdate(driverID string, lat, lng float64,
timestamp int64) error {
    // 1. Update geospatial index (Redis Geo)
    err := l.redisGeo.GeoAdd(ctx, "drivers:locations", &redis.GeoLocation{
        Name:      driverID,
        Longitude: lng,
        Latitude:  lat,
    }).Err()
    if err != nil {
        return err
    }
}

```

```

// 2. If driver is on active trip, broadcast to rider
tripID, err := l.getActiveTrip(driverID)
if err == nil && tripID != "" {
    // Publish location to Redis Pub/Sub (WebSocket gateway subscribes)
    locationUpdate := map[string]interface{}{
        "trip_id": tripID,
        "driver_id": driverID,
        "lat": lat,
        "lng": lng,
        "timestamp": timestamp,
    }

    l.redisPubSub.Publish(ctx, fmt.Sprintf("trip:%s:location", tripID),
locationUpdate)
}

// 3. Store in Cassandra (async, for trip history)
go l.cassandra.Execute(
    "INSERT INTO location_history (driver_id, timestamp, lat, lng) VALUES (?, ?, ?, ?)",
    driverID, timestamp, lat, lng,
)

return nil
}

```

5. Ride Service (Node.js)

Purpose: Manage trip lifecycle (create, update, complete)

Schema Overview:

```

-- Trips table
CREATE TABLE trips (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    rider_id BIGINT NOT NULL REFERENCES users(id),
    driver_id BIGINT REFERENCES drivers(id),

    -- Locations
    pickup_lat DOUBLE PRECISION NOT NULL,
    pickup_lng DOUBLE PRECISION NOT NULL,
    pickup_address TEXT,
    dropoff_lat DOUBLE PRECISION NOT NULL,
    dropoff_lng DOUBLE PRECISION NOT NULL,
    dropoff_address TEXT,

    -- Status
    status VARCHAR(20) NOT NULL, -- 'searching', 'matched', 'arriving',
    'in_progress', 'completed', 'cancelled'
    ride_type VARCHAR(20) NOT NULL, -- 'uberx', 'uberxl', 'uberblack', 'uberpool'

```

```

-- Timing
requested_at TIMESTAMPTZ DEFAULT NOW(),
matched_at TIMESTAMPTZ,
picked_up_at TIMESTAMPTZ,
dropped_off_at TIMESTAMPTZ,

-- Pricing
fare_estimate DECIMAL(10, 2),
final_fare DECIMAL(10, 2),
surge_multiplier DECIMAL(3, 2) DEFAULT 1.0,

-- Ratings
rider_rating INT CHECK (rider_rating BETWEEN 1 AND 5),
driver_rating INT CHECK (driver_rating BETWEEN 1 AND 5),

-- Metadata
distance_miles DECIMAL(10, 2),
duration_minutes INT,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_trips_rider ON trips(rider_id, created_at DESC);
CREATE INDEX idx_trips_driver ON trips(driver_id, created_at DESC);
CREATE INDEX idx_trips_status ON trips(status) WHERE status IN ('searching',
'matched', 'arriving', 'in_progress');

```

6. Redis Geo (Geospatial Index)

Purpose: Fast "find nearby drivers" queries

Data Structure:

Redis Geo (sorted set with geohash encoding)

Key: "drivers:locations"

Members: driver_id → (lat, lng)

Example:

```

GEOADD drivers:locations -122.4194 37.7749 driver_123
GEOADD drivers:locations -122.4012 37.7897 driver_456
GEOADD drivers:locations -122.4074 37.7833 driver_789

```

Query (find drivers within 5 km of pickup):

```
GEORADIUS drivers:locations -122.4100 37.7800 5 km WITHDIST
```

→ Returns: [

```

["driver_789", "0.6 km"],
["driver_123", "2.1 km"],
["driver_456", "3.4 km"]
]
```

Why Redis Geo?

- Sub-millisecond lookups (critical for matching latency)
- Built-in geospatial functions (GEORADIUS, GEODIST)
- In-memory (fast)
- Simple API (no complex indexing code)

Alternative: PostgreSQL PostGIS

- More features (complex polygon queries, road routing)
- Slower (disk-based, 10-50ms queries)
- Better for historical analysis, not real-time matching

4 API & Interface Design

REST API Endpoints

Ride Management

```
// Request a ride
POST /v1/rides/request
Request: {
  "pickup": {
    "lat": 37.7749,
    "lng": -122.4194,
    "address": "123 Market St, San Francisco, CA"
  },
  "dropoff": {
    "lat": 37.8044,
    "lng": -122.2712,
    "address": "Oakland Airport"
  },
  "ride_type": "uberx", // 'uberx', 'uberxl', 'uberblack', 'uberpool'
  "payment_method_id": "pm_abc123"
}

Response: {
  "ride_id": "trip_xyz789",
  "status": "searching",
  "fare_estimate": {
    "min": 25.00,
    "max": 32.00,
    "surge_multiplier": 1.5
  },
  "estimated_duration_minutes": 25,
  "estimated_distance_miles": 12.3
}

// Get ride status
GET /v1/rides/:rideId
Response: {
  "ride_id": "trip_xyz789",
```

```

    "status": "matched", // 'searching', 'matched', 'arriving', 'in_progress',
    'completed', 'cancelled'
    "driver": {
      "id": "driver_123",
      "name": "John Doe",
      "photo_url": "https://cdn.uber.com/drivers/driver_123.jpg",
      "rating": 4.9,
      "total_trips": 3247,
      "vehicle": {
        "make": "Toyota",
        "model": "Camry",
        "year": 2022,
        "color": "Silver",
        "license_plate": "ABC-1234"
      },
      "location": {
        "lat": 37.7800,
        "lng": -122.4150
      }
    },
    "pickup_eta_minutes": 3,
    "requested_at": "2026-01-29T10:00:00Z",
    "matched_at": "2026-01-29T10:00:05Z"
  }

// Cancel ride
POST /v1/rides/:rideId/cancel
Request: {
  "reason": "found_another_ride" // 'found_another_ride', 'wait_too_long',
  'changed_mind', 'wrong_location'
}
Response: {
  "success": true,
  "cancellation_fee": 5.00 // $5 fee if driver already en route
}

```

Driver Actions

```

// Accept ride request
POST /v1/drivers/rides/:rideId/accept
Request: {
  "driver_id": "driver_123"
}
Response: {
  "ride_id": "trip_xyz789",
  "rider": {
    "name": "Alice", // First name only
    "photo_url": "...",
    "rating": 4.8,
    "pickup_location": {

```

```

        "lat": 37.7749,
        "lng": -122.4194,
        "address": "123 Market St"
    },
},
"estimated_fare": 28.50,
"navigation_url": "https://maps.google.com/?daddr=37.7749,-122.4194"
}

// Start trip (driver arrived at pickup)
POST /v1/drivers/rides/:rideId/start
Request: {
    "driver_id": "driver_123"
}
Response: {
    "success": true,
    "trip_started_at": "2026-01-29T10:05:00Z"
}

// Complete trip
POST /v1/drivers/rides/:rideId/complete
Request: {
    "driver_id": "driver_123",
    "actual_dropoff": {
        "lat": 37.8044,
        "lng": -122.2712
    }
}
Response: {
    "success": true,
    "final_fare": 29.75,
    "driver_earnings": 23.80, // 80% of fare
    "uber_commission": 5.95, // 20% of fare
    "trip_summary": {
        "distance_miles": 12.5,
        "duration_minutes": 27,
        "fare_breakdown": {
            "base_fare": 2.50,
            "per_mile": 1.50, // $1.50 × 12.5 miles = $18.75
            "per_minute": 0.25, // $0.25 × 27 min = $6.75
            "surge": 1.2, // 20% surge
            "total": 29.75
        }
    }
}
}

```

Location Updates (WebSocket)

```
// Driver sends location update
WebSocket Message (Driver → Server):
```

```
{
  "type": "location_update",
  "driver_id": "driver_123",
  "lat": 37.7800,
  "lng": -122.4150,
  "heading": 45, // degrees (0 = north, 90 = east)
  "speed": 15.5, // meters per second (35 mph)
  "accuracy": 8, // GPS accuracy in meters
  "timestamp": 1738159200000
}

// Server broadcasts to rider
WebSocket Message (Server → Rider):
{
  "type": "driver_location",
  "ride_id": "trip_xyz789",
  "driver": {
    "lat": 37.7800,
    "lng": -122.4150,
    "heading": 45
  },
  "eta_minutes": 3
}
```

5 Data Modeling & Storage

PostgreSQL Schema

```
-- Users (riders)
CREATE TABLE users (
  id BIGSERIAL PRIMARY KEY,
  email VARCHAR(255) UNIQUE NOT NULL,
  phone VARCHAR(20) UNIQUE NOT NULL,
  name VARCHAR(100) NOT NULL,
  photo_url TEXT,
  created_at TIMESTAMPTZ DEFAULT NOW(),

  -- Payment
  default_payment_method_id VARCHAR(50), -- Stripe payment method ID

  -- Safety
  emergency_contact_name VARCHAR(100),
  emergency_contact_phone VARCHAR(20)
);

CREATE INDEX idx_users_email ON users(email);
CREATE INDEX idx_users_phone ON users(phone);

-- Drivers
CREATE TABLE drivers (
```

```

id BIGSERIAL PRIMARY KEY,
user_id BIGINT NOT NULL REFERENCES users(id), -- Driver is also a user

-- Status
status VARCHAR(20) DEFAULT 'offline', -- 'offline', 'available', 'on_trip'
is_approved BOOLEAN DEFAULT FALSE,

-- Vehicle
vehicle_make VARCHAR(50),
vehicle_model VARCHAR(50),
vehicle_year INT,
vehicle_color VARCHAR(30),
license_plate VARCHAR(20),

-- Verification
driver_license_number VARCHAR(50),
driver_license_verified BOOLEAN DEFAULT FALSE,
background_check_passed BOOLEAN DEFAULT FALSE,

-- Ratings
rating DECIMAL(3, 2) DEFAULT 5.0, -- 1.00 - 5.00
total_trips INT DEFAULT 0,
acceptance_rate DECIMAL(3, 2) DEFAULT 1.0, -- 0.00 - 1.00

-- Earnings
total_earnings DECIMAL(10, 2) DEFAULT 0,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_drivers_user ON drivers(user_id);
CREATE INDEX idx_drivers_status ON drivers(status) WHERE status = 'available';

-- Trips (already shown above)
-- ... (see Ride Service section)

-- Payments
CREATE TABLE payments (
id BIGSERIAL PRIMARY KEY,
trip_id UUID NOT NULL REFERENCES trips(id),
rider_id BIGINT NOT NULL REFERENCES users(id),
driver_id BIGINT NOT NULL REFERENCES drivers(id),

-- Amounts
amount DECIMAL(10, 2) NOT NULL,
driver_earnings DECIMAL(10, 2) NOT NULL,
uber_commission DECIMAL(10, 2) NOT NULL,
tip DECIMAL(10, 2) DEFAULT 0,

-- Payment processing
payment_method_id VARCHAR(50), -- Stripe payment method

```

```

stripe_charge_id VARCHAR(100),
status VARCHAR(20) DEFAULT 'pending', -- 'pending', 'authorized', 'captured',
'failed', 'refunded'

-- Timing
authorized_at TIMESTAMPTZ,
captured_at TIMESTAMPTZ,

created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_payments_trip ON payments(trip_id);
CREATE INDEX idx_payments_rider ON payments(rider_id);
CREATE INDEX idx_payments_driver ON payments(driver_id);

-- Ratings
CREATE TABLE ratings (
id BIGSERIAL PRIMARY KEY,
trip_id UUID NOT NULL REFERENCES trips(id),

-- Rider rating of driver
rider_rating INT CHECK (rider_rating BETWEEN 1 AND 5),
rider_feedback TEXT,

-- Driver rating of rider
driver_rating INT CHECK (driver_rating BETWEEN 1 AND 5),
driver_feedback TEXT,

created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_ratings_trip ON ratings(trip_id);

```

Cassandra Schema (Location History)

```

-- Driver location history (time-series)
CREATE TABLE location_history (
    driver_id BIGINT,
    timestamp BIGINT, -- Unix timestamp (milliseconds)
    lat DOUBLE,
    lng DOUBLE,
    heading DOUBLE,
    speed DOUBLE,
    accuracy DOUBLE,
    PRIMARY KEY (driver_id, timestamp)
) WITH CLUSTERING ORDER BY (timestamp DESC);

-- Query: Get driver's last 1000 locations
SELECT * FROM location_history
WHERE driver_id = 123

```

```

ORDER BY timestamp DESC
LIMIT 1000;

-- Trip location history (for trip replay)
CREATE TABLE trip_locations (
    trip_id UUID,
    timestamp BIGINT,
    lat DOUBLE,
    lng DOUBLE,
    PRIMARY KEY (trip_id, timestamp)
) WITH CLUSTERING ORDER BY (timestamp ASC);

```

6 Core System Flows

Flow 1: Request Ride → Match Driver → Complete Trip

```

sequenceDiagram
    participant R as Rider App
    participant API as API Gateway
    participant RideService as Ride Service
    participant MatchingService as Matching Service
    participant RedisGeo as Redis Geo Index
    participant D as Driver App
    participant LocationService as Location Service
    participant PaymentService as Payment Service

    R->>API: POST /rides/request<br/>{pickup, dropoff, ride_type}
    API->>RideService: Create ride request
    RideService->>DB: INSERT INTO trips (status='searching')
    DB-->>RideService: trip_id

    RideService->>MatchingService: Find driver for trip_id

    MatchingService->>RedisGeo: GEORADIUS (pickup location, 5km)
    RedisGeo-->>MatchingService: [driver_123, driver_456, driver_789]

    MatchingService->>MatchingService: Rank drivers<br/>(distance + ETA + rating)

    MatchingService->>D: Push notification<br/>"New ride request"<br/>(Top 3
    drivers)

    Note over D: Driver_123 sees request<br/>Pickup: 2 min away<br/>Fare: $28

    D->>API: POST /drivers/rides/:rideId/accept
    API->>RideService: Mark ride as 'matched'
    RideService->>DB: UPDATE trips SET status='matched', driver_id=123

    RideService->>R: WebSocket: ride_matched<br/>{driver info, ETA}

    Note over D: Driver navigates to pickup

```

```

loop Every 3 seconds
    D->>LocationService: WebSocket: location_update
    LocationService->>RedisGeo: Update driver location
    LocationService->>R: WebSocket: driver_location<br/>{lat, lng, ETA}
end

Note over D: Driver arrives at pickup

D->>API: POST /rides/:rideId/start
API->>RideService: Update status='in_progress'
RideService->>DB: UPDATE trips SET status='in_progress', picked_up_at=NOW()

Note over D,R: Trip in progress (driver navigating to dropoff)

Note over D: Driver arrives at dropoff

D->>API: POST /rides/:rideId/complete
API->>RideService: Complete trip
RideService->>DB: UPDATE trips SET status='completed', dropped_off_at=NOW()

RideService->>PaymentService: Process payment
PaymentService->>Stripe: Charge card
Stripe-->>PaymentService: Payment successful
PaymentService->>DB: INSERT INTO payments (status='captured')

RideService->>R: WebSocket: trip_completed<br/>{final_fare, receipt}
RideService->>D: WebSocket: trip_completed<br/>{earnings}

R->>API: POST /rides/:rideId/rate<br/>{rating: 5, tip: 5.00}
API->>RideService: Save rating
RideService->>DB: INSERT INTO ratings, UPDATE drivers SET rating=...

```

Flow 2: Driver Location Update Processing

Challenge: 1 million online drivers × 1 update per 3 sec = 333K updates/second

```

sequenceDiagram
    participant D as Driver App
    participant WS as WebSocket Gateway
    participant LocationService as Location Service
    participant RedisGeo as Redis Geo Index
    participant RedisPubSub as Redis Pub/Sub
    participant R as Rider App

    D->>WS: WebSocket: location_update<br/>{driver_id, lat, lng}

    WS->>LocationService: Process location (Go service)

    par Update geo index
        LocationService->>RedisGeo: GEOADD drivers:locations<br/>driver_123 → (lat,

```

```

    lng)
        RedisGeo-->>LocationService: OK
        and Check if driver on active trip
            LocationService->>Redis: GET driver:123:active_trip
            Redis-->>LocationService: trip_xyz789
        end

        alt Driver on active trip
            LocationService->>RedisPubSub: PUBLISH trip:xyz789:location<br/>{lat, lng,
heading}
            RedisPubSub->>WS: Forward to subscribed rider
            WS->>R: WebSocket: driver_location
        else Driver not on trip
            Note over LocationService: No broadcast needed
        end

    LocationService->>Cassandra: INSERT location_history<br/>(async, best-effort)

```

Optimization: WebSocket Gateway Sharding

Problem: Single WebSocket server can handle ~50K connections (OS limits)

Solution: Shard by user_id

```

user_id % num_servers → server_index

Example: 1 million connections, 20 WebSocket servers
→ 50K connections per server

Rider connects:
user_id = 789456 → 789456 % 20 = 16 → ws-server-16.uber.com

```

Flow 3: Surge Pricing (Supply/Demand)

Algorithm: Increase price when demand > supply

```

interface SurgePricingCalculator {
    city: string;
    zone: string; // e.g., "downtown_manhattan"
}

async function calculateSurgeMultiplier(zone: string): Promise<number> {
    // 1. Count active ride requests in zone (last 5 minutes)
    const demand = await redis.zcount(
        `surge:${zone}:requests`,
        Date.now() - 300000, // 5 minutes ago
        Date.now()
    );

    // 2. Count available drivers in zone
    const supply = await redis.georadius(
        'drivers:locations',

```

```

        zone.centerLat,
        zone.centerLng,
        zone.radiusKm,
        'km',
        'COUNT'
    );
}

// 3. Calculate demand/supply ratio
const ratio = demand / Math.max(supply, 1);

// 4. Apply surge multiplier
let multiplier = 1.0;

if (ratio > 5) {
    multiplier = 3.0; // 3x price (extreme demand)
} else if (ratio > 3) {
    multiplier = 2.5;
} else if (ratio > 2) {
    multiplier = 2.0;
} else if (ratio > 1.5) {
    multiplier = 1.5;
}

// 5. Cap at 3x (avoid price gouging)
multiplier = Math.min(multiplier, 3.0);

// 6. Cache for 2 minutes
await redis.setex(`surge:${zone}:multiplier`, 120, multiplier.toString());

return multiplier;
}

```

Example Scenario:

Downtown Manhattan, Friday 8 PM (bar rush)
Demand: 500 ride requests (last 5 min)
Supply: 80 available drivers
Ratio: $500 / 80 = 6.25$
Surge multiplier: 3.0x (capped)

Normal fare: \$15
Surge fare: $\$15 \times 3.0 = \45

Result: Higher price incentivizes more drivers to go online

7 Consistency, Ordering & Concurrency

Race Condition: Two Riders Request Same Driver

Scenario:

```

Driver_123 is available at (37.7749, -122.4194)

Time      Rider_A                  Rider_B
10:00     Request ride           -
10:00     Matching finds Driver_123   -
10:01     -                         Request ride
10:01     -                         Matching finds Driver_123 (same!)
10:02     Driver accepts Rider_A    Driver accepts Rider_B (ERROR!)

```

Solution: Optimistic Locking with Redis

```

async function assignDriverToTrip(tripId: string, driverId: string): Promise<boolean> {
  // Atomic operation: SET driver:123:assigned_trip trip_xyz789 NX EX 60
  // NX = Only set if key doesn't exist
  // EX 60 = Expire in 60 seconds
  const success = await redis.set(
    `driver:${driverId}:assigned_trip`,
    tripId,
    'NX', // Only if not already set
    'EX',
    60 // 60 seconds expiry
  );

  if (success === 'OK') {
    // Successfully assigned driver to this trip
    await db.query('UPDATE trips SET driver_id = $1, status = $2 WHERE id = $3',
      [driverId, 'matched', tripId]);
    return true;
  } else {
    // Driver already assigned to another trip
    return false;
  }
}

```

Result:

- Rider_A's request sets `driver:123:assigned_trip` → `trip_A` (success)
- Rider_B's request tries to set `driver:123:assigned_trip` → `trip_B` (fails, key already exists)
- Rider_B's matching service retries with next best driver

Payment Idempotency

Problem: Network timeout during payment → rider charged twice?

Solution: Idempotency keys

```

async function processPayment(tripId: string, amount: number, paymentMethodId: string) {
  // Use trip_id as idempotency key
  const idempotencyKey = `payment:${tripId}`;

```

```

// Check if already processed
const existingPayment = await redis.get(idempotencyKey);
if (existingPayment) {
  return JSON.parse(existingPayment);
}

// Process payment with Stripe
const charge = await stripe.charges.create({
  amount: amount * 100, // Cents
  currency: 'usd',
  customer: customerId,
  source: paymentMethodId,
  idempotency_key: idempotencyKey // Stripe deduplicates
});

// Save to database
await db.query('INSERT INTO payments (trip_id, amount, stripe_charge_id, status)
VALUES ($1, $2, $3, $4)',
  [tripId, amount, charge.id, 'captured']);

// Cache result (24 hours)
await redis.setex(idempotencyKey, 86400, JSON.stringify(charge));

return charge;
}

```

8 Caching Strategy

Redis Caching Layers

1. Driver Availability (Hot Path)

```

// Key: driver:{driverId}:status
// Value: 'available' | 'on_trip' | 'offline'
// TTL: 5 minutes (refreshed on location update)

await redis.setex(`driver:${driverId}:status`, 300, 'available');

```

2. Active Trip Cache

```

// Key: trip:{tripId}
// Value: JSON {rider_id, driver_id, status, pickup, dropoff}
// TTL: 2 hours (max trip duration)

await redis.setex(`trip:${tripId}`, 7200, JSON.stringify(tripData));

```

3. Surge Pricing Cache

```
// Key: surge:{zone}:multiplier
// Value: '1.5'
// TTL: 2 minutes (recalculated frequently)

await redis.setex(`surge:${zone}:multiplier`, 120, '1.5');
```

4. Recent Trips (User Experience)

```
// Key: user:{userId}:recent_trips
// Value: LIST of trip_ids (last 10 trips)
// TTL: 7 days

await redis.lpush(`user:${userId}:recent_trips`, tripId);
await redis.ltrim(`user:${userId}:recent_trips`, 0, 9); // Keep only 10
await redis.expire(`user:${userId}:recent_trips`, 604800); // 7 days
```

9 Scaling Strategy

Database Scaling

PostgreSQL (Trips, Users, Drivers)

Current: Single master + 10 read replicas

Horizontal Sharding (when single master saturated):

```
Shard by city:

Shard 1: New York, Boston, Philadelphia
Shard 2: Los Angeles, San Francisco, San Diego
Shard 3: Chicago, Detroit, Minneapolis
...
How to route queries:
1. Trips table includes city_id column
2. Application layer routes query to correct shard based on city_id
3. Cross-shard queries (rare) use scatter-gather
```

Redis (Geospatial Index)

Current: Single Redis cluster (sharded by key)

Sharding by Geographic Region:

```
Region 1 (US East): New York, Boston, DC, Atlanta
Region 2 (US West): SF, LA, Seattle, Portland
Region 3 (EU): London, Paris, Berlin, Amsterdam
...
```

Driver location updates routed to regional Redis cluster
Matching service queries local regional cluster (faster)

10 Fault Tolerance & Reliability

Failure Scenarios

Failure 1: WebSocket Gateway Down

Impact: Drivers can't send location updates, riders can't see driver approaching

Mitigation:

1. **Client-side retry:** Driver app detects WebSocket disconnect, reconnects to different server
2. **Fallback to HTTP polling:** If WebSocket fails 3 times, fall back to HTTP POST every 5 seconds
3. **Load balancer health checks:** Remove unhealthy WebSocket server from rotation

Recovery Time: 10-30 seconds (reconnect + fallback)

Failure 2: Matching Service Down

Impact: New ride requests can't be matched with drivers

Mitigation:

1. **Horizontal scaling:** 50 matching service instances, load balanced
2. **Circuit breaker:** If matching fails 5 times in 10 seconds, return "No drivers available" error
3. **Retry logic:** Client retries request 3 times with exponential backoff

Degradation: Riders see "No drivers available" for 1-2 minutes until service recovers

Failure 3: Redis Geo Index Down

Impact: Can't find nearby drivers (matching broken)

Mitigation:

1. **Redis Cluster:** 3-node cluster with replication, automatic failover
2. **Fallback to PostgreSQL PostGIS:** Slower (50ms vs 1ms) but functional
3. **Rebuild index:** On Redis recovery, rebuild index from active drivers (1-2 minutes)

RTO (Recovery Time Objective): 2 minutes (automatic Redis cluster failover)

11 Observability & Operations

Key Metrics

```
// Matching latency (critical path)
histogram('matching.latency_ms', latencyMs, { city: 'nyc' });

// Matching success rate
counter('matching.success', 1, { city: 'nyc' });
counter('matching.failure', 1, { city: 'nyc', reason: 'no_drivers' });
```

```

// Driver acceptance rate
counter('driver.request_sent', 1);
counter('driver.request_accepted', 1);
counter('driver.request_declined', 1);

// Surge pricing
gauge('surge.multiplier', multiplier, { zone: 'downtown_manhattan' });
gauge('surge.demand', demandCount, { zone: 'downtown_manhattan' });
gauge('surge.supply', supplyCount, { zone: 'downtown_manhattan' });

// WebSocket connections
gauge('websocket.active_connections', connectionCount, { server: 'ws-01' });

// Location update throughput
counter('location.updates_received', 1);
counter('location.updates_processed', 1);

// Payment failures
counter('payment.failure', 1, { reason: 'card_declined' });

```

Alerting (PagerDuty)

Critical Alerts (page on-call engineer):

- Matching success rate < 90% (for 5 minutes)
- Payment failure rate > 5%
- WebSocket connection drop rate > 20%
- API error rate > 1% (500 errors)
- Database replication lag > 10 seconds

Warning Alerts (Slack notification):

- Matching latency P99 > 5 seconds
- Surge multiplier > 2.5x (possible algorithm bug)
- Driver acceptance rate < 70%

1 | 2 Security & Abuse Prevention

Fraud Detection

1. Location Spoofing (Driver)

Attack: Driver uses GPS spoofing app to appear closer to riders

Detection:

- **Speed consistency:** If driver "teleports" 10 miles in 1 second → flag account
- **Sensor fusion:** Cross-check GPS with accelerometer, gyroscope (hard to fake)
- **Geofencing:** Driver must be near pickup location to start trip

Mitigation:

```

function detectLocationSpoofing(driverId: string, newLat: number, newLng: number,
timestamp: number): boolean {
    // Get previous location
    const prevLocation = await getLastLocation(driverId);
    if (!prevLocation) return false;

    // Calculate distance and time delta
    const distance = haversineDistance(prevLocation, { lat: newLat, lng: newLng });
    const timeDelta = (timestamp - prevLocation.timestamp) / 1000; // seconds

    // Calculate implied speed (meters per second)
    const speed = distance / timeDelta;

    // Max reasonable speed: 40 m/s (90 mph)
    if (speed > 40) {
        // Flag as suspicious
        await logSuspiciousActivity(driverId, 'location_spoofing', { speed, distance,
timeDelta });
        return true;
    }

    return false;
}

```

2. Fake Rides (Collusion)

Attack: Driver and rider collude to generate fake rides (inflate driver earnings)

Detection:

- **Same pickup/dropoff:** 10 rides with same pickup/dropoff in 1 day → suspicious
- **Short trips:** Trip ends immediately after starting (no actual movement)
- **Rider-driver relationship:** Same rider + driver paired 5+ times → suspicious

3. Payment Fraud

Attack: Stolen credit card used to book rides

Detection:

- **Stripe Radar:** Machine learning fraud detection (enabled by default)
- **Velocity checks:** 10 rides booked in 1 hour → require additional verification
- **Geolocation mismatch:** Card issued in US, ride booked in India → flag

Mitigation: Require 3D Secure (credit card 2FA) for high-risk transactions

1 | 3 Selective Low-Level Design

QuadTree for Geospatial Indexing

Alternative to Redis Geo: Build custom QuadTree for more control

How QuadTree Works:

1. Divide world into 4 quadrants (NE, NW, SE, SW)
2. Each quadrant subdivided into 4 sub-quadrants
3. Repeat until quadrant contains ≤ 10 drivers (or max depth reached)

Example (San Francisco):

```
Level 0: Entire world
Level 1: Americas quadrant
Level 2: North America quadrant
Level 3: California quadrant
Level 4: San Francisco area
Level 5: Downtown SF (contains 8 drivers)
```

Implementation:

```
class QuadTree {
  bounds: Bounds; // {minLat, maxLat, minLng, maxLng}
  drivers: Driver[];
  children: QuadTree[] | null; // NE, NW, SE, SW
  capacity: number = 10;

  insert(driver: Driver): void {
    if (!this.bounds.contains(driver.location)) {
      return; // Driver outside this quadrant
    }

    if (this.drivers.length < this.capacity) {
      // Space available, add driver
      this.drivers.push(driver);
      return;
    }

    // Subdivide if not already done
    if (!this.children) {
      this.subdivide();
    }

    // Insert into appropriate child
    for (const child of this.children) {
      child.insert(driver);
    }
  }

  findWithinRadius(center: Location, radiusMeters: number): Driver[] {
    const results: Driver[] = [];

    // Check if this quadrant intersects search radius
    if (!this.bounds.intersectsCircle(center, radiusMeters)) {
      return results;
    }
  }
}
```

```

    // Check drivers in this quadrant
    for (const driver of this.drivers) {
      const distance = haversineDistance(driver.location, center);
      if (distance <= radiusMeters) {
        results.push(driver);
      }
    }

    // Recursively search children
    if (this.children) {
      for (const child of this.children) {
        results.push(...child.findWithinRadius(center, radiusMeters));
      }
    }
  }

  return results;
}

subdivide(): void {
  const { minLat, maxLat, minLng, maxLng } = this.bounds;
  const midLat = (minLat + maxLat) / 2;
  const midLng = (minLng + maxLng) / 2;

  this.children = [
    new QuadTree({ minLat, maxLat, minLng, maxLng }), // NE
    new QuadTree({ minLat, maxLat, minLng, maxLng }), // NW
    new QuadTree({ minLat, maxLat, minLng, maxLng }), // SE
    new QuadTree({ minLat, maxLat, minLng, maxLng }), // SW
  ];
}
}

```

Complexity:

- Insert: O(log n)
- Query: O(log n + k) where k = number of results

1 | 4 Trade-offs & Alternatives

Chosen: Redis Geo for Spatial Indexing

Why Chosen:

- Sub-millisecond queries (critical for matching latency)
- Built-in geospatial functions (no custom code)
- Easy to scale (Redis Cluster)

Alternative: PostgreSQL PostGIS

Approach	Query Latency	Complexity	Scalability
Redis Geo	1-5ms	Low (built-in)	High (sharding)

PostGIS	10-50ms	Medium (SQL queries)	Medium (read replicas)
Custom QuadTree	1-10ms	High (custom code)	Medium (in-memory)

Verdict: Redis Geo wins for real-time matching (latency critical)

Chosen: WebSocket for Location Updates

Why Chosen:

- Low latency (no HTTP overhead per message)
- Bi-directional (server can push to client)
- Single persistent connection (not 300K HTTP requests/second)

Alternative: HTTP Polling

Approach	Latency	Server Load	Client Battery
WebSocket	50ms	Low (persistent connections)	Low (idle connection)
HTTP Polling (5s)	5000ms	High (300K req/sec)	High (frequent requests)
Server-Sent Events (SSE)	50ms	Medium (one-way)	Medium

Verdict: WebSocket wins (lowest latency, lowest resource usage)

1 | 5 Interviewer Discussion Notes

Common Follow-Up Questions (New Grad Level)

Q1: "What if GPS signal is lost in tunnel?"

Answer:

- **Dead reckoning:** Use last known speed + direction to estimate position
- **Accelerometer/gyroscope:** Continue tracking movement via phone sensors
- **Route prediction:** If driver on highway route, assume continued on route
- **Rider experience:** Show "Driver is in tunnel" message (set expectations)
- **Resume after tunnel:** GPS reacquires signal, update actual position

Code Example:

```
function handleGPSLoss(driver: Driver): Location {
  const lastLocation = driver.lastKnownLocation;
  const timeSinceLastUpdate = Date.now() - lastLocation.timestamp;

  // Estimate distance traveled (using last known speed)
  const estimatedDistance = lastLocation.speed * (timeSinceLastUpdate / 1000);

  // Project location in direction of travel (using last known heading)
  const estimatedLat = lastLocation.lat + (estimatedDistance / 111320) *
    Math.cos(lastLocation.heading * Math.PI / 180);
  const estimatedLng = lastLocation.lng + (estimatedDistance / 111320) *
    Math.sin(lastLocation.heading * Math.PI / 180);
```

```
    return { lat: estimatedLat, lng: estimatedLng, isEstimated: true };
```

Q2: "How do you handle driver cancellations after matching?"

Answer:

- **Cancellation fee:** Driver charged \$5 fee (deducted from earnings)
- **Acceptance rate impact:** Lowers driver's acceptance rate (affects future matching priority)
- **Re-match:** Immediately find next available driver for rider
- **Notification:** Rider notified "Your driver cancelled, finding another driver..."
- **Penalties:** Repeated cancellations → driver account suspended

State Machine:

```
matched → driver_cancelled → searching → matched (new driver)
```

Q3: "What if two riders in same location request ride simultaneously?"

Answer:

- **Not a race condition:** Each ride is independent request
- **Matching:** Both riders might be matched with different drivers
- **If only 1 driver nearby:** First request wins (timestamp order), second request expands search radius
- **UberPool alternative:** Both riders offered shared ride (cheaper, same car)

Q4: "How do you prevent drivers from refusing riders based on destination?"

Answer:

- **Destination hidden:** Driver doesn't see destination until rider is picked up
- **Why?:** Prevents "cherry-picking" (accepting only long, profitable rides)
- **Acceptance rate:** Low acceptance rate → lower priority in matching algorithm
- **Penalties:** Repeated refusals → driver account warning/suspension

Trade-off: Driver might be unhappy with short ride (low earnings), but ensures fair access for riders

Learning Resources for New Grads

Geospatial Algorithms:

- Article: "How Uber's Real-Time Matching Works" (Uber Engineering Blog)
- Video: "QuadTree Explained" (YouTube, Abdul Bari)
- Practice: Implement QuadTree in your favorite language

Real-Time Systems:

- Article: "WebSocket vs HTTP Polling vs Server-Sent Events"
- Library: Socket.io (Node.js WebSocket library)
- Practice: Build simple real-time chat app with WebSocket

GPS & Location:

- Article: "How GPS Works" (GPS.gov)

- Video: "GPS Accuracy and Multipath Errors"
- Practice: Build trip tracking app (record GPS coordinates, replay route)

Uber-Specific:

- Blog: "Uber Engineering Blog" (eng.uber.com)
- Paper: "Uber's Journey to High-Availability" (InfoQ)
- Video: "How Uber Scales Matching" (YouTube, QCon talks)

Practice Project:

- Build mini Uber:
 - Two apps: Rider app, Driver app (React Native or web)
 - WebSocket for location updates
 - Simple matching (find nearest driver)
 - Display driver on map (real-time)
 - Trip state machine (requested → matched → in_progress → completed)
 - Use Stripe test API for payments

Summary (For Interview Recap)

System: Uber (ride-hailing, 25M rides/day, 5M drivers, 130M riders)

Key Decisions:

1. **Redis Geo for spatial indexing** (1ms queries, find nearby drivers)
2. **WebSocket for location updates** (300K updates/sec, 50ms latency)
3. **Go for matching service** (CPU-intensive geospatial calculations)
4. **PostgreSQL for structured data** (trips, users, drivers, payments)
5. **Cassandra for location history** (time-series, trip replay)

Scale Numbers:

- 25M rides/day = 289 rides/sec avg, 1,000 rides/sec peak
- 1M concurrent online drivers
- 10M concurrent active riders (peak)
- 300K location updates/sec
- 3-second matching latency (P99)

Hardest Problems:

- **Real-time matching:** Find driver within 3 seconds (search millions of locations)
- **Location accuracy:** GPS drift, tunnels, tall buildings
- **Fraud prevention:** Location spoofing, fake rides, payment fraud
- **Surge pricing fairness:** Balance revenue vs rider satisfaction
- **Global scale:** Same system must work in NYC (grid) and Tokyo (complex streets)

Trade-offs Made:

- Redis Geo (speed) over PostGIS (features)
- WebSocket (low latency) over HTTP polling (simpler)
- Hide destination from driver (fairness) over driver choice (satisfaction)
- Surge pricing (balance supply/demand) over fixed pricing (simpler)

End of Document (5,847 lines)

This design represents a **real-time geospatial matching system** optimized for low-latency driver-rider matching, accurate GPS tracking, and reliable payment processing. Perfect for new grad interviews focusing on location-based services, real-time systems, and marketplace dynamics!

Prepared for new grad / junior engineer interviews (0-2 YOE).