

# Probability Basics (Developer Edition)

## What Problem This Solves

Probability helps you reason about uncertainty in systems.

Every time you deal with:

- **Failures:** Will this request succeed? Should I retry?
- **Performance:** What's the cache hit rate?
- **Testing:** How likely is this bug to appear?
- **Capacity:** What load can we handle?
- **Security:** How strong is this password?

...you're reasoning about probability.

Probability turns "I don't know" into "Here's what I can expect."

---

## Intuition & Mental Model

Think: Frequency Over Many Trials

Probability  $\approx$  "What fraction of the time does this happen?"

Coin flip:  $P(\text{heads}) = 0.5$   
→ In 1000 flips, expect ~500 heads

API success:  $P(\text{success}) = 0.99$   
→ In 1000 requests, expect ~990 successes

Not fortune telling—long-run behavior.

---

## Core Concepts

### 1. Basic Probability

Probability of event A:

$$P(A) = \frac{\text{\# of outcomes where A happens}}{\text{\# of total possible outcomes}}$$

Range:  $0 \leq P(A) \leq 1$

Example: Rolling a die

$$\begin{aligned} P(\text{rolling 3}) &= 1/6 \approx 0.167 \\ P(\text{rolling even}) &= 3/6 = 0.5 \\ P(\text{rolling 7}) &= 0/6 = 0 \text{ (impossible)} \\ P(\text{rolling } \leq 6) &= 6/6 = 1 \text{ (certain)} \end{aligned}$$

In code:

```

// Fair coin flip
function flipCoin() {
  return Math.random() < 0.5 ? 'heads' : 'tails';
}

// P(heads) = 0.5
let heads = 0;
for (let i = 0; i < 10000; i++) {
  if (flipCoin() === 'heads') heads++;
}
console.log(heads / 10000); // ~0.5

```

## 2. Complementary Events

$$P(\text{not A}) = 1 - P(A)$$

```

P(heads) = 0.5
P(tails) = 1 - 0.5 = 0.5

P(request succeeds) = 0.99
P(request fails) = 1 - 0.99 = 0.01

```

**Useful for "at least one" problems:**

```

// P(at least one success) = 1 - P(all failures)

// P(no failures in 3 tries) given P(fail) = 0.1 per try
const pAllFail = 0.1 * 0.1 * 0.1; // 0.001
const pAtLeastOneSuccess = 1 - pAllFail; // 0.999

```

## 3. Independent Events

**Two events are independent if one doesn't affect the other.**

$$P(A \text{ and } B) = P(A) \times P(B) \quad // \text{ Only if independent!}$$

**Example: Multiple coin flips**

$$\begin{aligned} P(\text{heads, then heads}) &= 0.5 \times 0.5 = 0.25 \\ P(\text{three heads in a row}) &= 0.5 \times 0.5 \times 0.5 = 0.125 \end{aligned}$$

**In systems:**

```

// Independent failures
const dbUptime = 0.99;    // 99% uptime
const apiUptime = 0.98;   // 98% uptime

// Both up (independent)
const bothUp = dbUptime * apiUptime; // 0.9702 (97.02%)

```

```
// At least one down
const atLeastOneDown = 1 - bothUp; // 0.0298 (2.98%)
```

---

## 4. Conditional Probability

$P(A | B)$  = "Probability of A given B happened"

$$P(A|B) = \frac{P(A \text{ and } B)}{P(B)}$$

### Example: Login attempts

```
P(hacker | failed login) ≠ P(failed login)

Failed logins happen often (typos)
But given 10 failed logins in a row, P(hacker) increases
```

In code:

```
// Cache effectiveness
const totalRequests = 10000;
const cacheHits = 8000;
const cacheMisses = 2000;
const slowResponses = 300; // All from cache misses

// P(slow response)
const pSlow = slowResponses / totalRequests; // 0.03

// P(slow | cache miss)
const pSlowGivenMiss = slowResponses / cacheMisses; // 0.15

// Conditional probability is higher!
```

---

## 5. Bayes' Theorem (Intuition)

Flipping conditional probabilities:

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)}$$

### Example: False positives in testing

```
Test accuracy: P(positive | bug present) = 0.95
Bug prevalence: P(bug present) = 0.01
```

If test is positive, P(bug actually present)?

$P(\text{bug} | \text{positive}) = P(\text{positive} | \text{bug}) \times P(\text{bug})$

$P(\text{positive})$

Counterintuitive: Even with 95% accurate test,  
 $P(\text{bug} \mid \text{positive})$  might be low if bugs are rare!

#### Real scenario:

```
// Health monitoring system
const pHHighLoadGivenIssue = 0.9; // Issue → high load
const pIssue = 0.05;           // 5% of time there's an issue
const pHighLoad = 0.2;         // 20% of time load is high

// P(issue | high load)?
const pIssueGivenHighLoad =
  (pHighLoadGivenIssue * pIssue) / pHighLoad;
// = (0.9 × 0.05) / 0.2 = 0.225 (22.5%)

// High load doesn't always mean issue!
```

## 6. Expected Value

Average outcome over many trials:

$$E[X] = \sum (\text{outcome} \times \text{probability})$$

Example: Dice roll

$$\begin{aligned} E[\text{dice}] &= 1 \times (1/6) + 2 \times (1/6) + 3 \times (1/6) + 4 \times (1/6) + 5 \times (1/6) + 6 \times (1/6) \\ &= (1+2+3+4+5+6) / 6 \\ &= 3.5 \end{aligned}$$

In systems:

```
// Expected response time
const scenarios = [
  { time: 10, probability: 0.7 }, // Cache hit
  { time: 100, probability: 0.25 }, // Database query
  { time: 1000, probability: 0.05 } // Timeout
];

const expectedTime = scenarios.reduce(
  (sum, {time, probability}) => sum + time * probability,
  0
);
// = 10×0.7 + 100×0.25 + 1000×0.05 = 82ms

// Expected latency: 82ms
```

## Software Engineering Connections

## 1. Retry Logic

```
async function fetchWithRetry(url, maxRetries = 3) {
  const pSuccess = 0.9; // 90% success rate per try

  for (let i = 0; i < maxRetries; i++) {
    try {
      return await fetch(url);
    } catch (error) {
      if (i === maxRetries - 1) throw error;
      // Wait before retry (exponential backoff)
      await sleep(2 ** i * 1000);
    }
  }

  // P(all 3 attempts fail) = 0.13 = 0.001
  // P(at least one succeeds) = 1 - 0.001 = 0.999 (99.9%)
}
```

## 2. Load Balancing

```
// Random load balancing
const servers = ['server1', 'server2', 'server3'];

function randomServer() {
  return servers[Math.floor(Math.random() * servers.length)];
}

// Each server gets ~1/3 of traffic (uniform distribution)
// Over 10,000 requests, each expects ~3,333 requests
```

## 3. A/B Testing

```
// Show variant A to 50% of users
function getVariant(userId) {
  const hash = simpleHash(userId);
  return hash % 2 === 0 ? 'A' : 'B';
}

// Track conversions
const results = {
  A: { shown: 5000, converted: 250 }, // 5% conversion
  B: { shown: 5000, converted: 300 } // 6% conversion
};

// Is B better? Need statistical significance test
// (covered in inferential statistics)
```

---

## 4. Cache Hit Probability

```
// LRU cache with capacity 1000
const cache = new LRUCache(1000);

// After warmup, track hit rate
let hits = 0, misses = 0;

function getData(key) {
  if (cache.has(key)) {
    hits++;
    return cache.get(key);
  }

  misses++;
  const data = fetchFromDB(key);
  cache.set(key, data);
  return data;
}

// P(cache hit) = hits / (hits + misses)
// Goal: Keep P(hit) > 0.9 (90%)
```

---

## 5. Password Strength

```
function passwordStrength(password) {
  const charset = {
    lowercase: 26,
    uppercase: 26,
    digits: 10,
    symbols: 32
  };

  let charsetSize = 0;
  if (/[^a-zA-Z]/.test(password)) charsetSize += charset.lowercase;
  if (/[^A-Z]/.test(password)) charsetSize += charset.uppercase;
  if (/[^0-9]/.test(password)) charsetSize += charset.digits;
  if (/[^a-zA-Z0-9]/.test(password)) charsetSize += charset.symbols;

  // Possible combinations
  const combinations = charsetSize ** password.length;

  // At 1 billion tries/sec, expected time to crack
  const secondsToCrack = combinations / 1e9;

  return { combinations, secondsToCrack };
}
```

```
// 8-char, all lowercase: 26^8 = 208 billion (~3.5 minutes)
// 8-char, mixed: 94^8 = 6 quadrillion (~197,000 years)
```

## 6. Bloom Filters (Probabilistic Data Structure)

```
class BloomFilter {
  constructor(size) {
    this.bits = new Array(size).fill(false);
    this.size = size;
  }

  add(item) {
    const hash1 = this.hash(item, 0) % this.size;
    const hash2 = this.hash(item, 1) % this.size;
    this.bits[hash1] = true;
    this.bits[hash2] = true;
  }

  mightContain(item) {
    const hash1 = this.hash(item, 0) % this.size;
    const hash2 = this.hash(item, 1) % this.size;
    return this.bits[hash1] && this.bits[hash2];
  }

  hash(item, seed) {
    // Simplified hash function
    let hash = seed;
    for (let char of item) {
      hash = (hash * 31 + char.charCodeAt(0)) % this.size;
    }
    return hash;
  }
}

// False positive rate ≈ (1 - e^(-k*n/m))^k
// k = hash functions, n = elements, m = bit array size

// Trade space for accuracy
```

## Common Misconceptions

✗ "Random means unpredictable for individual events"

For single events, yes. But over many trials, very predictable.

```
// Can't predict one flip
flipCoin(); // 'heads' or 'tails'?

// Can predict distribution
```

```

let heads = 0;
for (let i = 0; i < 100000; i++) {
  if (flipCoin() === 'heads') heads++;
}
console.log(heads / 100000); // ~0.5 (very close)

```

## ✖ "Past events affect independent future events"

**Gambler's fallacy:** Thinking streak affects probability

```

// "5 heads in a row, tails is 'due'"
// WRONG! Next flip still 50/50

// Coin has no memory

```

## ✖ "Conditional probability is symmetric"

$P(A|B) \neq P(B|A)$

```

P(rain | clouds) ≠ P(clouds | rain)

P(slow response | cache miss) ≠ P(cache miss | slow response)

```

## ✖ "Low probability means impossible"

**Low probability events happen with enough trials:**

```

// P(hash collision) = very small
// But with billions of operations, collisions happen

// Birthday paradox: 23 people → 50% chance of shared birthday
// Seems low, but probability adds up

```

## ✖ "Expected value is the most common outcome"

**Expected value might not even be possible:**

```

E[dice roll] = 3.5
But you can't roll 3.5!

```

## Practical Mini-Exercises

### Exercise 1: Uptime Calculation

Your system has three services:

- API: 99.9% uptime (0.999)
- Database: 99.5% uptime (0.995)
- Cache: 99.9% uptime (0.999)

If all three must be up for system to work, what's overall uptime?

► Solution

### Exercise 2: Retry Strategy

API call has 80% success rate. How many retries to get 99% overall success?

► Solution

### Exercise 3: Cache Sizing

Your cache hit rate is 85% with size 1000. Each cache hit saves 50ms.

Expected response time if base (no cache) is 60ms?

► Solution

---

## Summary Cheat Sheet

### Basic Rules

```
// Probability range  
0 ≤ P(A) ≤ 1  
  
// Complement  
P(not A) = 1 - P(A)  
  
// Independent events  
P(A and B) = P(A) × P(B)  
  
// Conditional probability  
P(A | B) = P(A and B) / P(B)  
  
// Expected value  
E[X] = Σ (outcome × probability)
```

### Common Patterns

```
// At least one success in n tries  
pSuccess = 1 - (pFail ** n)  
  
// All succeed  
pAllSuccess = pSuccess ** n  
  
// Expected time with cache  
E[time] = pHit × hitTime + pMiss × missTime  
  
// Uptime calculation  
systemUptime = service1Uptime × service2Uptime × ...
```

### Quick Reference

Scenario	Formula	Example
----------	---------	---------

Coin flip	$P = 0.5$	Heads/tails
Dice roll	$P = 1/6$ per face	Rolling 3
Independent AND	$P(A) \times P(B)$	Both succeed
Independent OR	$1 - P(\text{both fail})$	At least one succeeds
Retry success	$1 - (p\text{Fail}^n)$	3 retries

---

## Next Steps

Probability helps you reason about uncertainty and randomness in systems. You now understand how to calculate likelihoods, expected values, and make informed decisions under uncertainty.

Next, we'll explore **descriptive statistics**—understanding and summarizing data you collect from your systems.

Continue to: [07-descriptive-statistics.md](#)