

# 14. Optimization

Phase 5: Modern Software Math

~45 minutes | 🕵️ Decision Framework | ⚖️ Constraints vs Objectives

---

## What Problem This Solves

You're trying to:

- Minimize latency while staying within budget
- Maximize throughput given CPU/memory limits
- Find the best database index strategy (can't index everything)
- Allocate developers across projects to maximize productivity
- Price your SaaS to maximize profit (not just revenue)
- Route traffic to minimize response time
- Compress data as much as possible without losing quality

**Without optimization thinking**, you make decisions heuristically ("use more servers!") or arbitrarily ("let's split time equally"). You don't recognize when you're stuck in local optima or when constraints fundamentally limit solutions.

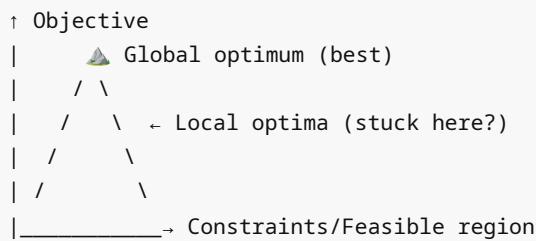
**With optimization**, you frame problems precisely, identify what you're really trying to maximize or minimize, and make decisions systematically.

---

## Intuition & Mental Model

### The Core Insight: Landscape Hiking

OPTIMIZATION = Finding the best point in a landscape



### Types of Optima:

Global Optimum: Absolute best possible solution  
Rarely achievable (NP-hard in most real problems)

Local Optimum: Best in neighborhood, but worse elsewhere  
Where greedy algorithms get stuck

Feasible: Satisfies all constraints (good enough)  
Often the real goal in practice

---

# Core Concepts

## 1. Framing the Problem

**Pattern: What do we optimize? What are constraints?**

```
function formulateProblem(scenario) {
    // Every optimization problem has 3 parts:

    return {
        objective: {
            maximize: 'revenue / cost ratio',
            or_minimize: 'latency / energy',
            description: 'What we want most'
        },
        constraints: {
            budget: '$100k',
            cpu_cores: 8,
            network: '1Gbps',
            description: 'Immovable limits'
        },
        decision_variables: {
            servers: 'how many?',
            instance_type: 'which type?',
            region: 'which region?',
            description: 'What we control'
        }
    };
}

// Example: E-commerce infrastructure
formulateProblem({
    maximize: 'Concurrent users (throughput)',
    constraints: [
        'Budget: $50k/month',
        'Team size: 2 engineers',
        'SLA: 99.9% uptime'
    ],
    variables: [
        'Instance count: 1-100',
        'Instance type: t3, m5, c5',
        'Region: us-east-1, eu-west-1'
    ]
});
```

**Three Key Questions:**

1. What are we really trying to optimize? (revenue, latency, accuracy, cost)
2. What constraints are actually hard? (regulatory, physical, financial)
3. What can we actually change? (budget, team, architecture)

## 2. Linear Programming (Intuition)

**Problem:** Maximize profit given constraints

```
function linearProgrammingExample() {
    // Bakery: Make brownies and cookies
    // Brownies: $2 profit, 1 hour baking time, 2 cups flour
    // Cookies: $1.50 profit, 0.5 hours baking time, 1 cup flour
    // Constraints: 40 hours/week, 80 cups flour/week

    // Decision variables: b = brownies, c = cookies
    // Maximize: 2b + 1.5c
    // Subject to:
    //   b + 0.5c ≤ 40 (time constraint)
    //   2b + c ≤ 80   (flour constraint)
    //   b, c ≥ 0      (non-negativity)

    // Feasible region = polygon
    // Optimum at a corner point (linear programming property)

    const corners = [
        { b: 0, c: 0, profit: 0 },
        { b: 40, c: 0, profit: 80 },           // Make only brownies
        { b: 0, c: 80, profit: 120 },         // Make only cookies
        { b: 20, c: 40, profit: 100 },        // Mix
    ];

    return corners
        .sort((a, b) => b.profit - a.profit)[0];
}

console.log(linearProgrammingExample());
// { b: 0, c: 80, profit: 120 }
// Make only cookies! (better margin given constraints)
```

### Real Example: Server Allocation

```
function allocateServers() {
    // Web servers: $500/month, 100 requests/sec, low-cost
    // API servers: $1000/month, 50 requests/sec, high-margin
    // Budget: $5000/month, need 300 req/sec minimum

    const budget = 5000;
    const minThroughput = 300;

    const solutions = [];

    for (let web = 0; web <= 10; web++) {
        for (let api = 0; api <= 5; api++) {
            const cost = web * 500 + api * 1000;
            const throughput = web * 100 + api * 50;
            if (cost <= budget && throughput >= minThroughput) {
                solutions.push({ web, api, cost, throughput });
            }
        }
    }
}
```

```

const profit = throughput * 10; // $10 per req

if (cost <= budget && throughput >= minThroughput) {
  solutions.push({
    webServers: web,
    apiServers: api,
    cost,
    throughput,
    profit
  });
}
}

return solutions.sort((a, b) => b.profit - a.profit)[0];
}

console.log(allocateServers());
/* {
  webServers: 10,
  apiServers: 0,
  cost: 5000,
  throughput: 1000,
  profit: 10000
}
// Web servers more cost-effective in this scenario */

```

### 3. Greedy Algorithms: Local Optimization

```

function greedyScheduling(tasks) {
  // Greedy: Always pick the task with highest value first
  // (Not optimal globally, but fast)

  const sorted = [...tasks].sort((a, b) => b.value - a.value);
  let time = 0;
  let totalValue = 0;
  const scheduled = [];

  for (const task of sorted) {
    if (time + task.duration <= 8) { // 8-hour day
      scheduled.push(task.name);
      time += task.duration;
      totalValue += task.value;
    }
  }

  return { scheduled, totalValue, timeUsed: time };
}

// Tasks: [name, duration, value]
const tasks = [

```

```

{ name: 'Feature A', duration: 3, value: 50 },
{ name: 'Feature B', duration: 2, value: 40 },
{ name: 'Feature C', duration: 4, value: 60 },
{ name: 'Bug fix', duration: 1, value: 30 }
];

greedyScheduling(tasks);
/* {
  scheduled: [ 'Feature C', 'Feature A', 'Bug fix' ],
  totalValue: 140,
  timeUsed: 8
}
// Greedy picked high-value tasks first, packed the 8-hour day well
// (Not always optimal: sometimes smaller high-value tasks beat bigger ones) */

```

### When Greedy Fails:

```

function greedyFailsExample() {
  const budget = 100;
  const items = [
    { name: 'Item A', cost: 50, value: 60 }, // 1.2x value/cost
    { name: 'Item B', cost: 30, value: 40 }, // 1.33x ← Greedy picks first
    { name: 'Item C', cost: 20, value: 20 }, // 1.0x
    { name: 'Item D', cost: 50, value: 70 } // 1.4x but can't fit with A
  ];

  // Greedy (by value/cost): A + B = cost 80, value 100
  // Optimal: B + D = cost 80, value 110

  // Greedy got stuck locally
}

```

## 4. Dynamic Programming: Optimal Substructure

**Key Idea:** If optimal solution contains smaller optimal solutions, use recursion + memoization

```

function fibonacciOptimization() {
  // Naive recursion: O(2^n) - exponential!
  function naiveFib(n) {
    if (n <= 1) return n;
    return naiveFib(n - 1) + naiveFib(n - 2); // Recalculates same values!
  }

  // Dynamic programming: O(n) - linear!
  function dpFib(n, memo = {}) {
    if (n in memo) return memo[n];
    if (n <= 1) return n;

    memo[n] = dpFib(n - 1, memo) + dpFib(n - 2, memo);
    return memo[n];
  }
}

```

```

// Time comparison for n=40
console.time('Naive');
naiveFib(35); // Takes ~3 seconds
console.timeEnd('Naive');

console.time('DP');
dpFib(35); // Takes ~1ms
console.timeEnd('DP');
}

// Key: Memoization avoids recalculating overlapping subproblems

```

### Real Example: Longest Common Subsequence

```

function longestCommonSubsequence(str1, str2) {
    // DP table: lcs[i][j] = LCS of str1[0..i] and str2[0..j]
    const m = str1.length, n = str2.length;
    const dp = Array(m + 1).fill(0).map(() => Array(n + 1).fill(0));

    for (let i = 1; i <= m; i++) {
        for (let j = 1; j <= n; j++) {
            if (str1[i - 1] === str2[j - 1]) {
                dp[i][j] = dp[i - 1][j - 1] + 1;
            } else {
                dp[i][j] = Math.max(dp[i - 1][j], dp[i][j - 1]);
            }
        }
    }

    return dp[m][n];
}

longestCommonSubsequence('ABCDGH', 'AEDFHR'); // 3 (ADH)

// Applications: Git diff, DNA sequence alignment, plagiarism detection

```

## 5. Constraint Satisfaction Problems (CSP)

**Question:** Find an assignment that satisfies all constraints

```

function solveScheduling(meetings, rooms) {
    // Meetings: [name, duration, participants]
    // Rooms: [name, capacity]
    // Constraints: No room conflicts, enough capacity

    function isValid(assignment, meeting, room, time) {
        // Check capacity
        if (room.capacity < meeting.participants) return false;

        // Check room availability
    }
}

```

```

for (const [m, r, t] of assignment) {
  if (r === room.name) {
    // Time conflict?
    if (!(time + meeting.duration <= t || t + m.duration <= time)) {
      return false;
    }
  }
}

return true;
}

function backtrack(assignment, meetingIndex) {
  if (meetingIndex === meetings.length) {
    return assignment; // Found valid assignment!
  }

  const meeting = meetings[meetingIndex];

  for (const room of rooms) {
    for (let time = 9; time < 17; time++) {
      if (isValid(assignment, meeting, room, time)) {
        assignment.push([meeting, room.name, time]);
        const result = backtrack(assignment, meetingIndex + 1);
        if (result) return result;
        assignment.pop(); // Backtrack
      }
    }
  }

  return null; // No solution
}

return backtrack([], 0);
}

const meetings = [
  { name: 'Standup', duration: 1, participants: 5 },
  { name: 'Pairing', duration: 2, participants: 2 },
  { name: 'Planning', duration: 3, participants: 8 }
];

const rooms = [
  { name: 'Conference A', capacity: 10 },
  { name: 'Conference B', capacity: 6 }
];

// Finds valid schedule respecting all constraints

```

## 6. Gradient Descent: Continuous Optimization

**Problem:** Minimize a smooth function by following the slope

```

function gradientDescent(startX, learningRate = 0.01, iterations = 1000) {
  // Objective: Minimize f(x) = (x - 5)^2 (has minimum at x=5)

  const f = x => Math.pow(x - 5, 2);
  const df = x => 2 * (x - 5); // Derivative

  let x = startX;
  const history = [{ x, fx: f(x), iter: 0 }];

  for (let i = 1; i <= iterations; i++) {
    const gradient = df(x);
    x = x - learningRate * gradient; // Move in opposite direction of gradient

    if (i % 100 === 0) {
      history.push({ x: x.toFixed(4), fx: f(x).toFixed(4), iter: i });
    }
  }

  return history;
}

gradientDescent(0);
/* [
  { x: 0, fx: 25, iter: 0 },
  { x: '3.62', fx: '1.91', iter: 100 },
  { x: '4.74', fx: '0.07', iter: 200 },
  { x: '4.96', fx: '0.0018', iter: 300 },
  ...
  → Converges to x=5, f(x)=0
]
// This is how neural networks learn! */

```

### Real Example: Price Optimization

```

function optimizePrice() {
  // Revenue = price × demand
  // Demand = 1000 - 20 × price (demand decreases with price)
  // Revenue(p) = p × (1000 - 20p) = 1000p - 20p²

  // Maximum at: dRevenue/dp = 0 → 1000 - 40p = 0 → p = 25

  const revenue = p => p * (1000 - 20 * p);
  const dRevenue = p => 1000 - 40 * p;

  let price = 0;
  for (let iter = 0; iter < 100; iter++) {
    const gradient = dRevenue(price);
    price += 0.01 * gradient; // Gradient ascent (maximize)
  }

  return {

```

```

        optimalPrice: price.toFixed(2),
        maxRevenue: revenue(price).toFixed(0)
    );
}

console.log(optimizePrice());
// { optimalPrice: '25.00', maxRevenue: '12500' }
// Price at $25 maximizes revenue (not always highest price!)

```

## Software Engineering Connections

### 1. Resource Allocation

```

function allocateDevResources(sprints, projects, developers) {
    // Problem: Assign developers to projects to maximize delivered value
    // Constraints: Each dev works one project, limited time

    const maxValue = {};

    function backtrack(devIndex, assignments, value) {
        if (devIndex === developers.length) {
            maxValue.current = Math.max(maxValue.current || 0, value);
            maxValue.assignment = [...assignments];
            return;
        }

        for (let project = 0; project < projects.length; project++) {
            const dev = developers[devIndex];
            const proj = projects[project];
            const capacity = dev.productivity * sprints;
            const contribution = Math.min(capacity, proj.work) * proj.priority;

            assignments.push({ dev: dev.name, project: proj.name });
            backtrack(devIndex + 1, assignments, value + contribution);
            assignments.pop();
        }
    }

    backtrack(0, [], 0);
    return maxValue;
}

const devs = [
    { name: 'Alice', productivity: 10 },
    { name: 'Bob', productivity: 8 },
    { name: 'Carol', productivity: 12 }
];

const projects = [
    { name: 'Feature A', work: 40, priority: 5 },

```

```

    { name: 'Feature B', work: 30, priority: 4 },
    { name: 'Technical Debt', work: 20, priority: 3 }
];

```

// Finds optimal assignment maximizing value delivered

## 2. Database Query Optimization

```

function optimizeQueryPlan(query, tables) {
  // Different join orders = different costs
  // Goal: Minimize total I/O operations

  function estimateCost(joinOrder) {
    let cost = 0;
    let resultSize = tables[joinOrder[0]].rows;

    for (let i = 1; i < joinOrder.length; i++) {
      const nextTable = tables[joinOrder[i]];
      cost += resultSize * nextTable.rows; // Cross product cost
      resultSize = Math.floor(resultSize * nextTable.selectivity);
    }

    return cost;
  }

  // Try all join orders (factorial complexity, but small in practice)
  const orders = [];
  function permute(arr, start = 0) {
    if (start === arr.length - 1) {
      orders.push([...arr]);
    } else {
      for (let i = start; i < arr.length; i++) {
        [arr[start], arr[i]] = [arr[i], arr[start]];
        permute(arr, start + 1);
        [arr[start], arr[i]] = [arr[i], arr[start]];
      }
    }
  }
}

permute([...Array(tables.length).keys()]);

return orders
  .map(order => ({
    order: order.map(i => tables[i].name),
    cost: estimateCost(order)
  }))
  .sort((a, b) => a.cost - b.cost)[0];
}

const tables = [
  { name: 'Users', rows: 1000, selectivity: 0.1 },

```

```

{ name: 'Orders', rows: 10000, selectivity: 0.2 },
{ name: 'Items', rows: 100000, selectivity: 0.05 }
];

// Finds optimal join order for query execution

```

### 3. Caching Strategy

```

function cachingStrategy(requests, cacheSize) {
    // Problem: Which items to keep in cache?
    // Constraints: Limited cache space
    // Objective: Minimize cache misses

    // Greedy approach: LRU (Least Recently Used)
    class LRUCache {
        constructor(size) {
            this.size = size;
            this.cache = new Map();
        }

        get(key) {
            if (this.cache.has(key)) {
                // Move to front (most recently used)
                const value = this.cache.get(key);
                this.cache.delete(key);
                this.cache.set(key, value);
                return value;
            }
            return null;
        }

        set(key, value) {
            if (this.cache.has(key)) {
                this.cache.delete(key);
            } else if (this.cache.size >= this.size) {
                // Evict least recently used
                const oldestKey = this.cache.keys().next().value;
                this.cache.delete(oldestKey);
            }
            this.cache.set(key, value);
        }
    }

    const cache = new LRUCache(cacheSize);
    let hits = 0, misses = 0;

    for (const key of requests) {
        if (!cache.get(key)) {
            misses++;
            cache.set(key, `value_${key}`);
        } else {

```

```

        hits++;
    }
}

return { hits, misses, hitRate: (hits / (hits + misses) * 100).toFixed(1) + '%' };
}

// Simulate: Access patterns [1, 2, 3, 2, 1, 4, 2, 3, ...]
const requests = [1, 2, 3, 2, 1, 4, 2, 3, 4, 1, 5, 2];
console.log(cachingStrategy(requests, 3));
// { hits: 6, misses: 6, hitRate: '50%' }

```

## 4. Load Balancing

```

function loadBalance(requests, servers) {
    // Objective: Minimize max server load (balance)

    function getLoad(assignment) {
        const loads = servers.map(s => s.capacity);
        assignment.forEach((serverIndex, requestIndex) => {
            loads[serverIndex] -= requests[requestIndex].weight;
        });
        return Math.min(...loads); // Bottleneck (max load)
    }

    // Greedy: Always assign to least-loaded server
    const assignment = Array(requests.length).fill(0);

    for (let i = 0; i < requests.length; i++) {
        let minLoadIdx = 0;
        let minLoad = servers[0].capacity;

        for (let s = 0; s < servers.length; s++) {
            const currentLoad = servers[s].capacity -
                assignment.filter(idx => idx === s).reduce((sum, _) => sum +
                    requests[_].weight, 0);

            if (currentLoad < minLoad) {
                minLoad = currentLoad;
                minLoadIdx = s;
            }
        }

        assignment[i] = minLoadIdx;
    }

    return assignment;
}

const requests = [
    { id: 1, weight: 10 },

```

```

{ id: 2, weight: 15 },
{ id: 3, weight: 12 },
{ id: 4, weight: 20 }
];

const servers = [
  { name: 'Server 1', capacity: 50 },
  { name: 'Server 2', capacity: 50 }
];

// Distributes load evenly

```

## 5. Algorithm Selection

```

function chooseAlgorithm(dataSize) {
  const algorithms = [
    { name: 'Bubble Sort', complexity: 'O(n2)', best: 100 },
    { name: 'Merge Sort', complexity: 'O(n log n)', best: 10000 },
    { name: 'Quick Sort', complexity: 'O(n log n)', best: 5000 }
  ];

  // Choose algorithm optimized for data size
  const choice = {
    small: algorithms[0],      // n2 is fine for small data
    medium: algorithms[2],     // Quick sort good average
    large: algorithms[1]       // Merge sort guaranteed O(n log n)
  };

  if (dataSize < 1000) return choice.small;
  if (dataSize < 100000) return choice.medium;
  return choice.large;
}

// Optimization = Choosing right tool for problem size

```

## Common Misconceptions

### ✗ "Optimization always means best possible"

**Wrong:** Optimal usually means "best under constraints"

```

// Global optimum: Hire infinite developers, ship all features, spend nothing
// Realistic optimum: Ship most valuable features within budget and team size

```

### ✗ "We should always optimize for performance"

**Sometimes cost/simplicity matters more:**

```

// Task: Sort 100 items
// Option A: Quick sort O(n log n) - complex, 1ms
// Option B: Bubble sort O(n2) - simple, 10ms

// For 100 items: B is better (10ms is fine, code is maintainable)
// For 1M items: A is necessary (B would take hours)

// Optimization must consider actual constraints

```

## X "Greedy always finds the best solution"

**Counter-example:** Activity selection with weights

```

// Greedy (by earliest end time): [Activity A, Activity C] = value 50
// Optimal: [Activity B, Activity D] = value 70

// Greedy is fast but often suboptimal

```

## X "More iterations always finds better optimum"

**Risk:** Overfitting, local minima

```

// Gradient descent can get stuck in local minima
// Iterations without improvement = likely found local optimum
// More iterations won't escape without changing approach

```

## Practical Mini-Exercises

- ▶ **Exercise 1: Knapsack Problem** (Click to expand)
- ▶ **Exercise 2: Server Capacity Planning** (Click to expand)
- ▶ **Exercise 3: Task Prioritization** (Click to expand)

## Summary Cheat Sheet

```

// OPTIMIZATION FRAMEWORK
1. Identify objective (minimize/maximize)
2. List all constraints
3. Define decision variables
4. Choose algorithm:
   - Linear programming: Convex, deterministic
   - Greedy: Fast, often suboptimal
   - Dynamic programming: Overlapping subproblems
   - CSP/Backtracking: Discrete constraints
   - Gradient descent: Continuous smooth functions

// COMMON PATTERNS
Knapsack:      Dynamic programming
Scheduling:     Greedy or backtracking

```

```
Resource alloc: Linear program or search
Shortest path: Dijkstra (greedy-like)
ML training: Gradient descent
```

#### When to optimize:

- Constraints are tight
- Trade-offs matter (cost vs performance)
- Problem repeats (worth effort)

#### When NOT to optimize:

- Good enough solves problem
- Premature optimization (measure first)
- Complexity outweighs benefit

---

## Next Steps

**You've completed:** Optimization foundations

**Up next:** [15. Information Theory](#) - Entropy, compression, hashing, why information matters

#### Before moving on:

```
// Challenge: Minimize latency given: $100 budget
// - t2.micro: $10/month, 50ms latency
// - t2.small: $25/month, 30ms latency
// - t2.medium: $40/month, 10ms latency
// How many of each to achieve <20ms average latency at minimum cost?

function latencyOptimization() {
    // Your solution
}
```