

18. Risk, Uncertainty & When Math Breaks Down

Phase 6: Real-World Applications & Limitations

⌚ ~45 minutes | 🕵️ Judgment Beyond Equations | ⚠️ Critical Final Perspective

What Problem This Solves

You've built models that:

- Predict latency (but actual > predicted)
- Estimate costs (but real costs blow up)
- Forecast growth (but it plateaus unexpectedly)
- Optimize for throughput (but causes mysterious failures)
- Assume normal distribution (but data has fat tails)

Without model skepticism, you treat math as gospel. "The model says X, so we should do X." You miss assumptions, ignore context, and are blindsided when reality diverges from equations.

With model skepticism, you treat math as a *tool*, not truth. You stress-test assumptions, perform sensitivity analysis, and maintain humility about what models can and cannot tell you.

Intuition & Mental Model

The Core Insight: All Models Are Wrong (Some Are Useful)

Model: Simplified representation of reality
"Assume normal distribution..."
"Assume independence..."
"Assume linear relationship..."

Reality: Complex, nonlinear, context-dependent
Fat tails, correlations, regime changes
Black swans, feedback loops, emergence

Mental Model: The Map vs The Territory

MODEL (Map)	REALITY (Territory)
_____	
City	/ \
.	Complex
_____	Changing
Useful for navigation	Unpredictable
but missing details	The map is NOT the territory

Core Concepts

1. Hidden Assumptions in Models

```

function modelAssumptions() {
  // Example: Linear growth model
  function predictUsers(monthlyGrowth, months) {
    // ASSUMPTION: Constant growth rate
    return Array.from({ length: months }, (_, i) =>
      1000 * Math.pow(1 + monthlyGrowth, i)
    );
  }

  // Reality check: What assumptions?
  const assumptions = [
    'Market size is unlimited (false: saturation exists)',
    'Growth rate is constant (false: changes over time)',
    'No competition (false: they eat your growth)',
    'No churn (false: users leave)',
    'No seasonality (false: holidays, cycles)',
    'Economic conditions stable (false: recessions happen)'
  ];

  // Model says: 1000 → 2000 → 4000 → 8000 users (exponential)
  // Reality: 1000 → 1800 → 2500 → 2800 users (logistic, plateau)

  return {
    modelPrediction: predictUsers(0.2, 12),
    realityCheck: 'Verify each assumption before trusting output'
  };
}

```

Exercise in Assumptions:

```

function checkAssumptions(model, data) {
  const checks = {
    independence: 'Are observations truly independent?',
    stationarity: 'Does distribution stay constant over time?',
    normality: 'Is Gaussian a good fit, or are there fat tails?',
    linearity: 'Is relationship actually linear?',
    noOutliers: 'Are extreme values possible? (Black swans)',
    completeness: 'Does data capture all relevant factors?'
  };

  // Example: Assuming independent coin flips
  // If flips come from same coin with wear: NOT independent

  return checks;
}

```

2. Sensitivity Analysis: What If Assumptions Are Wrong?

```

function sensitivityAnalysis(baseCase, parameters) {
  // Test: How much does output change if parameters vary?

  function calculateNPV(initialCost, annualRevenue, years, discountRate) {
    let npv = -initialCost;
    for (let year = 1; year <= years; year++) {
      npv += annualRevenue / Math.pow(1 + discountRate, year);
    }
    return npv;
  }

  // Base case: $100k cost, $40k/year revenue, 5 years, 10% discount
  const base = calculateNPV(100000, 40000, 5, 0.10);

  // Sensitivity: What if revenue is 20% lower?
  const revenueDown = calculateNPV(100000, 32000, 5, 0.10);

  // What if discount rate is 15% instead?
  const higherDiscount = calculateNPV(100000, 40000, 5, 0.15);

  // What if both?
  const bothWorse = calculateNPV(100000, 32000, 5, 0.15);

  return {
    base: Math.round(base),
    revenueDown20: Math.round(revenueDown),
    discountUp5: Math.round(higherDiscount),
    bothWorse: Math.round(bothWorse),
    insight: bothWorse < 0
      ? 'Project becomes NEGATIVE if both assumptions off by small amounts!'
      : 'Still profitable in pessimistic scenarios'
  };
}

console.log(sensitivityAnalysis());
/* {
  base: 51633,
  revenueDown20: 21306,
  discountUp5: 34194,
  bothWorse: 6955,
  insight: '...'
}
// Small assumption changes → Big outcome changes! */

```

Monte Carlo Simulation: Test thousands of scenarios

```

function monteCarloSensitivity(iterations = 10000) {
  // Assume revenue and discount rate are uncertain (ranges)

  const results = [];

```

```

for (let i = 0; i < iterations; i++) {
  // Random draw from distributions
  const revenue = 30000 + Math.random() * 20000; // $30k-$50k
  const discount = 0.08 + Math.random() * 0.07; // 8%-15%
  const years = Math.floor(3 + Math.random() * 4); // 3-6 years

  const npv = calculateNPV(100000, revenue, years, discount);
  results.push(npv);
}

// Analyze distribution
results.sort((a, b) => a - b);

return {
  min: Math.round(results[0]),
  p10: Math.round(results[Math.floor(0.10 * iterations)]),
  median: Math.round(results[Math.floor(0.50 * iterations)]),
  p90: Math.round(results[Math.floor(0.90 * iterations)]),
  max: Math.round(results[iterations - 1]),
  probNegative: (results.filter(x => x < 0).length / iterations * 100).toFixed(1)
+ '%'
};
}

function calculateNPV(cost, revenue, years, discount) {
  let npv = -cost;
  for (let y = 1; y <= years; y++) {
    npv += revenue / Math.pow(1 + discount, y);
  }
  return npv;
}

console.log(monteCarloSensitivity());
/* {
  min: -30123,
  p10: 5432,
  median: 42000,
  p90: 78654,
  max: 120345,
  probNegative: '12.3%'
}
// 12% chance of loss even though "expected" is positive */

```

3. Black Swans: Fat-Tailed Distributions

```

function blackSwanRisk() {
  // Normal distribution: Rare events are EXTREMELY rare
  // Real world: Rare events happen more often (fat tails)

  function normalSample(mean, stddev) {
    // Box-Muller transform

```

```

const u1 = Math.random();
const u2 = Math.random();
const z = Math.sqrt(-2 * Math.log(u1)) * Math.cos(2 * Math.PI * u2);
return mean + z * stddev;
}

// Simulate: Stock returns (assume normal)
const normalReturns = Array.from({ length: 1000 }, () => normalSample(0, 1));

// How many returns > 3 std deviations?
const extremeNormal = normalReturns.filter(x => Math.abs(x) > 3).length;

// Normal:  $P(|x| > 3\sigma) = 0.27\%$  → Expect 2.7 out of 1000

// Real market data: 10x-100x more common!
// 2008 crash: 25 $\sigma$  event (should happen once per universe lifetime)
// But it happened.

return {
  normalPrediction: '0.27% of events > 3 $\sigma$ ',
  normalCount: extremeNormal,
  realMarket: '~3-5% of events > 3 $\sigma$  (10x more!)',
  implication: 'Models underestimate tail risk dramatically'
};
}

```

Power Law vs Normal:

```

function compareTails() {
  // Normal: Thin tails (events drop off exponentially)
  // Power law: Fat tails (events drop off slowly)

  function normalTail(x) {
    return Math.exp(-x * x / 2);
  }

  function powerLawTail(x, alpha = 2) {
    return Math.pow(x, -alpha);
  }

  // At x=3:
  console.log('Normal at 3 $\sigma$ ', normalTail(3));      // 0.011 (very rare)
  console.log('Power law at 3:', powerLawTail(3)); // 0.111 (10x more likely)

  // At x=5:
  console.log('Normal at 5 $\sigma$ ', normalTail(5));      // 0.000001 (never)
  console.log('Power law at 5:', powerLawTail(5)); // 0.04 (still happens!)

  return {
    realWorldFatTails: [
      'Wealth distribution (1% owns 50%)',

```

```

        'Website traffic (few sites get all traffic)',
        'Network connections (hubs exist)',
        'Natural disasters (rare but catastrophic)',
        'Cyberattacks (occasional huge breaches)'
    ]
};

}

```

4. Feedback Loops: When Models Create Reality

```

function feedbackLoops() {
    // Your model affects reality, which invalidates the model

    // Example 1: Traffic prediction
    function trafficModel() {
        // Model: Route A is fastest
        // Everyone uses Route A → Route A becomes congested
        // Now Route B is fastest → Model wrong!
    }

    // Example 2: HFT algorithms
    function highFrequencyTrading() {
        // Model: "If price drops 1%, buy"
        // Everyone uses same model → All buy at same time
        // Price jumps up → Flash crash / rally
        // Model becomes self-fulfilling or self-defeating
    }

    // Example 3: Server load balancing
    function loadBalancingFeedback() {
        // Model: "Route to least-loaded server"
        // Server A becomes least-loaded → Gets all requests
        // Server A becomes most-loaded → Model redirects all traffic away
        // Oscillations!
    }

    return {
        lesson: 'Models that affect behavior create feedback loops',
        solution: 'Add randomization, hysteresis, or second-order thinking'
    };
}

```

5. Goodhart's Law: When a Measure Becomes a Target

```

function goodhartsLaw() {
    // "When a measure becomes a target, it ceases to be a good measure"

    const examples = [
        {

```

```

        metric: 'Lines of code',
        intent: 'Measure productivity',
        gamed: 'Developers write verbose, unnecessary code',
        failure: 'More code ≠ more value'
    },
    {
        metric: 'Bug count',
        intent: 'Track quality',
        gamed: 'Developers split 1 bug into 10 tickets or mark as "won't fix"',
        failure: 'Gaming the number, not fixing bugs'
    },
    {
        metric: 'Test coverage',
        intent: 'Ensure testing',
        gamed: 'Write tests that don't assert anything (just hit lines)',
        failure: 'Coverage ↑, quality unchanged'
    },
    {
        metric: 'Uptime SLA',
        intent: 'Ensure reliability',
        gamed: 'Label outages as "maintenance" or turn off monitoring during incidents',
        failure: '99.9% uptime on paper, but users experience downtime'
    }
];

return {
    examples,
    solution: 'Use multiple metrics, qualitative assessment, and rotate metrics'
};
}

```

6. Model Degradation: Concept Drift

```

function conceptDrift() {
    // Problem: Model trained on past data becomes less accurate over time

    function trainModel(historicalData) {
        // Train spam filter on emails from 2020
        return {
            accuracy: 0.95,
            trainingYear: 2020,
            features: ['Nigerian prince', 'lottery', 'click here']
        };
    }

    function applyModel2024(model) {
        // Spam in 2024: AI-generated phishing, personalized attacks
        // Model's features no longer discriminate
        return {
            accuracy: 0.70, // Degraded!

```

```

        reason: 'Spammers adapted, legitimate emails changed tone',
        solution: 'Retrain regularly, detect drift, online learning'
    );
}

return {
    lesson: 'Data distributions change over time',
    examples: [
        'User behavior shifts (pandemic, trends)',
        'Adversarial adaptation (spam, fraud)',
        'Market regime changes (recession → boom)',
        'Product evolution (new features change usage)'
    ]
};
}

```

Software Engineering Connections

1. Capacity Planning Under Uncertainty

```

function capacityPlanning() {
    // Question: How many servers do we need?

    // Model: Forecast traffic, compute required capacity
    function naiveModel(expectedTraffic, safetyMargin = 1.2) {
        const serversNeeded = Math.ceil(expectedTraffic / 1000 * safetyMargin);
        return serversNeeded;
    }

    // Reality check: What if wrong?
    function robustPlanning(expectedTraffic) {
        const scenarios = [
            { name: 'Base case', traffic: expectedTraffic, prob: 0.5 },
            { name: '2x spike', traffic: expectedTraffic * 2, prob: 0.3 },
            { name: '5x spike (viral)', traffic: expectedTraffic * 5, prob: 0.1 },
            { name: 'Attack (10x)', traffic: expectedTraffic * 10, prob: 0.1 }
        ];

        // Expected value:  $\sum p(\text{scenario}) \times \text{cost}(\text{scenario})$ 
        const expectedServers = scenarios.reduce((sum, s) =>
            sum + s.prob * Math.ceil(s.traffic / 1000), 0
        );

        // But also: What's worst-case you can tolerate?
        const worstCaseServers = Math.ceil(scenarios[scenarios.length - 1].traffic /
1000);

        return {
            naive: naiveModel(expectedTraffic),
            expected: Math.round(expectedServers),

```

```

        worstCase: worstCaseServers,
        recommendation: 'Autoscale between expected and worst-case'
    };
}

return robustPlanning(100000);
/* {
    naive: 120,
    expected: 245,
    worstCase: 1000,
    recommendation: 'Autoscale between 245 and 1000'
} */
}

```

2. SLA Design: When Percentiles Lie

```

function percentileTrap() {
    // P99 latency = 200ms looks great
    // But: User makes 100 requests per page load

    function userExperience(p99_latency, requestsPerPage) {
        // Probability of at least one slow request
        const probAllFast = Math.pow(0.99, requestsPerPage); // All hit P99
        const probOneSlow = 1 - probAllFast;

        return {
            p99Latency: p99_latency + 'ms',
            requestsPerPage,
            probFastPage: (probAllFast * 100).toFixed(1) + '%',
            probSlowPage: (probOneSlow * 100).toFixed(1) + '%',
            insight: probOneSlow > 0.5
                ? 'MAJORITY of page loads will be slow!'
                : 'Most page loads will be fast'
        };
    }

    console.log(userExperience(200, 100));
    /* {
        p99Latency: '200ms',
        requestsPerPage: 100,
        probFastPage: '36.6%',
        probSlowPage: '63.4%',
        insight: 'MAJORITY of page loads will be slow!'
    }
    // Even with P99=200ms, most USERS experience slowness! */

    return {
        lesson: 'P99 per-request ≠ P99 per-user experience',
        solution: 'Measure end-to-end latency (full page load)'
    }
}

```

```
};  
}
```

3. Cost Models: Hidden Variables

```
function hiddenCostVariables() {  
    // Model: Cloud cost = instances × hours × rate  
  
    function naiveCostModel(instances, hoursPerMonth, ratePerHour) {  
        return instances * hoursPerMonth * ratePerHour;  
    }  
  
    // What model IGNORES:  
    const hiddenCosts = {  
        dataTransfer: 'Outbound bandwidth charges',  
        iops: 'Disk operations charged separately',  
        snapshots: 'Backup storage accumulates',  
        reservedInstances: 'Upfront commitment reduces hourly',  
        spotInterruptions: 'Spot instances get killed → need fallback',  
        multiAZ: 'Cross-AZ traffic costs extra',  
        supportPlan: 'AWS Support: 10% of bill',  
        engineeringTime: 'Developer hours managing infrastructure'  
    };  
  
    // Naive: $100/month  
    // Reality: $100 × 1.5 (transfer) × 1.1 (support) + $50 (snapshots) = $215  
  
    return {  
        naiveEstimate: naiveCostModel(10, 720, 0.013), // ~$93  
        realityMultiplier: '1.5-3x',  
        lesson: 'Always add hidden cost buffer (50-100%)'  
    };  
}
```

Common Misconceptions

✗ "If the math says so, it must be right"

Math is only as good as assumptions:

```
// Model: "ROI = 200%"  
// Assumption 1: Market size unlimited (false)  
// Assumption 2: No competition (false)  
// Assumption 3: Conversion rate stays 5% (false, it drops)  
// Reality: ROI = 50% (still good, but not 200%)
```

✗ "We can model everything"

Some systems are fundamentally unpredictable:

```
// Chaotic systems: Tiny input differences → huge output differences  
// Examples: Weather (>2 weeks), stock market (short-term), viral content  
  
// No amount of data or compute will predict which tweet goes viral
```

✗ "Historical data predicts the future"

Until it doesn't:

```
// Housing prices only go up (until 2008)  
// VIX stays low (until it spikes 10x in a day)  
// Our system never crashes (until it does)  
  
// Survivorship bias: You only see data from systems that survived
```

✗ "Average is representative"

Averages hide distributions:

```
// Average salary: $100k (sounds great!)  
// Reality: $50k (50%), $80k (40%), $500k (10%) → Median $65k  
// Average response time: 100ms (sounds fine!)  
// Reality: P50=50ms, P99=2000ms (some users suffer)
```

Practical Mini-Exercises

- ▶ **Exercise 1: Assumption Audit** (Click to expand)
- ▶ **Exercise 2: Sensitivity Test** (Click to expand)

Summary Cheat Sheet

```
// CORE PRINCIPLES  
1. All models are simplifications (useful but incomplete)  
2. Assumptions ALWAYS exist (make them explicit)  
3. Test sensitivity (what if assumptions wrong?)  
4. Distributions matter (mean ≠ median, beware fat tails)  
5. Feedback loops exist (model affects reality)  
6. Models degrade (retrain, monitor drift)  
7. Measure gaming (Goodhart's Law)  
  
// RED FLAGS  
- "The model is 99% accurate" (on what data? In what context?)  
- "Expected value is positive" (what about variance? Tail risk?)  
- "Assume normal distribution" (test for fat tails)  
- "Historical trend continues" (until it doesn't)  
- "This metric perfectly captures X" (metrics get gamed)  
  
// BEST PRACTICES
```

- Document assumptions explicitly
- Perform sensitivity analysis (vary parameters ±20-50%)
- Simulate scenarios (Monte Carlo)
- Monitor model performance (retrain when accuracy drops)
- Use multiple metrics (avoid Goodhart's Law)
- Add buffers (reality worse than model)
- Stay humble (math ≠ reality)

Conclusion: Math as a Tool, Not Truth

You've completed **18 topics across 6 phases**:

Phase 1-2: Systems & Performance

- Discrete math, graphs, boolean algebra, Big-O, recursion, probability

Phase 3: Statistics

- Descriptive stats, inferential stats, data distributions

Phase 4: Finance & Decision Making

- Financial math, exponential growth, expected value

Phase 5: Modern Software Math

- Linear algebra, optimization, information theory

Phase 6: Real-World Limitations

- Numerical methods, randomized algorithms, **risk & uncertainty**

Final Wisdom

Mathematics gives you:

- ✓ **Structure** for thinking about problems
- ✓ **Language** for precise communication
- ✓ **Tools** for quantifying trade-offs
- ✓ **Models** for exploring scenarios

Mathematics does NOT give you:

- ✗ **Certainty** about the future
- ✗ **Perfect** predictions
- ✗ **Complete** representations of reality
- ✗ **Immunity** to unforeseen events

Use math as a flashlight, not a GPS.

It illuminates the path, but you still need judgment to navigate.

Where to Go From Here

Continue practicing:

1. Apply these concepts to your daily work

2. Question assumptions in models you encounter
3. Run sensitivity analyses before big decisions
4. Build intuition through toy problems
5. Read real-world case studies (2008 financial crisis, AWS outages, etc.)

Resources for deeper study:

- Fooled by Randomness (Nassim Taleb) - Black swans, fat tails
- Thinking in Systems (Donella Meadows) - Feedback loops, emergence
- The Signal and the Noise (Nate Silver) - Forecasting limitations
- How to Measure Anything (Douglas Hubbard) - Quantifying uncertainty

Remember:

"All models are wrong, but some are useful." — George Box

Use math wisely. Stay humble. Question everything.

Thank You

Congratulations on completing the **Applied Math for Software Engineers** curriculum!

You now have:

- 18 foundational topics
- Practical code examples in JavaScript/TypeScript
- Real-world software engineering connections
- Mental models for thinking about uncertainty
- Healthy skepticism about models

Go build something amazing (and question your assumptions along the way).

Applied Math for Software Engineers • Phase 6 • [Previous: Randomized Algorithms](#) |  **Curriculum Complete!**