

Subqueries, CTEs, and Window Functions: Advanced Query Composition

Subqueries: Queries Within Queries

A subquery is a SELECT statement nested inside another query. They let you build complex logic in stages.

Types of Subqueries

1. **Scalar subqueries:** Return a single value
2. **Row subqueries:** Return a single row (uncommon)
3. **Table subqueries:** Return multiple rows/columns
4. **Correlated subqueries:** Reference outer query

Scalar Subqueries: Single Value

Use case: "For each user, show how many orders they have"

```
SELECT
    name,
    (SELECT COUNT(*) FROM orders WHERE user_id = users.id) AS order_count
FROM users;
```

What happens:

- For each user, the subquery runs
- It counts orders for that specific user
- Returns a single number

Performance note: The subquery runs **once per user**. If you have 10,000 users, it runs 10,000 times.

This is called a **correlated subquery** because it references `users.id` from the outer query.

Table Subqueries: Multiple Rows

Use case: "Show orders placed by users in New York"

```
SELECT *
FROM orders
WHERE user_id IN (
    SELECT id FROM users WHERE city = 'New York'
);
```

What happens:

- Inner query returns a set of user IDs
- Outer query filters orders against that set

This is uncorrelated: The subquery doesn't reference the outer query.

Subqueries in FROM: Derived Tables

You can treat a subquery as a table:

```
SELECT avg_price, COUNT(*) AS product_count
FROM (
    SELECT category_id, AVG(price) AS avg_price
    FROM products
    GROUP BY category_id
) AS category_avg
GROUP BY avg_price;
```

What happens:

1. Inner query computes average price per category
2. Outer query treats that result as a table
3. Groups by those averages

Note: You must alias subqueries in `FROM (AS category_avg)`.

Correlated vs Uncorrelated Subqueries

Uncorrelated Subqueries

Independent of outer query:

```
SELECT name
FROM products
WHERE category_id IN (
    SELECT id FROM categories WHERE active = true
);
```

Execution:

1. Inner query runs once
2. Outer query uses the result

Performance: Usually fast (inner query runs once).

Correlated Subqueries

Depends on outer query:

```
SELECT name
FROM products p
WHERE price > (
    SELECT AVG(price) FROM products WHERE category_id = p.category_id
);
```

Execution:

1. For each product, inner query runs
2. Computes average for that product's category
3. Compares product price to that average

Performance: Can be slow (inner query runs once per outer row).

When the optimizer saves you: Sometimes the database optimizes correlated subqueries into joins. Sometimes it doesn't. Use EXPLAIN to check.

Converting Correlated to Uncorrelated

Correlated (slow):

```
SELECT
    u.name,
    (SELECT COUNT(*) FROM orders WHERE user_id = u.id) AS order_count
FROM users u;
```

Uncorrelated with JOIN (faster):

```
SELECT
    u.name,
    COUNT(o.id) AS order_count
FROM users u
LEFT JOIN orders o ON u.id = o.user_id
GROUP BY u.id, u.name;
```

Why it's faster: Single pass through both tables instead of N queries.

CTEs (Common Table Expressions): WITH Clause

CTEs let you name subqueries for readability and reuse.

Basic Syntax

```
WITH high_value_orders AS (
    SELECT * FROM orders WHERE total > 1000
)
SELECT user_id, COUNT(*) AS big_order_count
FROM high_value_orders
GROUP BY user_id;
```

Equivalent to:

```
SELECT user_id, COUNT(*) AS big_order_count
FROM (
    SELECT * FROM orders WHERE total > 1000
) AS high_value_orders
GROUP BY user_id;
```

Why use CTEs?

- **Readability:** Named stages are clearer than nested subqueries
- **Reusability:** Reference the CTE multiple times in the main query
- **Maintainability:** Easier to modify one piece

Multiple CTEs

```
WITH
active_users AS (
    SELECT id, name FROM users WHERE active = true
),
recent_orders AS (
    SELECT user_id, total FROM orders WHERE created_at > NOW() - INTERVAL '30 days'
)
SELECT
    u.name,
    SUM(o.total) AS recent_revenue
FROM active_users u
JOIN recent_orders o ON u.id = o.user_id
GROUP BY u.id, u.name;
```

Each CTE is defined once, used as needed.

Recursive CTEs: For Hierarchical Data

Use case: Org chart (employee → manager hierarchy).

```
WITH RECURSIVE employee_hierarchy AS (
    -- Base case: top-level employees (no manager)
    SELECT id, name, manager_id, 1 AS level
    FROM employees
    WHERE manager_id IS NULL

    UNION ALL

    -- Recursive case: employees with managers
    SELECT e.id, e.name, e.manager_id, eh.level + 1
    FROM employees e
    JOIN employee_hierarchy eh ON e.manager_id = eh.id
)
SELECT * FROM employee_hierarchy
ORDER BY level, name;
```

How it works:

1. Base case seeds the recursion (top-level managers)
2. Recursive case joins employees to previous level
3. Continues until no more rows are added
4. Returns the full hierarchy

Common use cases:

- Org charts
- Bill of materials (parts and subparts)
- File systems, category trees
- Graph traversal

Warning: Recursive CTEs can infinite loop if your data has cycles. Add a depth limit:

```
WHERE eh.level < 10
```

The CTE "Optimization Fence" Myth

There's a widespread belief that CTEs in PostgreSQL create an "optimization fence" that prevents the optimizer from inlining them.

The Nuance

Before Postgres 12: CTEs were always materialized (computed once, stored temporarily). The optimizer couldn't push predicates into them.

Since Postgres 12: CTEs are inlined by default (treated like subqueries). The optimizer can optimize across them.

Exception: If a CTE is referenced multiple times, OR if you use `MATERIALIZED`, it's materialized.

Example: Inline vs Materialized

```
WITH recent_orders AS (
    SELECT * FROM orders WHERE created_at > NOW() - INTERVAL '7 days'
)
SELECT * FROM recent_orders WHERE user_id = 123;
```

Postgres 12+: The optimizer inlines this and executes:

```
SELECT * FROM orders
WHERE created_at > NOW() - INTERVAL '7 days'
AND user_id = 123;
```

It pushes the `user_id = 123` filter into the scan.

Forcing Materialization

If you want the old behavior (compute once, reuse):

```
WITH recent_orders AS MATERIALIZED (
    SELECT * FROM orders WHERE created_at > NOW() - INTERVAL '7 days'
)
SELECT * FROM recent_orders WHERE user_id = 123;
```

Now Postgres computes `recent_orders` once, stores it, then filters.

When to use `MATERIALIZED`:

- CTE is expensive and referenced multiple times
- You want to protect a carefully tuned subquery from being rewritten

When NOT to use `MATERIALIZED`:

- When you want the optimizer to push predicates down
- When the CTE is cheap or only used once

MySQL and Other Databases

- **MySQL 8.0+:** Has CTEs, but materialization behavior differs. Check docs.
- **SQL Server:** CTEs are inlined (like subqueries).
- **Oracle:** CTEs can be hinted with materialize/inline directives.

Window Functions: The Game Changer

Window functions let you compute aggregates **without collapsing rows**.

The Problem They Solve

Goal: Show each order and the user's total order count.

Without window functions (requires self-join):

```
SELECT
    o.id,
    o.total,
    user_counts.order_count
FROM orders o
JOIN (
    SELECT user_id, COUNT(*) AS order_count
    FROM orders
    GROUP BY user_id
) user_counts ON o.user_id = user_counts.user_id;
```

Messy.

With window functions:

```
SELECT
    id,
    total,
    COUNT(*) OVER (PARTITION BY user_id) AS order_count
FROM orders;
```

Magic: Each row keeps its identity, but you get aggregate info.

Window Function Syntax

```
<aggregate_function>() OVER (
    PARTITION BY <column>      -- Optional: groups (like GROUP BY, but non-collapsing)
    ORDER BY <column>          -- Optional: ordering within groups
    ROWS/RANGE <frame>        -- Optional: which rows to include in calculation
)
```

Key Window Functions

1. Aggregate Functions (COUNT, SUM, AVG, etc.)

```
SELECT
    name,
    salary,
    AVG(salary) OVER (PARTITION BY department_id) AS dept_avg_salary
FROM employees;
```

Result: Each employee sees their department's average salary, but rows aren't collapsed.

2. ROW_NUMBER(): Assign Row Numbers

```
SELECT
    name,
    salary,
    ROW_NUMBER() OVER (ORDER BY salary DESC) AS salary_rank
FROM employees;
```

Result:

name	salary	salary_rank
Alice	120000	1
Bob	110000	2
Charlie	100000	3

Use case: Ranking, pagination.

3. RANK() and DENSE_RANK(): Handle ties

RANK: Leaves gaps after ties.

```
SELECT
    name,
    score,
    RANK() OVER (ORDER BY score DESC) AS rank
FROM contestants;
```

Result:

name	score	rank
Alice	100	1
Bob	100	1 -- Tie
Charlie	90	3 -- Gap!

DENSE_RANK: No gaps.

name	score	dense_rank
Alice	100	1

```
Bob      | 100    | 1
Charlie | 90     | 2    -- No gap
```

4. LAG() and LEAD(): Access Previous/Next Rows

LAG: Access previous row.

```
SELECT
  date,
  revenue,
  LAG(revenue) OVER (ORDER BY date) AS prev_day_revenue
FROM daily_sales;
```

Result:

date	revenue	prev_day_revenue
2024-01-01	1000	NULL
2024-01-02	1200	1000
2024-01-03	1100	1200

LEAD: Access next row (same idea).

Use case: "Compare today's revenue to yesterday's."

5. FIRST_VALUE() and LAST_VALUE(): First/Last in Window

```
SELECT
  name,
  salary,
  FIRST_VALUE(name) OVER (PARTITION BY department_id ORDER BY salary DESC) AS
highest_paid
FROM employees;
```

Result: Each row shows who's the highest paid in their department.

PARTITION BY: Grouping Without Collapsing

```
SELECT
  product_id,
  sale_date,
  amount,
  SUM(amount) OVER (PARTITION BY product_id) AS total_product_sales
FROM sales;
```

What happens:

- Rows are logically grouped by `product_id`
- For each row, `total_product_sales` shows the sum for that product
- But all rows remain individual

Think of it as: "GROUP BY without the collapsing."

ORDER BY in Window Functions: Running Totals

```
SELECT
    sale_date,
    amount,
    SUM(amount) OVER (ORDER BY sale_date) AS running_total
FROM sales;
```

Result:

sale_date	amount	running_total
2024-01-01	100	100
2024-01-02	150	250
2024-01-03	200	450

Key insight: ORDER BY in OVER makes the aggregate cumulative.

Frame Clauses: ROWS and RANGE

By default, ORDER BY in a window function uses RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW (running total).

You can customize:

ROWS BETWEEN

Sliding window of 3 rows:

```
SELECT
    date,
    amount,
    AVG(amount) OVER (ORDER BY date ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) AS
    moving_avg_3day
FROM sales;
```

Result: 3-day moving average (current + 2 previous).

Frame options:

- ROWS BETWEEN 2 PRECEDING AND CURRENT ROW : Last 3 rows
- ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW : From start to current (running total)
- ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING : Current row + 1 before + 1 after

RANGE BETWEEN

RANGE operates on values, not physical rows.

Example: "Average salary within \$10k of current salary."

```

SELECT
    name,
    salary,
    AVG(salary) OVER (ORDER BY salary RANGE BETWEEN 10000 PRECEDING AND 10000
FOLLOWING) AS nearby_avg
FROM employees;

```

ROWS vs RANGE:

- **ROWS:** Physical row positions
- **RANGE:** Logical value ranges

Real-World Window Function Examples

Example 1: Top 3 Products per Category

```

SELECT *
FROM (
    SELECT
        category_id,
        product_name,
        sales,
        ROW_NUMBER() OVER (PARTITION BY category_id ORDER BY sales DESC) AS rank
    FROM products
) ranked
WHERE rank <= 3;

```

How it works:

1. Assign row numbers within each category, ordered by sales
2. Filter to top 3

Example 2: Month-over-Month Growth

```

SELECT
    month,
    revenue,
    LAG(revenue) OVER (ORDER BY month) AS prev_month,
    revenue - LAG(revenue) OVER (ORDER BY month) AS growth
FROM monthly_revenue;

```

Result:

month	revenue	prev_month	growth
2024-01	10000	NULL	NULL
2024-02	12000	10000	2000
2024-03	11000	12000	-1000

Example 3: Cumulative Sum (Running Total)

```
SELECT
  date,
  amount,
  SUM(amount) OVER (ORDER BY date) AS cumulative_revenue
FROM daily_sales;
```

Example 4: Latest Order Per User (Without Self-Join)

Remember this from the JOINs chapter?

Old way (self-join):

```
SELECT u.name, o.created_at
FROM users u
JOIN orders o ON u.id = o.user_id
WHERE o.created_at = (
  SELECT MAX(created_at) FROM orders WHERE user_id = u.id
);
```

Window function way:

```
SELECT name, created_at
FROM (
  SELECT
    u.name,
    o.created_at,
    ROW_NUMBER() OVER (PARTITION BY u.id ORDER BY o.created_at DESC) AS rn
  FROM users u
  JOIN orders o ON u.id = o.user_id
) sub
WHERE rn = 1;
```

Cleaner, often faster.

When to Use Subqueries vs CTEs vs Window Functions

Use Subqueries When:

- Simple, one-off logic
- Inline filtering (EXISTS, IN)
- Scalar computations

Use CTEs When:

- Complex multi-stage queries
- Readability matters (almost always)
- Recursive logic
- Reusing the same logic multiple times

Use Window Functions When:

- You need aggregates but want to keep row detail

- Ranking, row numbering
- Running totals, moving averages
- Comparing rows to previous/next rows

Avoid:

- **Deep nesting of subqueries:** Use CTEs for readability
- **Correlated subqueries in tight loops:** Convert to JOINs or window functions
- **Overusing MATERIALIZED:** Let the optimizer do its job

Performance Considerations

Subquery Performance

Good:

```
-- Uncorrelated, runs once
WHERE user_id IN (SELECT id FROM active_users)
```

Bad:

```
-- Correlated, runs per row
WHERE total > (SELECT AVG(total) FROM orders WHERE user_id = outer.user_id)
```

Fix: Use JOIN or window function.

CTE Performance

Good:

```
WITH expensive_calc AS (
  SELECT ... -- Complex aggregation
)
SELECT * FROM expensive_calc
UNION ALL
SELECT * FROM expensive_calc; -- Reuses result
```

Bad:

```
-- Forcing materialization when not needed
WITH trivial_filter AS MATERIALIZED (
  SELECT * FROM orders WHERE id = 123
)
SELECT * FROM trivial_filter;
```

Window Function Performance

Good:

```
-- Single pass
SELECT
```

```
id,  
SUM(amount) OVER (ORDER BY date),  
AVG(amount) OVER (ORDER BY date)  
FROM sales;
```

Bad (forces multiple sorts):

```
-- Different ORDER BY in each window = multiple sorts  
SELECT  
    id,  
    SUM(amount) OVER (ORDER BY date),  
    AVG(amount) OVER (ORDER BY product_id)  
FROM sales;
```

Fix: Group compatible window functions together.

Pro Tip: EXPLAIN Your CTEs and Window Functions

Always check the execution plan. The optimizer might surprise you (good or bad).

```
EXPLAIN ANALYZE  
WITH ...
```

Common Pitfalls

Pitfall 1: Forgetting PARTITION BY

Problem:

```
SELECT  
    user_id,  
    order_id,  
    ROW_NUMBER() OVER (ORDER BY created_at) AS order_num  
FROM orders;
```

Bug: Row numbers are global, not per user.

Fix:

```
ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY created_at)
```

Pitfall 2: Using LAST_VALUE Without Frame Clause

Problem:

```
SELECT  
    date,  
    amount,
```

```
LAST_VALUE(amount) OVER (ORDER BY date) AS last_amount  
FROM sales;
```

Bug: `last_amount` is just `amount` (current row is the last in the default frame).

Fix:

```
LAST_VALUE(amount) OVER (ORDER BY date ROWS BETWEEN UNBOUNDED PRECEDING AND  
UNBOUNDED FOLLOWING)
```

Or just use `FIRST_VALUE` with `DESC`.

Pitfall 3: Filtering on Window Functions in WHERE

Problem:

```
SELECT  
    name,  
    ROW_NUMBER() OVER (ORDER BY salary DESC) AS rank  
FROM employees  
WHERE rank <= 5; -- ERROR!
```

Why it fails: Window functions are evaluated after WHERE.

Fix: Use a subquery:

```
SELECT * FROM (  
    SELECT  
        name,  
        ROW_NUMBER() OVER (ORDER BY salary DESC) AS rank  
    FROM employees  
) sub  
WHERE rank <= 5;
```

Key Takeaways

1. **Subqueries let you nest logic.** Prefer uncorrelated over correlated.
2. **CTEs improve readability.** Use them for complex multi-stage queries.
3. **Recursive CTEs solve hierarchical problems.** Watch for infinite loops.
4. **Window functions keep row detail while computing aggregates.** Massive power.
5. **PARTITION BY is like GROUP BY without collapsing rows.**
6. **ORDER BY in window functions enables running totals and LAG/LEAD.**
7. **CTEs in Postgres 12+ are inlined unless materialized or referenced multiple times.**
8. **Always EXPLAIN complex queries** to see what the optimizer is doing.

Next up: GROUP BY and aggregations in depth—because that's where things get weird.