

# Transactions, Isolation, and Locking: Concurrency Without Panic

## What Transactions Actually Are

A transaction is a **unit of work** that's either fully completed or fully rolled back. No half-done states.

```
BEGIN;
  UPDATE accounts SET balance = balance - 100 WHERE id = 1;
  UPDATE accounts SET balance = balance + 100 WHERE id = 2;
COMMIT;
```

Either both **UPDATEs happen, or neither does**. This is **atomicity**.

## ACID Without the Buzzwords

### Atomicity: All or Nothing

**What it means:** A transaction's changes are indivisible.

**Example:**

```
BEGIN;
  INSERT INTO orders (user_id, total) VALUES (1, 100);
  UPDATE users SET total_spent = total_spent + 100 WHERE id = 1;
COMMIT;
```

If the **UPDATE** fails (e.g., constraint violation), the **INSERT** is rolled back too.

**No partial commits.**

### Consistency: Constraints Are Enforced

**What it means:** The database moves from one valid state to another.

**Example:**

- Before transaction:  $\text{SUM(balances)} = 10000$
- After transaction:  $\text{SUM(balances)} = 10000$  (money moved, not created)

Constraints (foreign keys, checks) are enforced at commit time.

### Isolation: Transactions Don't Interfere (Mostly)

**What it means:** Concurrent transactions appear to run serially (depending on isolation level).

**The problem:** Transaction A reads data while Transaction B is modifying it. What does A see?

**Answer:** Depends on the **isolation level** (more on this below).

### Durability: Committed Data Survives Crashes

**What it means:** Once COMMIT returns, the data is safe (even if the server crashes 1ms later).

**How:** Write-ahead log (WAL). Changes are written to disk before COMMIT returns.

**Trade-off:** Durability = disk I/O = slower commits.

## Isolation Levels: The Core Concept

SQL defines **four isolation levels** to balance consistency vs performance:

1. **READ UNCOMMITTED** (lowest isolation, highest concurrency)
2. **READ COMMITTED** (most common)
3. **REPEATABLE READ** (stricter)
4. **SERIALIZABLE** (strictest, lowest concurrency)

**Higher isolation = more safety, lower performance.**

### The Concurrency Anomalies

Different isolation levels protect against different anomalies:

1. **Dirty Read:** Reading uncommitted data from another transaction
2. **Non-Repeatable Read:** Same query returns different results within a transaction
3. **Phantom Read:** New rows appear in a range query
4. **Serialization Anomaly:** Interleaved execution produces results impossible in serial execution

**Isolation level chart:**

Level	Dirty Read	Non-Repeatable Read	Phantom Read	Serialization Anomaly
READ UNCOMMITTED	✗ Allowed	✗ Allowed	✗ Allowed	✗ Allowed
READ COMMITTED	✓ Prevented	✗ Allowed	✗ Allowed	✗ Allowed
REPEATABLE READ	✓ Prevented	✓ Prevented	✗ Allowed (Postgres: ✓)	✗ Allowed
SERIALIZABLE	✓ Prevented	✓ Prevented	✓ Prevented	✓ Prevented

## Isolation Level Examples

### READ UNCOMMITTED: The Wild West

**Not supported in Postgres.** MySQL supports it (don't use it).

**What happens:**

- Transaction A can see uncommitted changes from Transaction B
- If B rolls back, A saw data that **never existed**

**Use case:** Almost never. Maybe for reporting where perfect accuracy doesn't matter.

### READ COMMITTED: The Default

## **Postgres and MySQL default.**

**What it does:** You only see committed data. But the data can change between queries.

### **Example: Non-Repeatable Read**

```
Timeline:  
-----  
T1: BEGIN;  
T1: SELECT balance FROM accounts WHERE id = 1;  
    --> balance = 100  
  
T2: BEGIN;  
T2: UPDATE accounts SET balance = 200 WHERE id = 1;  
T2: COMMIT;  
  
T1: SELECT balance FROM accounts WHERE id = 1;  
    --> balance = 200 (changed!)  
T1: COMMIT;
```

**Result:** Same SELECT in T1 returns different values.

**When it's a problem:** If you're making decisions based on the first read (e.g., "if balance > 50, withdraw 50").

**When it's fine:** Most web apps. You're not making complex multi-step decisions.

## **REPEATABLE READ: Consistent Snapshot**

### **Postgres default for explicit transactions.**

**What it does:** You see a snapshot of the database as of the transaction start.

### **Example: Preventing Non-Repeatable Read**

```
Timeline:  
-----  
T1: BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
T1: SELECT balance FROM accounts WHERE id = 1;  
    --> balance = 100  
  
T2: BEGIN;  
T2: UPDATE accounts SET balance = 200 WHERE id = 1;  
T2: COMMIT;  
  
T1: SELECT balance FROM accounts WHERE id = 1;  
    --> balance = 100 (unchanged in T1's snapshot!)  
T1: COMMIT;
```

**Result:** T1 sees a consistent view. T2's changes are invisible to T1.

**When it's useful:** Multi-step reads that must be consistent (reports, analytics).

**Postgres-specific:** Prevents phantom reads too (stricter than SQL standard).

## **SERIALIZABLE: Total Isolation**

**What it does:** Guarantees that concurrent transactions produce the same result as if they ran one at a time.

**How:** Detects conflicts and aborts transactions.

**Example: Serialization Anomaly Prevented**

Scenario: Two users booking the last seat on a flight.

```
T1: BEGIN TRANSACTION ISOLATION LEVEL SERIALIZABLE;
T1: SELECT COUNT(*) FROM bookings WHERE flight_id = 123;
    --> 99 (one seat left)
```

```
T2: BEGIN TRANSACTION ISOLATION LEVEL SERIALIZABLE;
T2: SELECT COUNT(*) FROM bookings WHERE flight_id = 123;
    --> 99 (one seat left)
```

```
T1: INSERT INTO bookings (flight_id, user_id) VALUES (123, 1);
T1: COMMIT; --> Success
```

```
T2: INSERT INTO bookings (flight_id, user_id) VALUES (123, 2);
T2: COMMIT; --> ERROR: could not serialize access
```

**Result:** T2 is aborted. Only one user gets the seat.

**Cost:** Higher chance of aborts, lower concurrency.

**When to use:** Financial transactions, inventory management, anything where correctness > performance.

## MVCC: How Postgres Does Isolation

**MVCC (Multi-Version Concurrency Control):** Instead of locking rows, Postgres keeps multiple versions.

**How it works:**

1. Each transaction has a transaction ID (xid)
2. Each row has metadata: `xmin` (created by xid), `xmax` (deleted/updated by xid)
3. When you read a row, Postgres checks: "Is this version visible to my transaction?"

**Example:**

Initial state:

```
Row: id=1, balance=100, xmin=10, xmax=NULL (visible to all)
```

```
T1 (xid=20): UPDATE accounts SET balance = 200 WHERE id = 1;
New row: id=1, balance=200, xmin=20, xmax=NULL
Old row: id=1, balance=100, xmin=10, xmax=20
```

```
T2 (xid=21, started before T1 committed):
Sees old row (xmin=10, xmax=20, 21 > 20 but T1 hasn't committed yet)
```

```
T2 (after T1 commits):
Sees old row (REPEATABLE READ snapshot rule)
```

**Benefits:**

- Readers don't block writers
- Writers don't block readers (unless updating the same row)

#### **Cost:**

- Old row versions accumulate (need VACUUM)
- Slightly more complex internals

## **MySQL InnoDB: Different Approach**

**InnoDB uses MVCC too**, but with different defaults:

- **READ COMMITTED** by default
- **REPEATABLE READ** available, but with caveats (phantom reads possible)
- **SERIALIZABLE** uses locking (not true MVCC)

## **Locking: When Transactions Collide**

### **Implicit Locks**

**What happens when you UPDATE a row:**

```
UPDATE users SET balance = balance - 100 WHERE id = 1;
```

**Postgres acquires a row-level lock** on the row with `id = 1`.

**Other transactions:**

- Can still SELECT the row (MVCC)
- Can't UPDATE/DELETE it (blocked until first transaction commits/rolls back)

**This prevents lost updates.**

### **Explicit Locks: SELECT FOR UPDATE**

**Use case:** "Lock this row so I can update it later."

```
BEGIN;
SELECT balance FROM accounts WHERE id = 1 FOR UPDATE;
-- (balance = 100)
-- Now row is locked. Other transactions can't update it.
UPDATE accounts SET balance = 90 WHERE id = 1;
COMMIT;
```

**What FOR UPDATE does:**

- Acquires a row lock
- Other transactions trying to SELECT FOR UPDATE / UPDATE block
- Ensures no one changes the row between your SELECT and UPDATE

**When to use:**

- Read-modify-write patterns
- Preventing race conditions (e.g., inventory checks)

## SELECT FOR SHARE: Shared Lock

```
SELECT * FROM orders WHERE id = 123 FOR SHARE;
```

### What it does:

- Allows other transactions to read (and also lock FOR SHARE)
- Blocks UPDATEs/DELETEs

**Use case:** Rare. Mostly used with referential integrity.

## SKIP LOCKED: Non-Blocking Queue

```
SELECT * FROM jobs WHERE status = 'pending' ORDER BY created_at LIMIT 1 FOR UPDATE  
SKIP LOCKED;
```

### What it does:

- Tries to lock a row
- If row is already locked, **skips it** instead of blocking

**Use case:** Job queues (multiple workers, each grabs an available job).

## NOWAIT: Fail Fast

```
SELECT * FROM accounts WHERE id = 1 FOR UPDATE NOWAIT;
```

### What it does:

- Tries to lock a row
- If row is locked, **immediately errors** instead of blocking

**Use case:** User-facing apps where you don't want to wait (show "Resource busy" error).

## Deadlocks: When Transactions Block Each Other

### Scenario:

```
T1: BEGIN;  
T1: UPDATE accounts SET balance = balance - 100 WHERE id = 1; (locks row 1)  
  
T2: BEGIN;  
T2: UPDATE accounts SET balance = balance + 100 WHERE id = 2; (locks row 2)  
  
T1: UPDATE accounts SET balance = balance + 100 WHERE id = 2; (waits for T2)  
T2: UPDATE accounts SET balance = balance - 100 WHERE id = 1; (waits for T1)
```

**Result:** Deadlock. Each transaction waits for the other.

**Database action:** Detects deadlock, aborts one transaction.

**Error:** ERROR: deadlock detected

## Avoiding Deadlocks

1. **Consistent lock order:** Always lock rows in the same order (e.g., by ID ascending)
2. **Keep transactions short:** Less time holding locks = less chance of conflict
3. **Use timeouts:** `SET lock_timeout = '5s';`

## Retrying After Deadlock

Application code:

```
async function transferMoney(from, to, amount) {  
  let retries = 3;  
  while (retries > 0) {  
    try {  
      await db.query('BEGIN');  
      await db.query('UPDATE accounts SET balance = balance - $1 WHERE id = $2',  
[amount, from]);  
      await db.query('UPDATE accounts SET balance = balance + $1 WHERE id = $2',  
[amount, to]);  
      await db.query('COMMIT');  
      return;  
    } catch (err) {  
      await db.query('ROLLBACK');  
      if (err.code === '40P01') { // Deadlock  
        retries--;  
        if (retries === 0) throw err;  
        await sleep(100 * Math.random()); // Backoff  
      } else {  
        throw err;  
      }  
    }  
  }  
}
```

## Transaction Isolation in ORMs

### The Problem: ORMs Hide Transactions

Prisma example:

```
await prisma.user.update({ where: { id: 1 }, data: { balance: 100 } });
```

This is **auto-committed**. No explicit transaction.

If you need a transaction:

```
await prisma.$transaction(async (tx) => {  
  await tx.user.update({ where: { id: 1 }, data: { balance: 100 } });  
  await tx.user.update({ where: { id: 2 }, data: { balance: 200 } });  
});
```

## ORM Transaction Pitfalls

### Pitfall 1: Long transactions with network round-trips

```
await prisma.$transaction(async (tx) => {
  const user = await tx.user.findUnique({ where: { id: 1 } });

  // Expensive API call (network I/O)
  const result = await fetch('https://external-api.com/slow');

  await tx.user.update({ where: { id: 1 }, data: { balance: result.value } });
});
```

**Problem:** Transaction is open for **seconds** (API call), holding locks.

**Fix:** Do external I/O **outside** the transaction:

```
const result = await fetch('https://external-api.com/slow');

await prisma.$transaction(async (tx) => {
  await tx.user.update({ where: { id: 1 }, data: { balance: result.value } });
});
```

### Pitfall 2: Implicit isolation level

Prisma uses the database default (READ COMMITTED for Postgres). If you need REPEATABLE READ:

```
await prisma.$executeRaw`SET TRANSACTION ISOLATION LEVEL REPEATABLE READ`;
```

Or set it at the connection level.

### Pitfall 3: Forgotten rollback on error

```
await prisma.$transaction(async (tx) => {
  await tx.user.update({ where: { id: 1 }, data: { balance: 100 } });
  throw new Error('Oops'); // Prisma auto-rolls back
});
```

Prisma handles this, but **manual SQL doesn't always**:

```
const client = await pool.connect();
try {
  await client.query('BEGIN');
  await client.query('UPDATE ...');
  throw new Error('Oops'); // You must ROLLBACK!
  await client.query('COMMIT');
} catch (err) {
  await client.query('ROLLBACK'); // Don't forget this!
  throw err;
} finally {
```

```
    client.release();
}
```

## Real-World Transaction Patterns

### Pattern 1: Optimistic Locking (Version Column)

**Goal:** Prevent lost updates without explicit locks.

**Schema:**

```
CREATE TABLE products (
  id INT PRIMARY KEY,
  name TEXT,
  stock INT,
  version INT NOT NULL DEFAULT 0
);
```

**Application code:**

```
// Read product
const product = await db.query('SELECT * FROM products WHERE id = $1', [productId]);

// User modifies stock (in UI)
const newStock = product.stock - 1;

// Update with version check
const result = await db.query(
  'UPDATE products SET stock = $1, version = version + 1 WHERE id = $2 AND version = $3',
  [newStock, productId, product.version]
);

if (result.rowCount === 0) {
  throw new Error('Conflict: Product was modified by another user');
}
```

**How it works:**

- If another user updated the product, `version` has changed
- Your UPDATE matches 0 rows (version mismatch)
- You detect the conflict and retry or notify the user

**Benefits:**

- No locks held between read and write
- Scales well (no blocking)

### Pattern 2: SELECT FOR UPDATE for Inventory

**Goal:** Ensure stock doesn't go negative.

```

await db.query('BEGIN');

const product = await db.query(
  'SELECT stock FROM products WHERE id = $1 FOR UPDATE',
  [productId]
);

if (product.stock < quantity) {
  await db.query('ROLLBACK');
  throw new Error('Insufficient stock');
}

await db.query(
  'UPDATE products SET stock = stock - $1 WHERE id = $2',
  [quantity, productId]
);

await db.query('COMMIT');

```

**Benefits:** No one can change stock between SELECT and UPDATE.

### Pattern 3: Two-Phase Commit (Distributed Transactions)

**Goal:** Coordinate transactions across multiple databases.

**Example:** Transfer money between accounts in different databases.

#### Phase 1: Prepare

- DB1: BEGIN; UPDATE accounts SET balance = balance - 100; PREPARE TRANSACTION 'tx1';
- DB2: BEGIN; UPDATE accounts SET balance = balance + 100; PREPARE TRANSACTION 'tx2';

#### Phase 2: Commit or Rollback

- If both succeed: COMMIT PREPARED 'tx1'; COMMIT PREPARED 'tx2';
- If either fails: ROLLBACK PREPARED 'tx1'; ROLLBACK PREPARED 'tx2';

**Complexity:** High. Most apps avoid this (use eventual consistency instead).

## Practical Guidelines

### When to Use Transactions

#### ✓ Always:

- Multi-step writes that must be atomic (e.g., order + payment)
- Read-modify-write patterns (check stock, then decrement)

#### ✓ Sometimes:

- Batch inserts (faster, but not always necessary)
- Complex queries with consistency requirements

#### ✗ Never:

- Long-running transactions (holding locks)

- Transactions spanning external API calls
- Read-only queries (unless you need REPEATABLE READ snapshot)

## When to Use Higher Isolation Levels

**READ COMMITTED (default):** Most web apps.

**REPEATABLE READ:** Reports, analytics, multi-step reads.

**SERIALIZABLE:** Financial transactions, inventory with strict accuracy.

## Keep Transactions Short

**Bad:**

```
await db.query('BEGIN');
await processHeavyLogic(); // 10 seconds
await db.query('UPDATE ...');
await db.query('COMMIT');
```

**Good:**

```
await processHeavyLogic(); // Do this first
await db.query('BEGIN');
await db.query('UPDATE ...');
await db.query('COMMIT');
```

**Rule:** Only database I/O should be inside transactions.

## Key Takeaways

1. **Transactions are all-or-nothing.** COMMIT or ROLLBACK, no middle ground.
2. **Isolation levels trade consistency for performance.** READ COMMITTED is usually fine.
3. **MVCC (Postgres) allows readers and writers to coexist** without blocking.
4. **SELECT FOR UPDATE locks rows** for read-modify-write patterns.
5. **Deadlocks happen.** Detect and retry with backoff.
6. **ORMs hide transaction details.** Be explicit when you need atomicity.
7. **Keep transactions short.** Don't hold locks during external I/O.
8. **Use optimistic locking (version columns)** for long-running user interactions.
9. **SERIALIZABLE prevents all anomalies** but has lower concurrency.
10. **Vacuum regularly (Postgres)** to clean up old MVCC versions.

**Next up:** NULLs and three-valued logic—because NULL isn't what you think it is.