

# NULLs and Three-Valued Logic: The Source of Countless Bugs

## NULL Is Not a Value

NULL means "unknown" or "not applicable."

It's not:

- An empty string ( '' )
- Zero ( 0 )
- False ( false )

It's the absence of a value.

This distinction breaks intuition for developers coming from most programming languages.

## Three-Valued Logic

In most languages, boolean expressions are **binary**: TRUE or FALSE.

In SQL, they're **ternary**: TRUE, FALSE, or UNKNOWN.

```
SELECT * FROM users WHERE email = NULL; -- Returns 0 rows!
```

Why? `email = NULL` evaluates to UNKNOWN, not TRUE.

UNKNOWN rows are **filtered out** in WHERE clauses.

## Truth Tables

AND

| A       | B       | A AND B |
|---------|---------|---------|
| TRUE    | TRUE    | TRUE    |
| TRUE    | FALSE   | FALSE   |
| TRUE    | UNKNOWN | UNKNOWN |
| FALSE   | TRUE    | FALSE   |
| FALSE   | FALSE   | FALSE   |
| FALSE   | UNKNOWN | FALSE   |
| UNKNOWN | TRUE    | UNKNOWN |
| UNKNOWN | FALSE   | FALSE   |
| UNKNOWN | UNKNOWN | UNKNOWN |

Key insight: FALSE AND UNKNOWN = FALSE (short-circuit).

**OR**

| A       | B       | A OR B  |
|---------|---------|---------|
| TRUE    | TRUE    | TRUE    |
| TRUE    | FALSE   | TRUE    |
| TRUE    | UNKNOWN | TRUE    |
| FALSE   | TRUE    | TRUE    |
| FALSE   | FALSE   | FALSE   |
| FALSE   | UNKNOWN | UNKNOWN |
| UNKNOWN | TRUE    | TRUE    |
| UNKNOWN | FALSE   | UNKNOWN |
| UNKNOWN | UNKNOWN | UNKNOWN |

**Key insight:** TRUE OR UNKNOWN = TRUE (short-circuit).

**NOT**

| A       | NOT A   |
|---------|---------|
| TRUE    | FALSE   |
| FALSE   | TRUE    |
| UNKNOWN | UNKNOWN |

**Key insight:** NOT UNKNOWN = UNKNOWN (not FALSE!).

## Comparing to NULL: Always UNKNOWN

### Equality

```
SELECT * FROM users WHERE email = NULL; -- Wrong!
```

**Result:** 0 rows, always.

**Why?** email = NULL is UNKNOWN (even if email IS NULL).

**Correct:**

```
SELECT * FROM users WHERE email IS NULL;
```

IS NULL is a special operator that returns TRUE/FALSE (not UNKNOWN).

### Inequality

```
SELECT * FROM users WHERE email != NULL; -- Wrong!
```

**Result:** 0 rows.

**Why?** `email != NULL` is UNKNOWN.

**Correct:**

```
SELECT * FROM users WHERE email IS NOT NULL;
```

## The Trap: NULL in Subqueries

```
SELECT * FROM orders WHERE user_id NOT IN (SELECT id FROM users WHERE active = false);
```

**If any user.id is NULL, this returns 0 rows!**

**Why?**

- `user_id NOT IN (1, 2, NULL)` expands to:
- `user_id != 1 AND user_id != 2 AND user_id != NULL`
- `user_id != NULL` is UNKNOWN
- `TRUE AND TRUE AND UNKNOWN = UNKNOWN`
- UNKNOWN rows are filtered out

**Fix:** Exclude NULLs explicitly:

```
SELECT * FROM orders WHERE user_id NOT IN (SELECT id FROM users WHERE active = false AND id IS NOT NULL);
```

Or use `NOT EXISTS` (safer):

```
SELECT * FROM orders o
WHERE NOT EXISTS (
    SELECT 1 FROM users u WHERE u.id = o.user_id AND u.active = false
);
```

## NULL in Aggregates

### COUNT: The Special Case

```
SELECT
    COUNT(*),      -- Counts rows
    COUNT(email), -- Counts non-NULL emails
    COUNT(*)
FROM users;
```

**If `email` is nullable:**

- `COUNT(*)` = total rows (including NULLs)
- `COUNT(email)` = rows with non-NULL email

**Different values!**

### SUM, AVG: Ignore NULLs

```
SELECT AVG(rating) FROM reviews;
```

### Data: `rating`

5 4 NULL 3

```
**Result:** `AVG(rating) = 4` (not 3, not 2.4).

**Why?** NULLs are excluded from both numerator and denominator.

- Sum: 5 + 4 + 3 = 12
- Count: 3 (not 4)
- Avg: 12 / 3 = 4

### MIN, MAX: Ignore NULLs

```sql
SELECT MIN(price), MAX(price) FROM products;
```

**If all prices are NULL:** Returns NULL.

### NULL Propagation in Expressions

```
SELECT 1 + NULL;          -- Returns NULL
SELECT 'Hello' || NULL; -- Returns NULL
SELECT NULL * 100;       -- Returns NULL
```

**Rule:** Any arithmetic or string operation involving NULL returns NULL (except COALESCE).

### COALESCE: The NULL Handler

**COALESCE returns the first non-NULL value:**

```
SELECT COALESCE(email, phone, 'No contact') AS contact FROM users;
```

**If `email` is NULL but `phone` is not:** Returns `phone`.

**If both are NULL:** Returns 'No contact'.

### Common Uses

#### Default Values

```
SELECT name, COALESCE(discount, 0) AS discount FROM products;
```

If discount is NULL: Returns 0.

### Safe Aggregation

```
SELECT user_id, COALESCE(SUM(amount), 0) AS total FROM transactions GROUP BY user_id;
```

If a user has no transactions (or all amounts are NULL): Returns 0 instead of NULL.

### COALESCE vs IFNULL / NVL

SQL Standard: COALESCE(a, b, c)

MySQL: IFNULL(a, b) (only 2 arguments)

Oracle: NVL(a, b) (only 2 arguments)

Use COALESCE for portability.

### NULLIF: The Reverse of COALESCE

NULLIF returns NULL if two values are equal:

```
SELECT NULLIF(value, 0) FROM table;
```

If value = 0: Returns NULL.

If value != 0: Returns value.

Use case: Avoiding division by zero:

```
SELECT total / NULLIF(count, 0) AS average FROM stats;
```

If count = 0 , division returns NULL (not error).

### NULL in JOINS

#### LEFT JOIN: NULLs Indicate Missing Matches

```
SELECT u.name, o.id AS order_id
FROM users u
LEFT JOIN orders o ON u.id = o.user_id;
```

If a user has no orders:

| name  | order_id |
|-------|----------|
| ----- | -----    |

```
Alice | 101
Bob   | NULL      ← Bob has no orders
```

### Filtering NULLs in WHERE turns LEFT JOIN into INNER JOIN:

```
SELECT u.name, o.id
FROM users u
LEFT JOIN orders o ON u.id = o.user_id
WHERE o.id IS NOT NULL; -- Excludes users without orders
```

This is effectively an INNER JOIN. Use INNER JOIN instead for clarity.

### NULL as a Join Key: Doesn't Match

```
SELECT * FROM users u
JOIN orders o ON u.optional_field = o.optional_field;
```

If both fields are NULL: They don't match.

Why? `NULL = NULL` is UNKNOWN, not TRUE.

Workaround (if you want NULL to match NULL):

```
ON (u.optional_field = o.optional_field OR (u.optional_field IS NULL AND
o.optional_field IS NULL))
```

Or use `IS NOT DISTINCT FROM` (Postgres):

```
ON u.optional_field IS NOT DISTINCT FROM o.optional_field
```

## NULL in ORDER BY

### Default Behavior

Postgres: NULLs sort **last** (ascending) or **first** (descending).

MySQL: NULLs sort **first** (ascending) or **last** (descending).

Example:

```
SELECT name, rating FROM products ORDER BY rating;
```

Postgres:

| name | rating |
|------|--------|
| A    | 1      |
| B    | 2      |
| C    | NULL   |

**MySQL:**

| name | rating |
|------|--------|
| C    | NULL   |
| A    | 1      |
| B    | 2      |

## Explicit NULL Handling

**Postgres:**

```
ORDER BY rating NULLS FIRST;  
ORDER BY rating NULLS LAST;
```

**MySQL:** Use conditional sorting:

```
ORDER BY rating IS NULL, rating; -- NULLs last
```

## NULL in DISTINCT

**NULL values are considered duplicates:**

```
SELECT DISTINCT email FROM users;
```

**If 3 users have NULL email:** Only one NULL appears in the result.

This is inconsistent with `NULL = NULL` (which is UNKNOWN), but pragmatic.

## NULL in UNIQUE Constraints

**Postgres: Multiple NULLs Allowed**

```
CREATE TABLE products (  
    id INT PRIMARY KEY,  
    sku TEXT UNIQUE  
);  
  
INSERT INTO products (id, sku) VALUES (1, NULL);  
INSERT INTO products (id, sku) VALUES (2, NULL); -- Allowed!
```

**Why?** `NONE != NULL`, so no uniqueness violation.

**SQL Standard: Same behavior (Postgres follows it)**

**MySQL: Depends on storage engine**

- **InnoDB:** Multiple NULLs allowed (like Postgres).
- **MyISAM:** Multiple NULLs allowed.

## Workaround: Use COALESCE with a sentinel value

```
CREATE UNIQUE INDEX idx_products_sku ON products (COALESCE(sku, ''));
```

Now only one NULL (treated as '') is allowed.

## NULL in CHECK Constraints

CHECK constraints allow NULL:

```
CREATE TABLE users (
    age INT CHECK (age >= 18)
);

INSERT INTO users (age) VALUES (NULL); -- Allowed!
```

Why? `NULL >= 18` is UNKNOWN, and CHECK constraints allow UNKNOWN.

If you want to enforce non-NULL:

```
age INT NOT NULL CHECK (age >= 18)
```

## Interview Trap Examples

### Trap 1: Counting NULLs

Question: How many users have no email?

Wrong:

```
SELECT COUNT(email) FROM users WHERE email = NULL;
```

Result: 0 (always).

Right:

```
SELECT COUNT(*) FROM users WHERE email IS NULL;
```

Or:

```
SELECT COUNT(*) - COUNT(email) FROM users;
```

### Trap 2: NOT IN with NULLs

Question: Find orders not placed by inactive users.

Wrong:

```
SELECT * FROM orders WHERE user_id NOT IN (SELECT id FROM users WHERE active = false);
```

If any inactive user has `id = NULL`, this returns 0 rows!

Right:

```
SELECT * FROM orders o
WHERE NOT EXISTS (
    SELECT 1 FROM users u WHERE u.id = o.user_id AND u.active = false
);
```

### Trap 3: NULL in Arithmetic

Question: Calculate total revenue.

Naive:

```
SELECT SUM(quantity * price) FROM order_items;
```

Problem: If any `quantity` or `price` is NULL, that row contributes NULL (ignored by SUM, but might not be intended).

Safe:

```
SELECT SUM(COALESCE(quantity, 0) * COALESCE(price, 0)) FROM order_items;
```

Or enforce NOT NULL constraints.

### Trap 4: WHERE col != value Doesn't Include NULLS

Question: Find all products not in category 'Books'.

Wrong:

```
SELECT * FROM products WHERE category != 'Books';
```

Problem: Products with `category = NULL` are excluded!

Why? `NULL != 'Books'` is UNKNOWN.

Right:

```
SELECT * FROM products WHERE category != 'Books' OR category IS NULL;
```

Or:

```
SELECT * FROM products WHERE NOT (category = 'Books');
```

(But this is confusing. Explicit IS NULL is clearer.)

## Trap 5: MAX/MIN with All NULLs

**Data:** ` price

NULL NULL

```
```sql
SELECT MAX(price) FROM products;
```

**Result:** NULL (not 0, not error).

**Safe:**

```
SELECT COALESCE(MAX(price), 0) FROM products;
```

## Practical NULL Handling Strategies

### Strategy 1: Avoid NULLs When Possible

**Use NOT NULL + DEFAULT:**

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    name TEXT NOT NULL,
    email TEXT NOT NULL,
    phone TEXT NOT NULL DEFAULT '',
    active BOOLEAN NOT NULL DEFAULT true
);
```

**Benefits:**

- No NULL checks needed
- Clearer semantics
- Fewer bugs

**When NULL is appropriate:**

- Optional foreign keys
- Truly unknown values (e.g., end\_date for ongoing events)

### Strategy 2: Use COALESCE for Defaults

```
SELECT name, COALESCE(discount, 0) AS discount FROM products;
```

**Instead of checking IS NULL everywhere.**

### Strategy 3: Use NOT EXISTS Instead of NOT IN

**Safer:**

```
WHERE NOT EXISTS (SELECT 1 FROM other_table WHERE ...)
```

Than:

```
WHERE id NOT IN (SELECT id FROM other_table)
```

## Strategy 4: Explicit NULL Handling in Application

Don't assume SQL will do what you expect:

```
const result = await db.query('SELECT email FROM users WHERE id = $1', [userId]);

if (result.rows[0].email === null) {
    // Handle NULL explicitly
    email = 'no-email@example.com';
} else {
    email = result.rows[0].email;
}
```

## Strategy 5: Use Partial Indexes to Exclude NULLS

```
CREATE INDEX idx_users_phone ON users(phone) WHERE phone IS NOT NULL;
```

Benefits:

- Smaller index (doesn't index NULLs)
- Faster queries on non-NULL values

## NULL vs Empty String vs Zero

### Empty String ( '')

**Is a value.** Not the same as NULL.

```
SELECT '' IS NULL;          -- FALSE
SELECT '' = '';            -- TRUE
SELECT LENGTH('');         -- 0
```

**Use case:** Explicitly empty text fields.

### Zero ( 0 )

**Is a value.** Not the same as NULL.

```
SELECT 0 IS NULL;          -- FALSE
SELECT 0 = 0;              -- TRUE
SELECT 0 + 10;             -- 10
```

**Use case:** Numeric zero (count, balance, etc.).

## NULL

**Is not a value.**

```
SELECT NULL IS NULL;      -- TRUE
SELECT NULL = NULL;       -- UNKNOWN (!)
SELECT NULL + 10;         -- NULL
```

**Use case:** Unknown or inapplicable values.

## Key Takeaways

1. **NULL means "unknown," not a value.** It's not 0, not '', not false.
2. **Comparing to NULL with = is always UNKNOWN.** Use IS NULL / IS NOT NULL .
3. **Three-valued logic:** TRUE, FALSE, UNKNOWN. WHERE filters out UNKNOWN.
4. **COUNT(\*) counts all rows; COUNT(column) counts non-NULL values.**
5. **Aggregates (SUM, AVG, MIN, MAX) ignore NULLs.**
6. **NULL propagates in expressions:** 1 + NULL = NULL .
7. **COALESCE returns the first non-NULL value.**
8. **NOT IN with NULLs is dangerous.** Use NOT EXISTS.
9. **LEFT JOIN produces NULLs for missing matches.** Don't filter them in WHERE unless you mean INNER JOIN.
10. **Avoid NULLs when possible** with NOT NULL + DEFAULT.

**Next up:** Constraints and schema design—how to make the database enforce correctness.