

# Constraints and Schema Design: Let the Database Do the Work

## The Philosophy: Constraints > Application Logic

**Bad approach:** Validate data in application code.

```
if (email.includes('@')) {  
    await db.query('INSERT INTO users (email) VALUES ($1)', [email]);  
}
```

**Problem:** Another app (or SQL client) can bypass validation.

**Good approach:** Enforce in the database.

```
CREATE TABLE users (  
    email TEXT NOT NULL CHECK (email ~ '^.+@.+\\..+$')  
) ;
```

**Benefits:**

- **Single source of truth:** One place to enforce rules
- **Can't be bypassed:** Even raw SQL respects constraints
- **Self-documenting:** Schema describes invariants

**Rule:** If something **must** be true, enforce it in the schema.

## NOT NULL: The Foundation

### Default Behavior: Columns Are Nullable

```
CREATE TABLE users (  
    name TEXT  
) ;  
  
INSERT INTO users (name) VALUES (NULL); -- Allowed
```

### Enforce Non-NULL

```
CREATE TABLE users (  
    name TEXT NOT NULL  
) ;  
  
INSERT INTO users (name) VALUES (NULL); -- ERROR: null value in column "name"
```

### When to Use NOT NULL

**Use on almost everything.** NULL complicates logic.

### Especially:

- Primary keys (always)
- Foreign keys (unless the relationship is optional)
- Business-critical fields (email, username, etc.)

### Don't use when:

- Value is genuinely optional (e.g., middle name)
- Foreign key is optional (e.g., manager\_id for CEO)

## Setting Defaults

```
CREATE TABLE users (
    active BOOLEAN NOT NULL DEFAULT true,
    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);
```

### INSERT without specifying values:

```
INSERT INTO users (name) VALUES ('Alice');
```

**Result:** active = true , created\_at = current timestamp .

## PRIMARY KEY: Unique Row Identifier

### What It Does

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    name TEXT NOT NULL
);
```

### Enforces:

1. **Uniqueness:** No two rows with the same `id`
2. **Non-NULL:** `id` cannot be NULL

### Equivalent to:

```
CREATE TABLE users (
    id SERIAL UNIQUE NOT NULL,
    name TEXT NOT NULL
);
```

But `PRIMARY KEY` is clearer and has special meaning (e.g., for foreign keys).

## Choosing Primary Keys

### Option 1: Serial / Auto-Increment (Surrogate Key)

```
-- Postgres:  
id SERIAL PRIMARY KEY  
  
-- MySQL:  
id INT AUTO_INCREMENT PRIMARY KEY
```

#### Benefits:

- Simple, guaranteed unique
- Small (integer)
- No business meaning (won't change)

#### Drawbacks:

- Sequential (predictable, can leak count)
- Not globally unique (per-table)

### Option 2: UUID (Surrogate Key)

```
id UUID PRIMARY KEY DEFAULT gen_random_uuid()
```

#### Benefits:

- Globally unique (across tables, databases, servers)
- Non-sequential (less predictable)

#### Drawbacks:

- Larger (16 bytes vs 4/8 for INT/BIGINT)
- Slightly slower inserts (random B-tree inserts)

### Option 3: Natural Key

```
CREATE TABLE countries (  
    code CHAR(2) PRIMARY KEY, -- 'US', 'CA', etc.  
    name TEXT NOT NULL  
) ;
```

#### Benefits:

- Self-explanatory (no need to join to get the code)

#### Drawbacks:

- Business keys can change (email, username, etc.)
- Composite keys are awkward

**Best practice:** Use surrogate keys (SERIAL or UUID) unless the natural key is immutable and stable.

## Composite Primary Keys

```
CREATE TABLE order_items (  
    order_id INT,  
    product_id INT,
```

```
quantity INT NOT NULL,  
PRIMARY KEY (order_id, product_id)  
);
```

**Enforces:** Each (order\_id, product\_id) pair is unique.

**Use case:** Junction tables (many-to-many relationships).

## UNIQUE Constraints: Enforce Uniqueness

```
CREATE TABLE users (  
    id SERIAL PRIMARY KEY,  
    email TEXT NOT NULL UNIQUE  
);
```

**Equivalent to:**

```
CREATE TABLE users (  
    id SERIAL PRIMARY KEY,  
    email TEXT NOT NULL,  
    CONSTRAINT users_email_unique UNIQUE (email)  
);
```

## Multiple NULLs Are Allowed

```
CREATE TABLE users (  
    id SERIAL PRIMARY KEY,  
    phone TEXT UNIQUE  
);  
  
INSERT INTO users (phone) VALUES (NULL);  
INSERT INTO users (phone) VALUES (NULL); -- Allowed!
```

**Why?** `NULL != NULL`, so no uniqueness violation.

**If you want only one NULL:** Use a partial unique index (Postgres):

```
CREATE UNIQUE INDEX users_phone_unique ON users (phone) WHERE phone IS NOT NULL;
```

## Composite UNIQUE Constraints

```
CREATE TABLE reservations (  
    id SERIAL PRIMARY KEY,  
    user_id INT NOT NULL,  
    date DATE NOT NULL,  
    UNIQUE (user_id, date)  
);
```

**Enforces:** Each user can have at most one reservation per date.

## FOREIGN KEY: Referential Integrity

### Basic Syntax

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT NOT NULL REFERENCES users(id)
);
```

**Enforces:**

1. `user_id` must exist in `users.id` (or be NULL, if nullable)
2. Can't delete a user if they have orders (unless cascade rules apply)

### ON DELETE / ON UPDATE Actions

#### ON DELETE CASCADE

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT NOT NULL REFERENCES users(id) ON DELETE CASCADE
);
```

**Behavior:** If a user is deleted, all their orders are deleted too.

**Use case:** Dependent data that should be cleaned up (e.g., user sessions, logs).

**Warning:** Powerful, use carefully (accidental deletes cascade).

#### ON DELETE SET NULL

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT REFERENCES users(id) ON DELETE SET NULL
);
```

**Behavior:** If a user is deleted, `user_id` in orders becomes NULL.

**Use case:** Historical data where you want to preserve orders but anonymize the user.

#### ON DELETE RESTRICT (Default)

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT NOT NULL REFERENCES users(id) ON DELETE RESTRICT
);
```

**Behavior:** Can't delete a user if they have orders (error).

**Use case:** Prevent accidental deletion.

### ON DELETE SET DEFAULT

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT NOT NULL DEFAULT 1 REFERENCES users(id) ON DELETE SET DEFAULT
);
```

**Behavior:** If a user is deleted, `user_id` becomes the default (e.g., "Unknown User").

**Use case:** Rare. Setting a sentinel value.

### ON UPDATE: Same Options

Usually less important (primary keys rarely change), but `ON UPDATE CASCADE` can be useful:

```
REFERENCES users(id) ON UPDATE CASCADE
```

If `users.id` changes, `orders.user_id` updates automatically.

### Composite Foreign Keys

```
CREATE TABLE order_items (
    order_id INT,
    product_id INT,
    quantity INT NOT NULL,
    FOREIGN KEY (order_id, product_id) REFERENCES order_product_combos(order_id,
    product_id)
);
```

Rare, but useful for complex relationships.

### Deferrable Constraints (Postgres)

By default, foreign keys are checked immediately (on each INSERT/UPDATE).

**Problem:** Circular dependencies.

```
-- Can't insert user without org, can't insert org without user!
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    org_id INT NOT NULL REFERENCES orgs(id)
);

CREATE TABLE orgs (
    id SERIAL PRIMARY KEY,
    owner_id INT NOT NULL REFERENCES users(id)
);
```

**Solution: Deferrable constraints:**

```

CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    org_id INT REFERENCES orgs(id) DEFERRABLE INITIALLY DEFERRED
);

CREATE TABLE orgs (
    id SERIAL PRIMARY KEY,
    owner_id INT REFERENCES users(id) DEFERRABLE INITIALLY DEFERRED
);

```

**Behavior:** Constraints are checked at **COMMIT** time, not immediately.

**Use case:** Rare. Mostly for complex migrations or circular dependencies.

## CHECK Constraints: Custom Validation

### Basic Syntax

```

CREATE TABLE users (
    age INT CHECK (age >= 18)
);

```

**Enforces:** Age must be at least 18.

```
INSERT INTO users (age) VALUES (17); -- ERROR: check constraint violated
```

### Named Constraints

```

CREATE TABLE products (
    price NUMERIC CHECK (price > 0),
    discount NUMERIC CHECK (discount >= 0 AND discount <= 1),
    CONSTRAINT price_after_discount CHECK (price * (1 - discount) > 0)
);

```

**Naming constraints helps with error messages:**

```
ERROR: new row violates check constraint "price_after_discount"
```

### Multi-Column Checks

```

CREATE TABLE events (
    start_date DATE NOT NULL,
    end_date DATE,
    CHECK (end_date IS NULL OR end_date >= start_date)
);

```

**Enforces:** If `end_date` is set, it must be after `start_date`.

## NULLs Pass CHECK Constraints

```
CREATE TABLE users (
    age INT CHECK (age >= 18)
);

INSERT INTO users (age) VALUES (NULL); -- Allowed!
```

**Why?** `NULL >= 18` is UNKNOWN, and CHECK allows UNKNOWN.

**To enforce non-NULL:**

```
age INT NOT NULL CHECK (age >= 18)
```

## Limitations of CHECK

**What you can't do:**

- Reference other tables (use triggers or application logic)
- Use subqueries (Postgres allows it but it's not efficient)

**What you can do:**

- Column-level or table-level checks
- Use functions (e.g., `CHECK (email ~ '^.+@.+')`)

## EXCLUSION Constraints (Postgres Only)

**Use case:** Ensure no overlapping ranges (e.g., room reservations).

```
CREATE TABLE reservations (
    room_id INT,
    reserved_at TSTZRANGE,
    EXCLUDE USING GIST (room_id WITH =, reserved_at WITH &&)
);
```

**Enforces:** No two reservations for the same room can overlap in time.

**Example:**

```
INSERT INTO reservations VALUES (1, '[2024-01-01 10:00, 2024-01-01 12:00)');
INSERT INTO reservations VALUES (1, '[2024-01-01 11:00, 2024-01-01 13:00>'); -- ERROR
```

**Why it's powerful:** Prevents double-booking at the database level.

**Requires:** `btree_gist` extension for `=` operator.

## Default Values

### Static Defaults

```
CREATE TABLE users (
    status TEXT NOT NULL DEFAULT 'active',
    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);
```

## Dynamic Defaults (Functions)

```
CREATE TABLE users (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    created_at TIMESTAMPTZ DEFAULT NOW()
);
```

**Note:** `NOW()` evaluates once per transaction, not per row.

For per-row timestamps (Postgres):

```
created_at TIMESTAMPTZ DEFAULT CLOCK_TIMESTAMP()
```

Or use triggers.

## Generated Columns (Postgres 12+, MySQL 5.7+)

Computed automatically:

```
CREATE TABLE products (
    price NUMERIC NOT NULL,
    tax_rate NUMERIC NOT NULL,
    price_with_tax NUMERIC GENERATED ALWAYS AS (price * (1 + tax_rate)) STORED
);
```

Behavior:

- `price_with_tax` is computed from `price` and `tax_rate`
- Updated automatically when `price` or `tax_rate` changes
- Can be indexed

STORED vs VIRTUAL:

- **STORED:** Computed on write, stored physically (faster reads)
- **VIRTUAL:** Computed on read (saves space, slower reads; not in Postgres 15-, only MySQL)

## Naming Conventions

Good constraint names help debugging:

```
-- Bad:
CONSTRAINT users_check CHECK (age >= 18)

-- Good:
CONSTRAINT users_age_min_18 CHECK (age >= 18)
```

### Error message:

```
ERROR: new row violates check constraint "users_age_min_18"
```

Clear what failed.

### Standard Naming Pattern

```
{table}_{column}_{type}
```

#### Examples:

- users\_email\_unique
- orders\_user\_id\_fkey
- products\_price\_check

## Constraints vs Application Logic

### Use Constraints When:

#### Invariants that must never be violated:

- Email must be unique
- Price must be positive
- Foreign keys must exist

#### Data integrity:

- Referential integrity (foreign keys)
- Uniqueness
- Not null

### Use Application Logic When:

#### Business rules that change frequently:

- Password strength requirements
- Coupon discount rules

#### Complex validation requiring external data:

- "User must have valid credit card" (external API)

#### Performance-sensitive checks:

- "User must not have >100 orders" (expensive query)

**Rule of thumb:** If it's a **hard invariant**, use a constraint. If it's a **soft rule**, use application logic.

## Schema Design Patterns

### Pattern 1: Soft Deletes

Instead of DELETE, mark as deleted:

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    email TEXT NOT NULL UNIQUE,
    deleted_at TIMESTAMPTZ
);
```

"Delete" user:

```
UPDATE users SET deleted_at = NOW() WHERE id = 123;
```

Query active users:

```
SELECT * FROM users WHERE deleted_at IS NULL;
```

Benefits:

- Preserve history
- Undo deletions
- Audit trail

Drawbacks:

- Complicates queries (always filter deleted\_at IS NULL )
- UNIQUE constraints don't respect soft deletes (email collision)

Fix unique constraint issue:

```
CREATE UNIQUE INDEX users_email_unique ON users (email) WHERE deleted_at IS NULL;
```

## Pattern 2: Audit Columns

```
CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    user_id INT NOT NULL REFERENCES users(id),
    total NUMERIC NOT NULL,
    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    updated_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    created_by INT REFERENCES users(id),
    updated_by INT REFERENCES users(id)
);
```

Update updated\_at with a trigger:

```
CREATE OR REPLACE FUNCTION update_updated_at()
RETURNS TRIGGER AS $$
BEGIN
    NEW.updated_at = NOW();
    RETURN NEW;
END;
```

```

$$ LANGUAGE plpgsql;

CREATE TRIGGER orders_updated_at
BEFORE UPDATE ON orders
FOR EACH ROW
EXECUTE FUNCTION update_updated_at();

```

### Pattern 3: Enum Types (Postgres)

**Instead of CHECK constraints:**

```

CREATE TYPE order_status AS ENUM ('pending', 'processing', 'shipped', 'delivered',
'cancelled');

CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    status order_status NOT NULL DEFAULT 'pending'
);

```

**Benefits:**

- Type-safe (typos are errors)
- Self-documenting
- Can't insert invalid values

**Drawbacks:**

- Hard to change (adding enum values requires ALTER TYPE)
- Postgres-specific (MySQL has ENUM, but it's less strict)

**Alternative: Lookup table:**

```

CREATE TABLE order_statuses (
    id SERIAL PRIMARY KEY,
    name TEXT NOT NULL UNIQUE
);

CREATE TABLE orders (
    id SERIAL PRIMARY KEY,
    status_id INT NOT NULL REFERENCES order_statuses(id)
);

```

More flexible, but requires joins.

### Pattern 4: JSONB for Semi-Structured Data

**Use case:** Flexible attributes.

```

CREATE TABLE products (
    id SERIAL PRIMARY KEY,
    name TEXT NOT NULL,

```

```
    metadata JSONB NOT NULL DEFAULT '{}'  
);
```

#### Insert:

```
INSERT INTO products (name, metadata) VALUES  
( 'Laptop', '{"brand": "Dell", "ram_gb": 16});
```

#### Query:

```
SELECT * FROM products WHERE metadata->>'brand' = 'Dell';
```

#### Index JSONB:

```
CREATE INDEX idx_products_metadata ON products USING GIN(metadata);
```

#### When to use:

- Attributes vary by product type
- Frequent schema changes

#### When NOT to use:

- Core business data (use proper columns)
- Need strong typing

## Database Portability vs Features

### Postgres-Specific Features

- EXCLUDE constraints
- ENUM types
- DEFERRABLE constraints
- ARRAY types
- JSONB with GIN indexes

**If you need portability:** Stick to ANSI SQL (PRIMARY KEY, FOREIGN KEY, UNIQUE, CHECK, NOT NULL).

**If you're committed to Postgres:** Use these features—they're powerful.

## Practical Guidelines

### Always Index Foreign Keys

```
CREATE TABLE orders (  
    id SERIAL PRIMARY KEY,  
    user_id INT NOT NULL REFERENCES users(id)  
);  
  
CREATE INDEX idx_orders_user_id ON orders(user_id);
```

**Why?** JOINS on foreign keys are common. Indexes make them fast.

## Use NOT NULL Liberally

**Default to NOT NULL unless the field is truly optional.**

## Validate in Both App and Database

**Layered defense:**

1. **Client-side:** Immediate feedback (UX)
2. **Server-side:** Security (don't trust client)
3. **Database:** Final enforcement (can't be bypassed)

## Use Migrations, Not Manual DDL

**Good:**

```
// Prisma migration
await prisma.$executeRaw`  
  ALTER TABLE users ADD CONSTRAINT users_age_min_18 CHECK (age >= 18);  
`;
```

**Bad:**

```
-- Manual SQL (no version control, no rollback)
ALTER TABLE users ADD CONSTRAINT ...;
```

**Use tools:** Prisma Migrate, Drizzle Kit, Flyway, Liquibase.

## Key Takeaways

1. **NOT NULL + defaults eliminate most NULL issues.**
2. **PRIMARY KEY enforces uniqueness + NOT NULL.** Use surrogate keys (SERIAL/UUID).
3. **FOREIGN KEY enforces referential integrity.** Use ON DELETE/UPDATE actions wisely.
4. **CHECK constraints enforce custom rules.** Use for simple validations.
5. **UNIQUE constraints prevent duplicates.** Multiple NULLs are allowed (usually fine).
6. **Naming constraints helps debugging.** Follow a consistent pattern.
7. **Constraints > application logic** for hard invariants.
8. **Always index foreign keys** for JOIN performance.
9. **Use migrations for schema changes.** Never manual DDL in production.
10. **Postgres-specific features are powerful.** Use them if you're committed to Postgres.

**Next up:** Postgres-specific features you should know about.